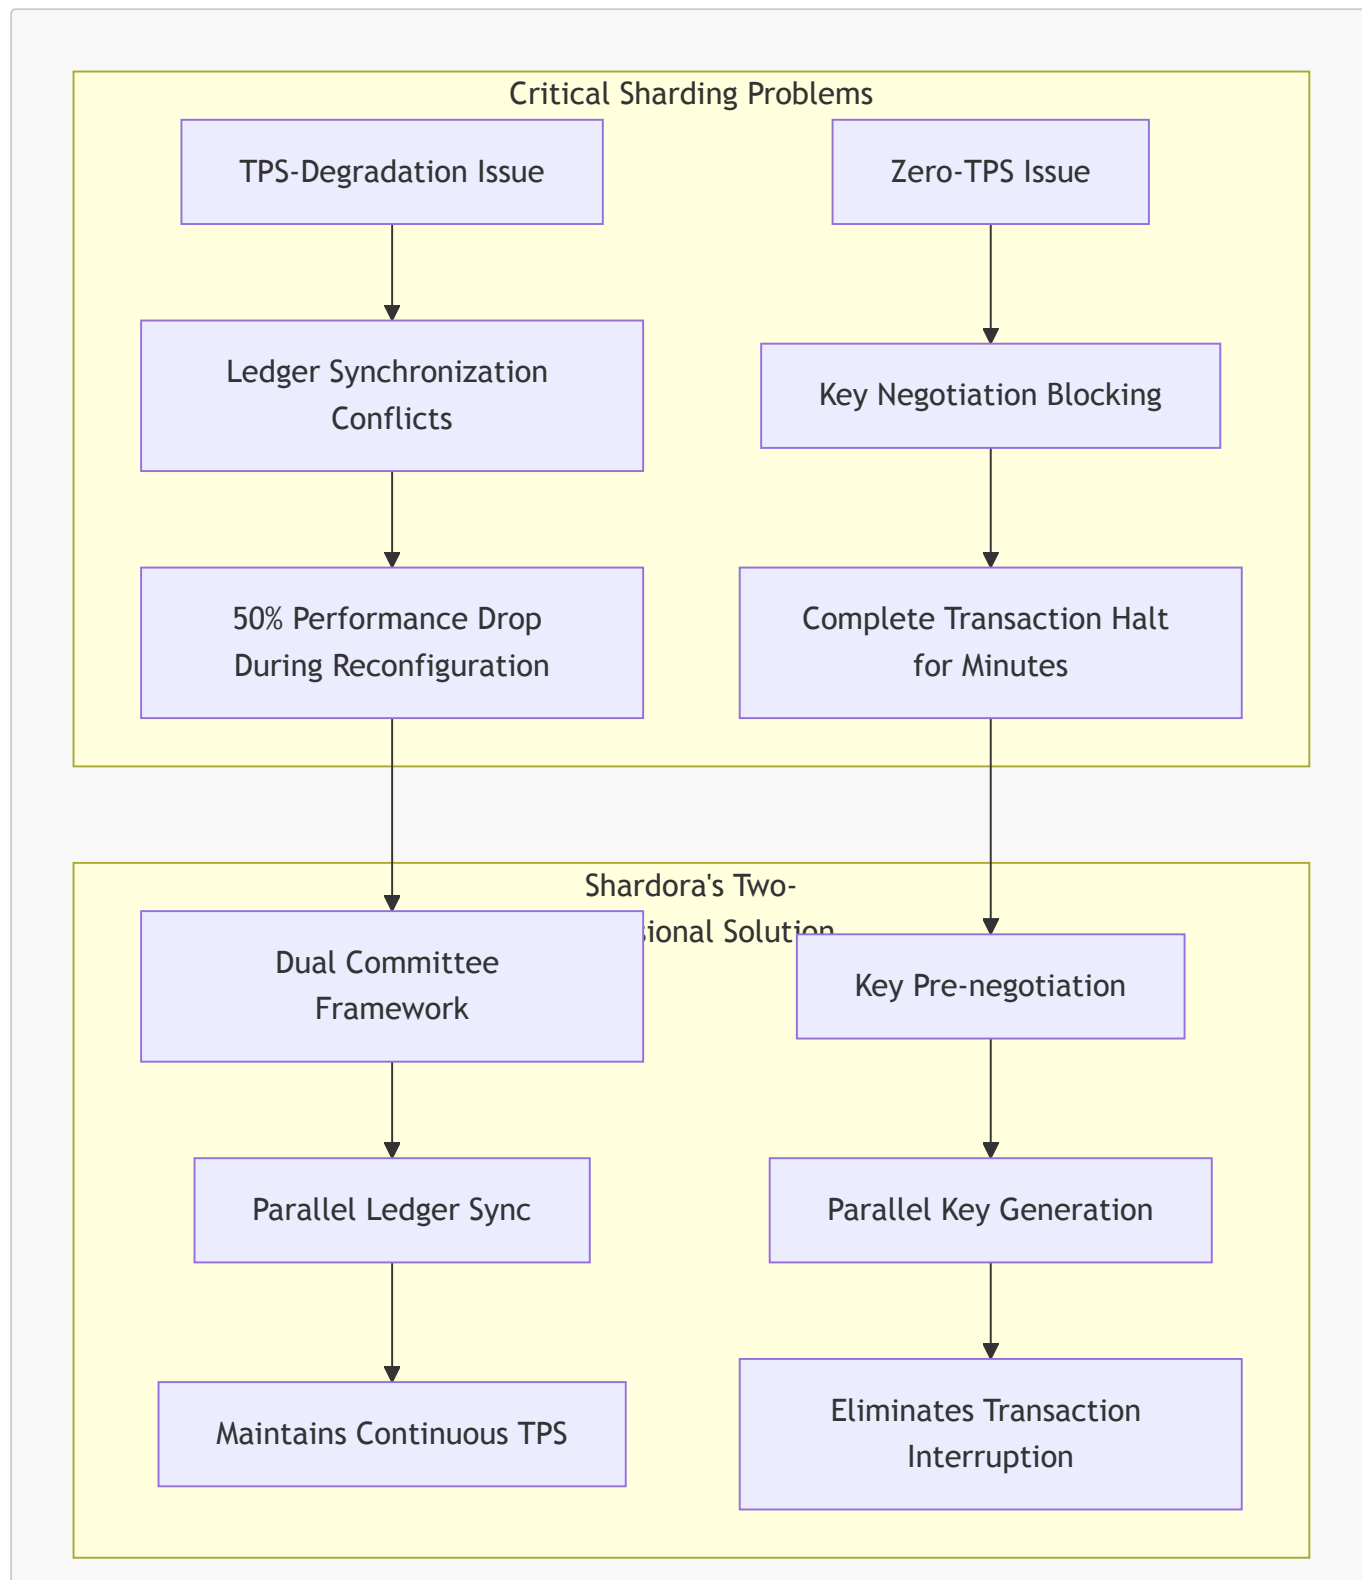


1. Introduction

Blockchain technology has revolutionized decentralized systems, but sharding implementations face two critical bottlenecks that severely limit their practical scalability. While sharding promises to increase throughput by processing transactions in parallel across multiple shards, existing systems suffer from performance degradation during essential network operations.

Shardora addresses these fundamental challenges through innovative two-dimensional parallelism, specifically targeting the **TPS-Degradation** and **Zero-TPS** issues that plague current sharded blockchain architectures.



2. Problems

The pursuit of scalability in blockchain networks, particularly through sharding, introduces a new set of complex challenges that can undermine the very performance benefits they aim to deliver. LogaShard specifically addresses two critical issues that emerge during the dynamic reconfiguration inherent in sharded architectures: TPS-Degradation and Zero-TPS.

2.1. Detailed Analysis of the TPS-Degradation Issue

The TPS-Degradation issue represents a significant limitation in existing sharding systems, directly impacting the network's ability to maintain consistent high throughput.

Root Causes: Ledger Synchronization Conflicts during Reconfiguration

The primary cause of TPS-Degradation is identified as "ledger synchronization conflicts during shard reconfiguration".¹¹ In dynamic sharding systems, nodes are periodically shuffled or reassigned among different shards to enhance security and prevent malicious actors from gaining prolonged control over a specific shard.¹² This auto-rotation mechanism introduces unpredictability, making it difficult for attackers to collude or target specific network segments.²³

However, this necessary security measure comes with a performance cost. When nodes are newly assigned to a shard, or when existing nodes are shuffled, they are required to fully synchronize the ledger state of their newly assigned shard before they can actively participate in the consensus process and validate transactions.¹² This synchronization process is crucial for ensuring data consistency and integrity across the network. Without a complete and up-to-date view of the shard's state, a validator cannot reliably process transactions or contribute to consensus. This period of intensive data synchronization, effectively a "cold start" problem for newly assigned validators, directly conflicts with the network's ability to maintain continuous transaction processing. During this "warm-up" phase, the effective number of active validators contributing to the network's TPS is reduced, leading to a noticeable drop in the overall throughput. This challenge is particularly pronounced in large-scale networks with frequent reconfigurations.

2.2. Detailed Analysis of the Zero-TPS Issue

Beyond mere degradation, sharding systems can experience a complete halt in transaction processing, a phenomenon termed the "Zero-TPS issue."

Root Causes: Disruptions from Key Negotiation during Reconfiguration

The Zero-TPS issue is primarily caused by "disruptions in transaction processing due to key negotiation" during shard reconfiguration. In many sharded blockchain architectures, especially those employing threshold signature schemes for enhanced scalability and reduced communication complexity (compared to quadratic PBFT), a critical step during each reconfiguration phase is the generation and negotiation of new cryptographic keys among the consensus nodes within the newly formed shards. These new keys are essential for the security of intra-shard consensus and for enabling secure cross-shard interactions.

However, this key negotiation procedure is often a **synchronous and time-consuming process**. It typically requires all participating nodes in a new shard to collectively agree on and generate shared

cryptographic material before they can resume their primary function of validating transactions. The process can last for "hundreds of seconds", during which the affected shard, or potentially the entire network if the process is centrally coordinated, essentially ceases to process transactions. This hard stop in throughput, resulting in "Zero-TPS," is a severe operational bottleneck. It means that the network, or parts of it, becomes temporarily unavailable for user activity, leading to service interruptions.

2.3. Other General Challenges of Sharding

Beyond the specific TPS-Degradation and Zero-TPS issues, sharding introduces several other general challenges that must be addressed for a robust and efficient system.

2.3.1. Cross-Shard Transaction Complexity and Overhead

One of the most significant challenges in sharded blockchains is managing "cross-shard transactions" (CTXs).³ These are transactions that involve accounts or smart contracts located on different shards. While intra-shard transactions are processed efficiently within a single shard, CTXs require coordination and communication between multiple shards to ensure atomicity—the "all-or-nothing" principle, where either all components of the transaction succeed or none do.

Traditional approaches to CTXs, such as two-phase commit (2PC) protocols, often introduce significant overhead and latency due to multiple rounds of consensus and communication between shards.³⁴ Some schemes may even record cross-shard transactions redundantly across all involved shards, increasing storage burden. The efficiency of processing CTXs is critical because, as networks scale, the proportion of cross-shard transactions can become very high, potentially exceeding 99% in networks with many shards. If not optimized, the overhead of CTXs can negate the scalability benefits of sharding, leading to higher transaction confirmation latency and reduced overall throughput.

2.3.2. Hot Shard Problem

The "hot shard problem" arises when transaction distributions among shards become imbalanced.¹⁹ This occurs when a disproportionately high volume of transactions targets accounts or smart contracts within a single shard, making it a "hot spot". This imbalance can be due to popular DApps, frequently accessed accounts, or inefficient transaction allocation mechanisms.

When a shard becomes "hot," it experiences excessive workload, leading to congestion, increased processing times, and potential bottlenecks, while other shards may be underutilized. This phenomenon can significantly reduce the actual throughput of the entire sharded blockchain compared to its theoretical performance. The shard containing the popular account experiences longer authentication times, and other shards involved in cross-shard transactions with the hot shard also face delays as they wait for confirmation. This effectively creates a centralized bottleneck within a decentralized network, undermining the core benefits of sharding. Addressing the hot shard issue requires dynamic load balancing and intelligent transaction routing to ensure an even distribution of workload across all shards.

2.3.3. Security Concerns in Sharded Environments

While sharding enhances scalability, it can introduce new security vulnerabilities. Dividing the network into smaller shards means that each shard has a smaller validator set compared to the entire network. This makes individual shards potentially more susceptible to "shard takeover attacks" (also known as 1%

attacks or 51% attacks on a single shard), where an attacker might only need to control a smaller percentage of the

total network's stake to compromise a single shard. If a malicious actor gains control over a majority of validators within a shard (e.g., $>2/3$), they could potentially validate fraudulent transactions or prevent legitimate ones from being processed within that shard.

Ensuring that validators are randomly and unpredictably assigned to shards and frequently rotated is crucial to mitigate this risk. Without robust shuffling mechanisms, an attacker could predict validator assignments and concentrate resources to compromise a target shard. The security of cross-shard communication is also a concern, as any flaws in data reconciliation or message transfer protocols could lead to inconsistencies or vulnerabilities across the entire network.

2.3.4. Data Availability Challenges

In sharded blockchains, ensuring "data availability" is critical. Data availability refers to the ability of all network participants to access and verify the transaction data needed to validate blocks and maintain consensus. In a sharded environment, where each node may only store a portion of the total state, ensuring that all necessary data is accessible for verification, especially for cross-shard transactions or historical queries, becomes complex.

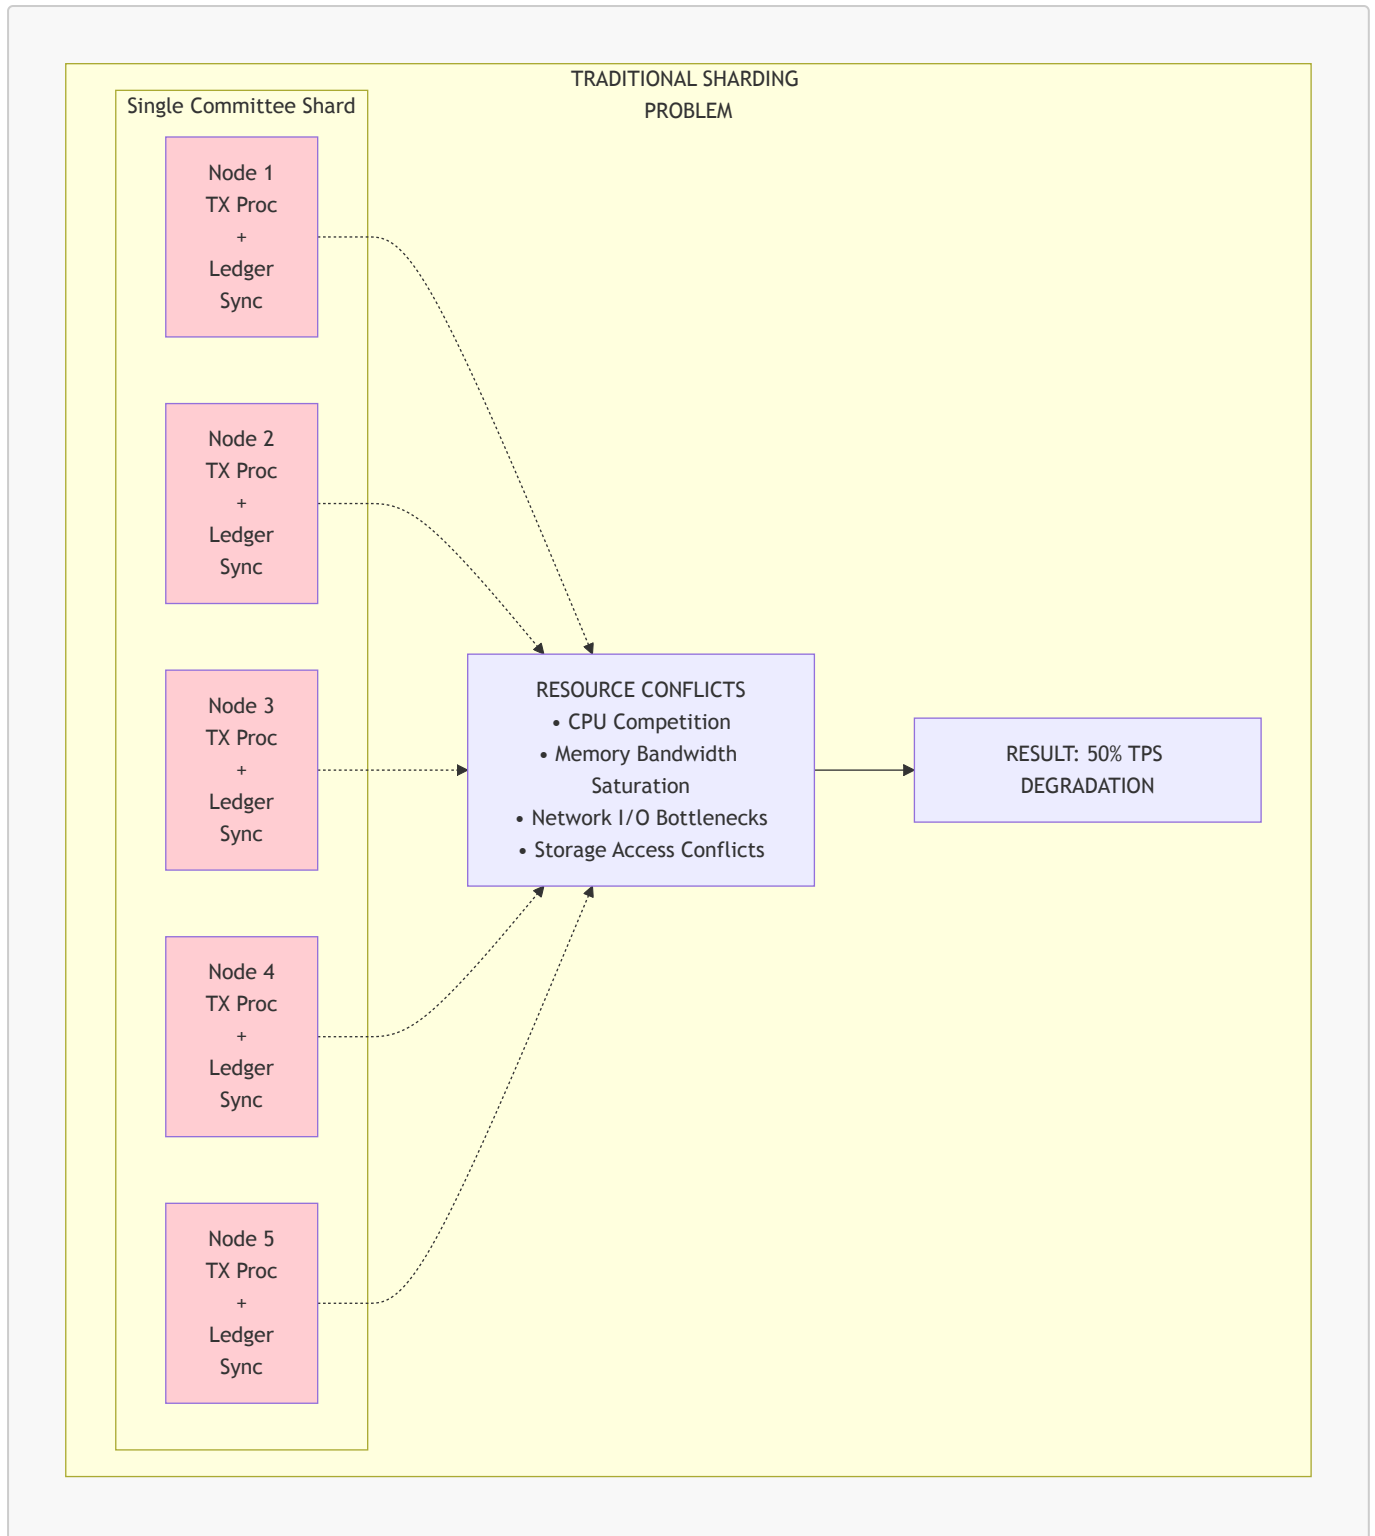
Challenges include "data withholding attacks," where malicious actors intentionally fail to share critical data, leading to inconsistencies. As transaction data accumulates, "storage bloat" can make it difficult for individual nodes to store and synchronize the entire ledger, potentially leading to centralization if only powerful nodes can keep up. "Verification overhead" also increases as data grows, hindering network speed. Solutions often involve "Data Availability Layers (DALs)" or "Data Availability Sampling (DAS)", but these add architectural complexity.

3. LoraShard Solutions

LoraShard introduces innovative architectural mechanisms specifically designed to address the critical TPS-Degradation and Zero-TPS issues that plague dynamic sharding systems. These solutions are rooted in parallelization and strategic resource management during shard reconfiguration.

3.1. Parallelized Dual Committee Framework for TPS-Degradation

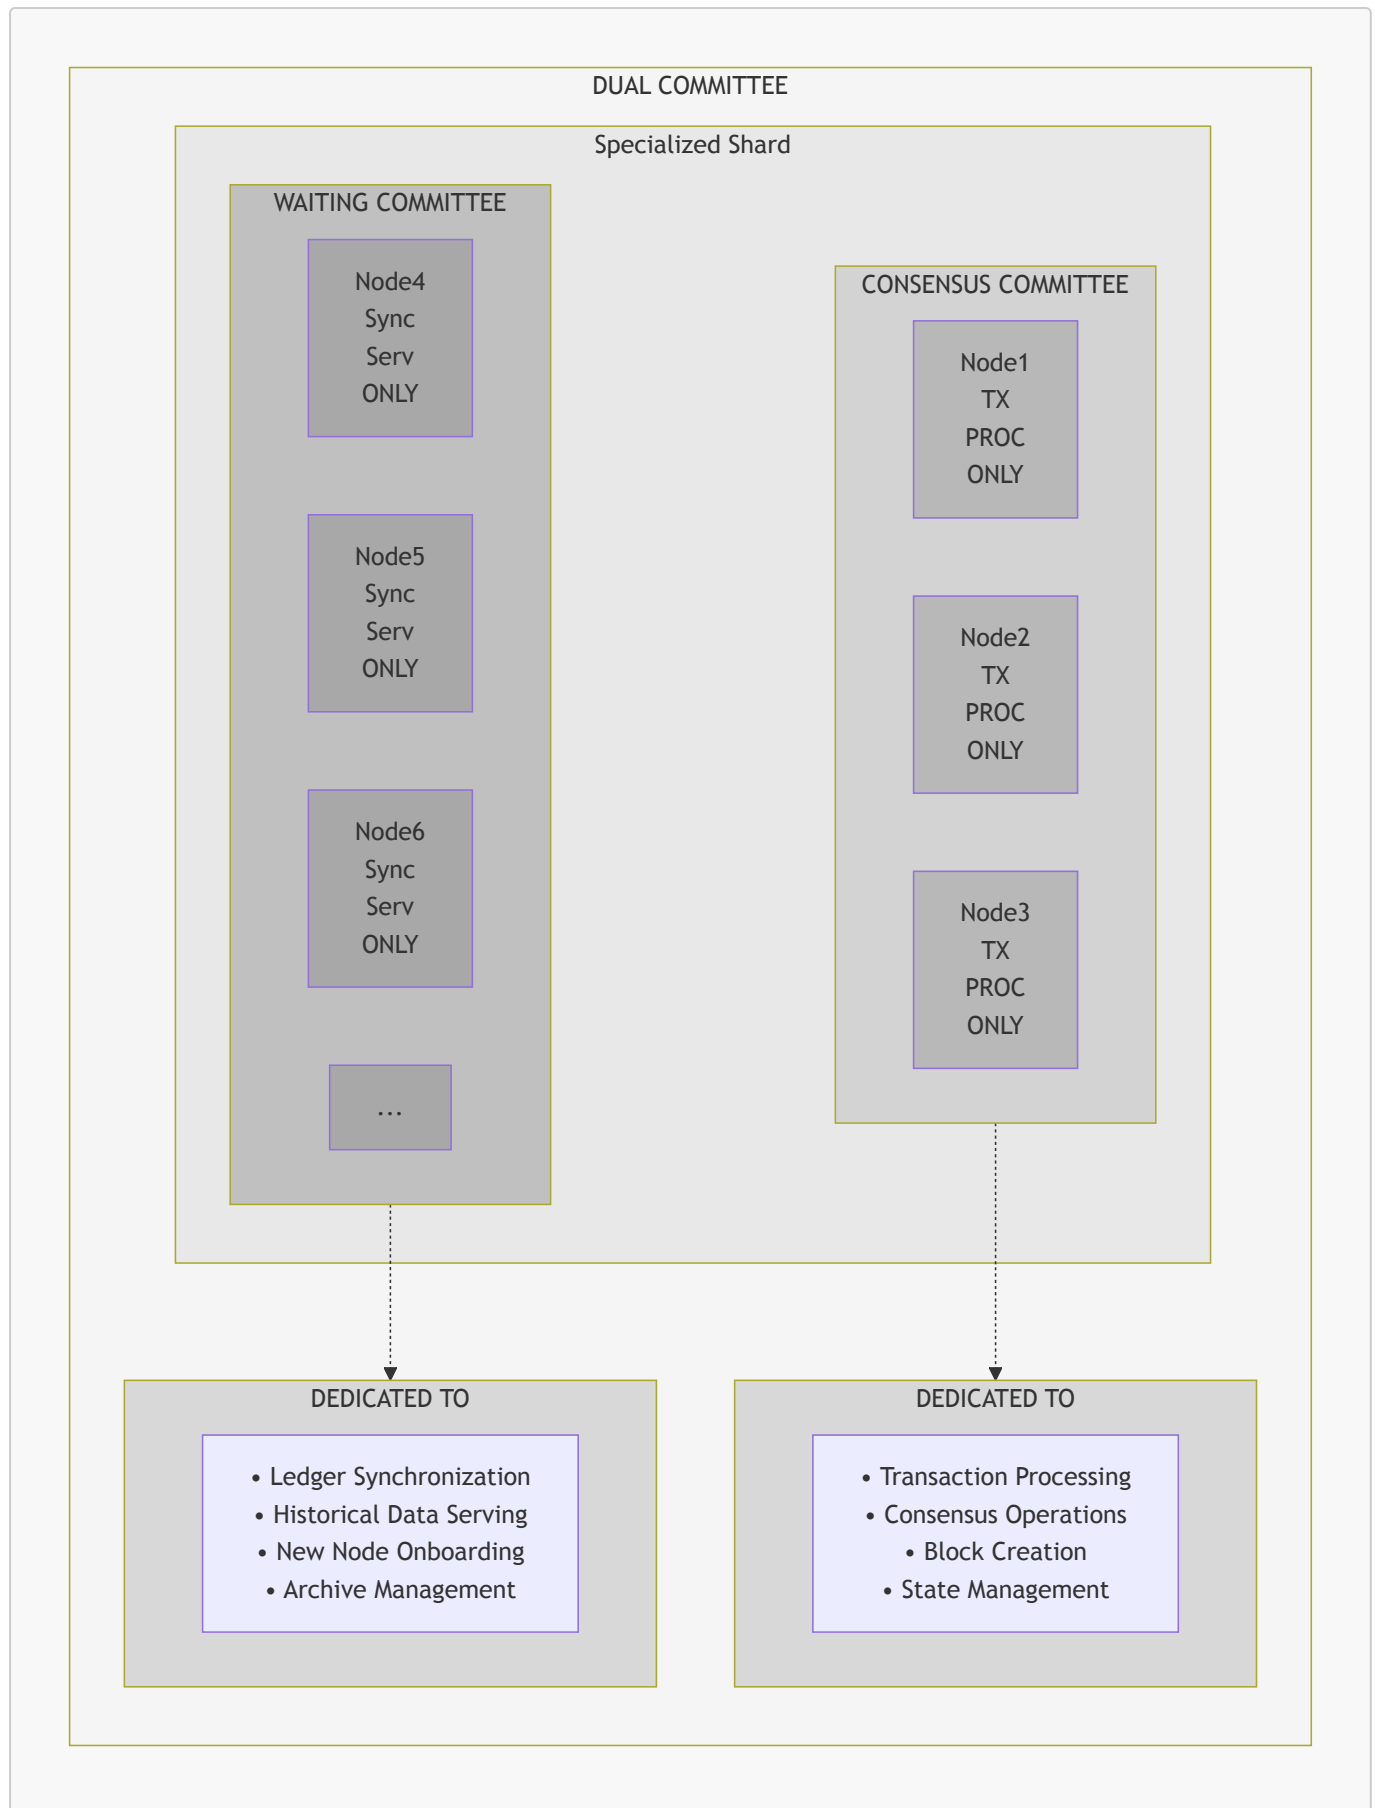
The **Parallelized Dual Committee Framework** represents LoraShard's revolutionary approach to solving the TPS-Degradation issue through architectural separation of concerns. This innovative design fundamentally reimagines how blockchain shards organize their nodes, moving away from the traditional monolithic approach where all nodes perform identical functions toward a **specialized, role-based architecture** that optimizes performance through intelligent workload distribution.



3.1.1. Architecture and Operation

3.1.1.1. Dual Committee Structure

The dual committee architecture divides each shard into two **functionally specialized groups** that collaborate to provide comprehensive blockchain services while maintaining operational independence for performance optimization.



Consensus Committee Structure:

- Size: Typically nc nodes (100–300 nodes depending on shard requirements)
- Selection Criteria: Highest reputation nodes within the shard for optimal performance
- Primary Function: Exclusive focus on transaction processing and consensus operations

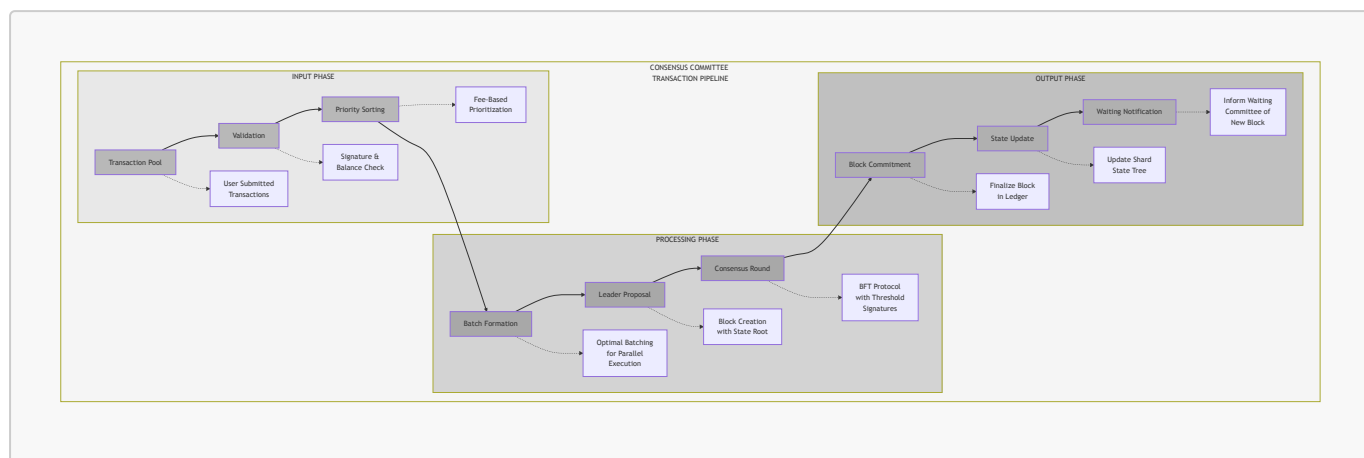
- Performance Optimization: Dedicated resources without synchronization overhead
- Leadership Model: Rotating leadership among committee members for fairness

Waiting Committee Structure:

- Size: All remaining shard nodes not in consensus committee (typically 300-500 nodes)
- Membership: Mix of established nodes and newly arrived nodes requiring synchronization
- Primary Function: Historical data maintenance and new node synchronization services
- Resource Allocation: Optimized for high-bandwidth data transfer and storage operations
- Scalability: Can accommodate variable numbers of nodes without affecting consensus performance

3.1.1.2. Consensus Committee Operations

The consensus committee operates with **dedicated resources and streamlined processes** optimized exclusively for high-throughput transaction processing without synchronization distractions.

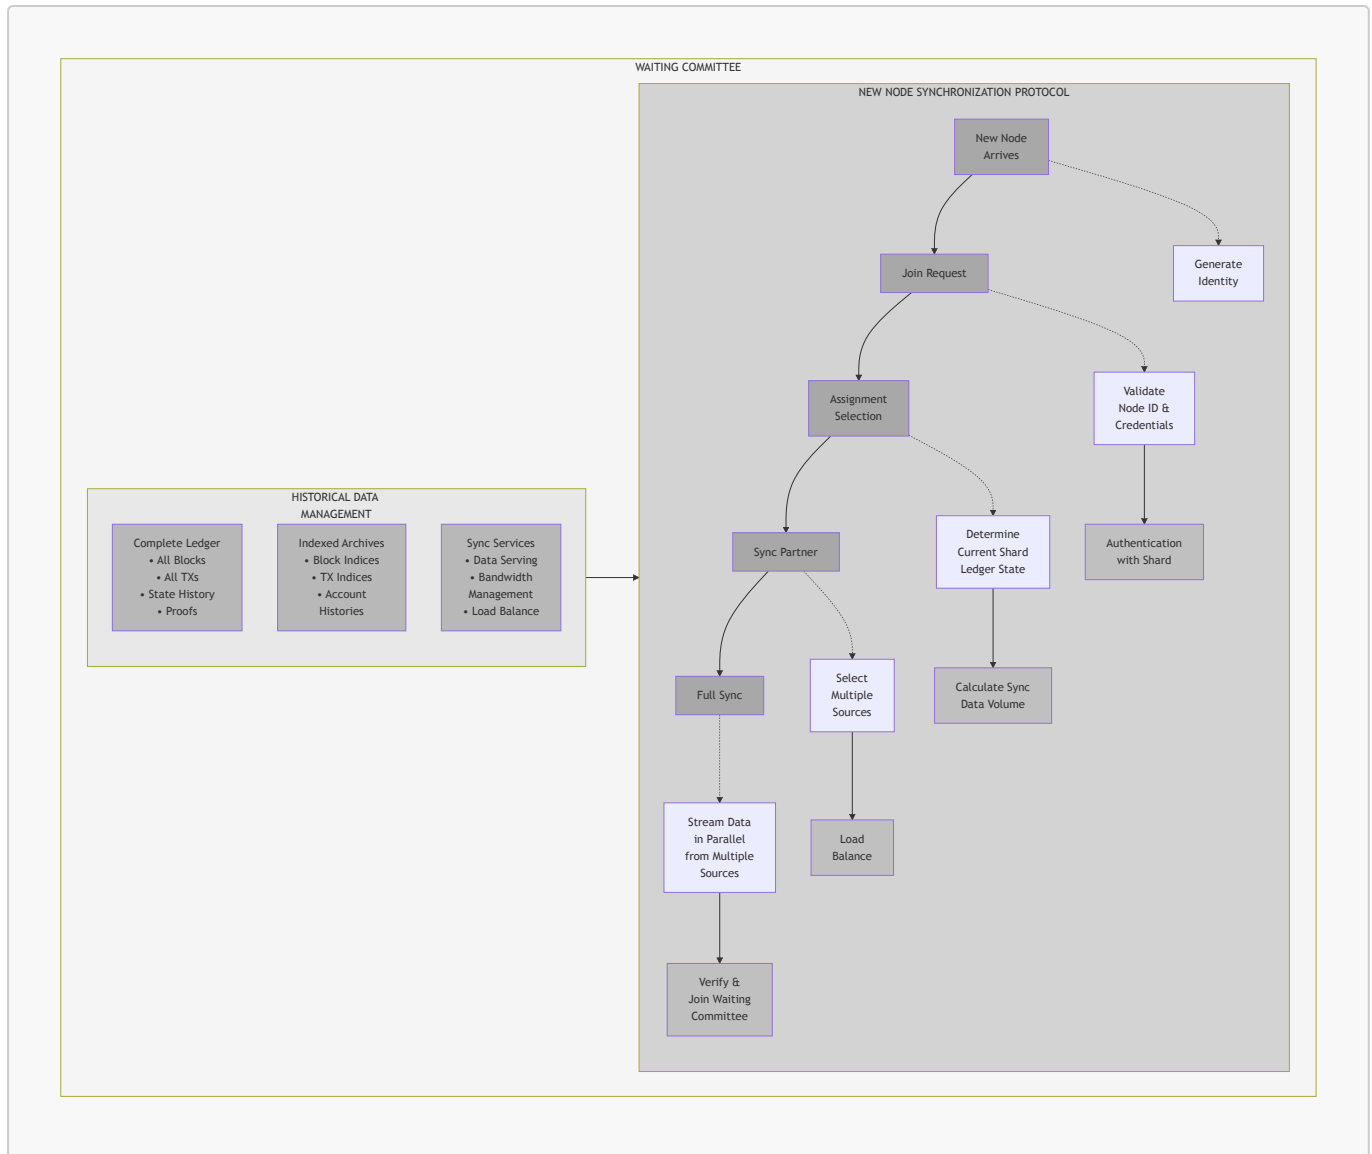


Performance Optimization Features:

- Dedicated Resource Access: No resource sharing with synchronization operations
- Optimized Batch Processing: Intelligent transaction batching for maximum throughput
- Parallel Validation: Multi-threaded transaction validation for faster processing
- Streamlined Consensus: BFT protocol optimized for consistent committee membership

3.1.1.3. Waiting Committee Operations

The waiting committee serves as the **comprehensive historical archive** for the shard, maintaining complete blockchain history and serving synchronization requests without impacting consensus performance.

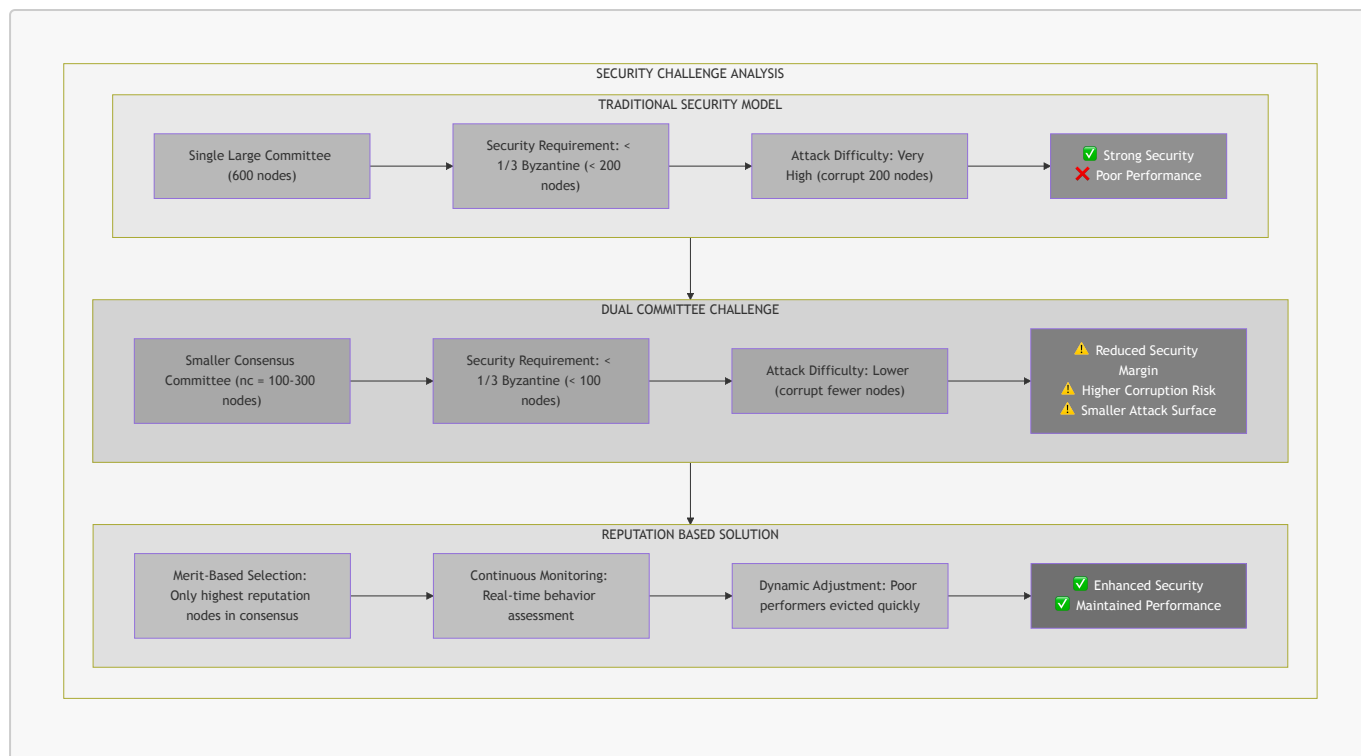


Synchronization Optimization Strategies:

- Segmented Transfer: Historical data divided into segments for parallel downloading
- Multiple Sources: Each new node connects to multiple waiting nodes for redundancy
- Bandwidth Optimization: Transfer rates adjusted based on network conditions and node capacity
- Progress Tracking: Real-time monitoring of synchronization progress and error recovery

3.1.2. Reputation Mechanism for Enhanced Security

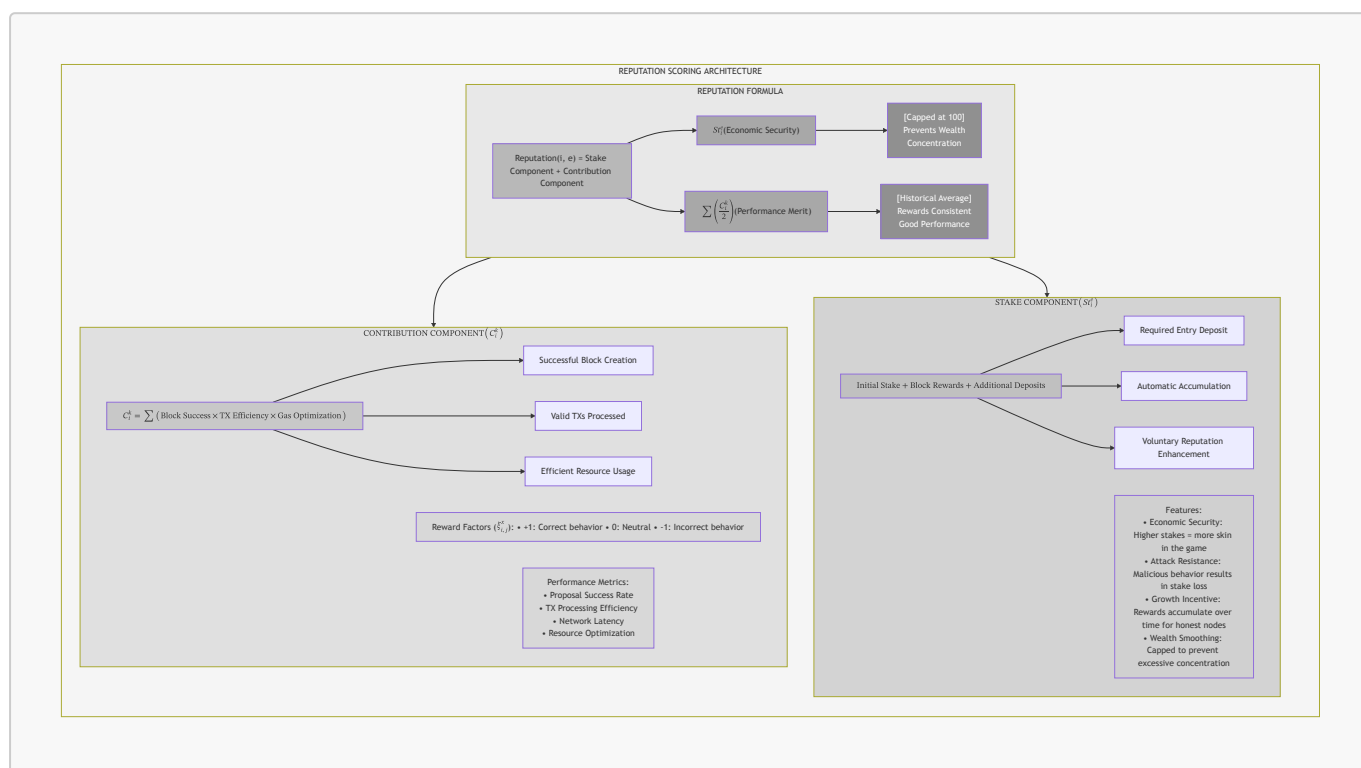
While the dual committee architecture solves the TPS-Degradation issue, it introduces new **security considerations** that must be addressed through sophisticated reputation-based mechanisms to maintain network security guarantees.



- **Reduced Committee Size:** Smaller consensus committees are more vulnerable to corruption
- **Concentrated Attack Surface:** Adversaries need to corrupt fewer nodes to impact consensus
- **Performance vs Security Trade-off:** Balance between efficiency and security guarantees
- **Dynamic Threat Landscape:** Evolving attack strategies requiring adaptive defenses

5.1.2.1. Reputation Scoring System

LoraShard implements a **sophisticated reputation scoring system** that evaluates nodes across multiple dimensions to ensure only the most reliable and capable nodes participate in critical consensus operations.

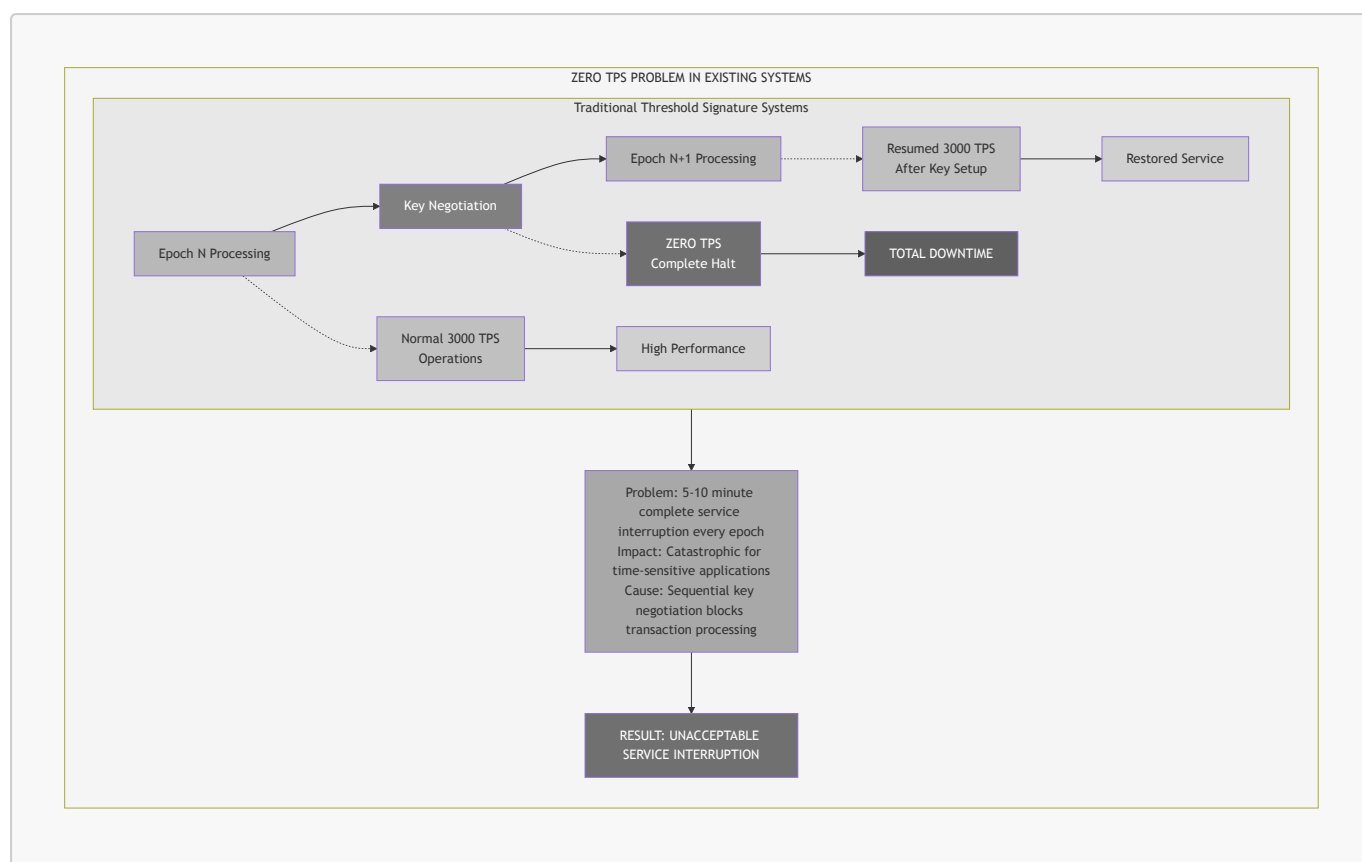


Reputation Calculation Algorithm:

- Stake Component: Economic security through financial commitment
- Contribution Component: Performance-based merit through block creation success
- Historical Weighting: Recent performance weighted more heavily than distant history
- Penalty System: Malicious behavior results in immediate reputation penalties

3.2. Parallelized Key Pre-negotiation Mechanism for Zero-TPS

The **Parallelized Key Pre-negotiation Mechanism** represents LoraShard's innovative solution to the Zero-TPS issue that plagues threshold signature-based blockchain systems. This mechanism addresses a fundamental bottleneck where consensus nodes must halt all transaction processing to negotiate cryptographic keys whenever committee membership changes, resulting in complete network unavailability during reconfiguration periods.



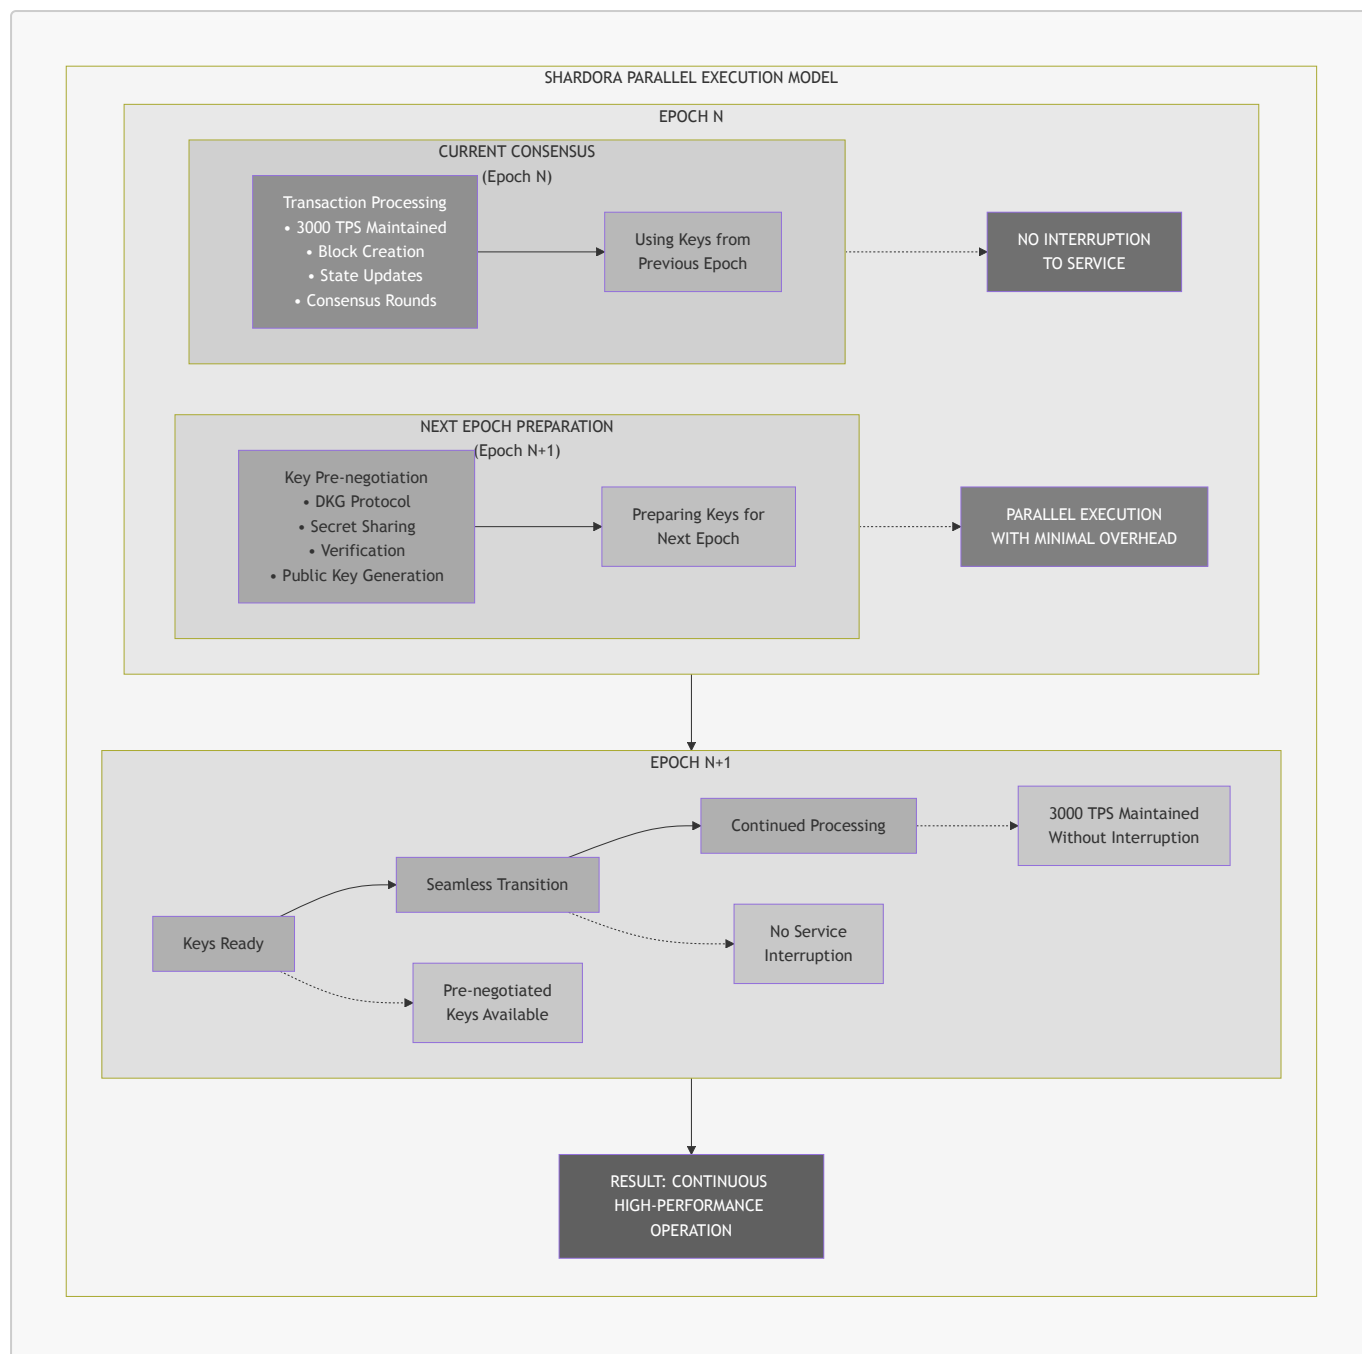
The Zero-TPS issue emerges because **threshold signature schemes** require all consensus participants to engage in complex Distributed Key Generation (DKG) protocols before they can participate in consensus. Unlike traditional PBFT systems where nodes can immediately join with their existing cryptographic identities, threshold signatures demand that all signers share compatible cryptographic keys generated through coordinated multi-party computation.

LoraShard's pre-negotiation mechanism **fundamentally restructures the temporal relationship** between key generation and consensus operations, enabling these critical processes to occur in parallel rather than sequentially, thereby maintaining continuous transaction processing capabilities throughout network reconfiguration.

3.2.1. Architecture and Operation

3.2.1.1. Temporal Parallelization Strategy

The core innovation of LoraShard's approach lies in **temporal decoupling** of key negotiation from consensus operations through parallel execution of current-epoch consensus and next-epoch key preparation.

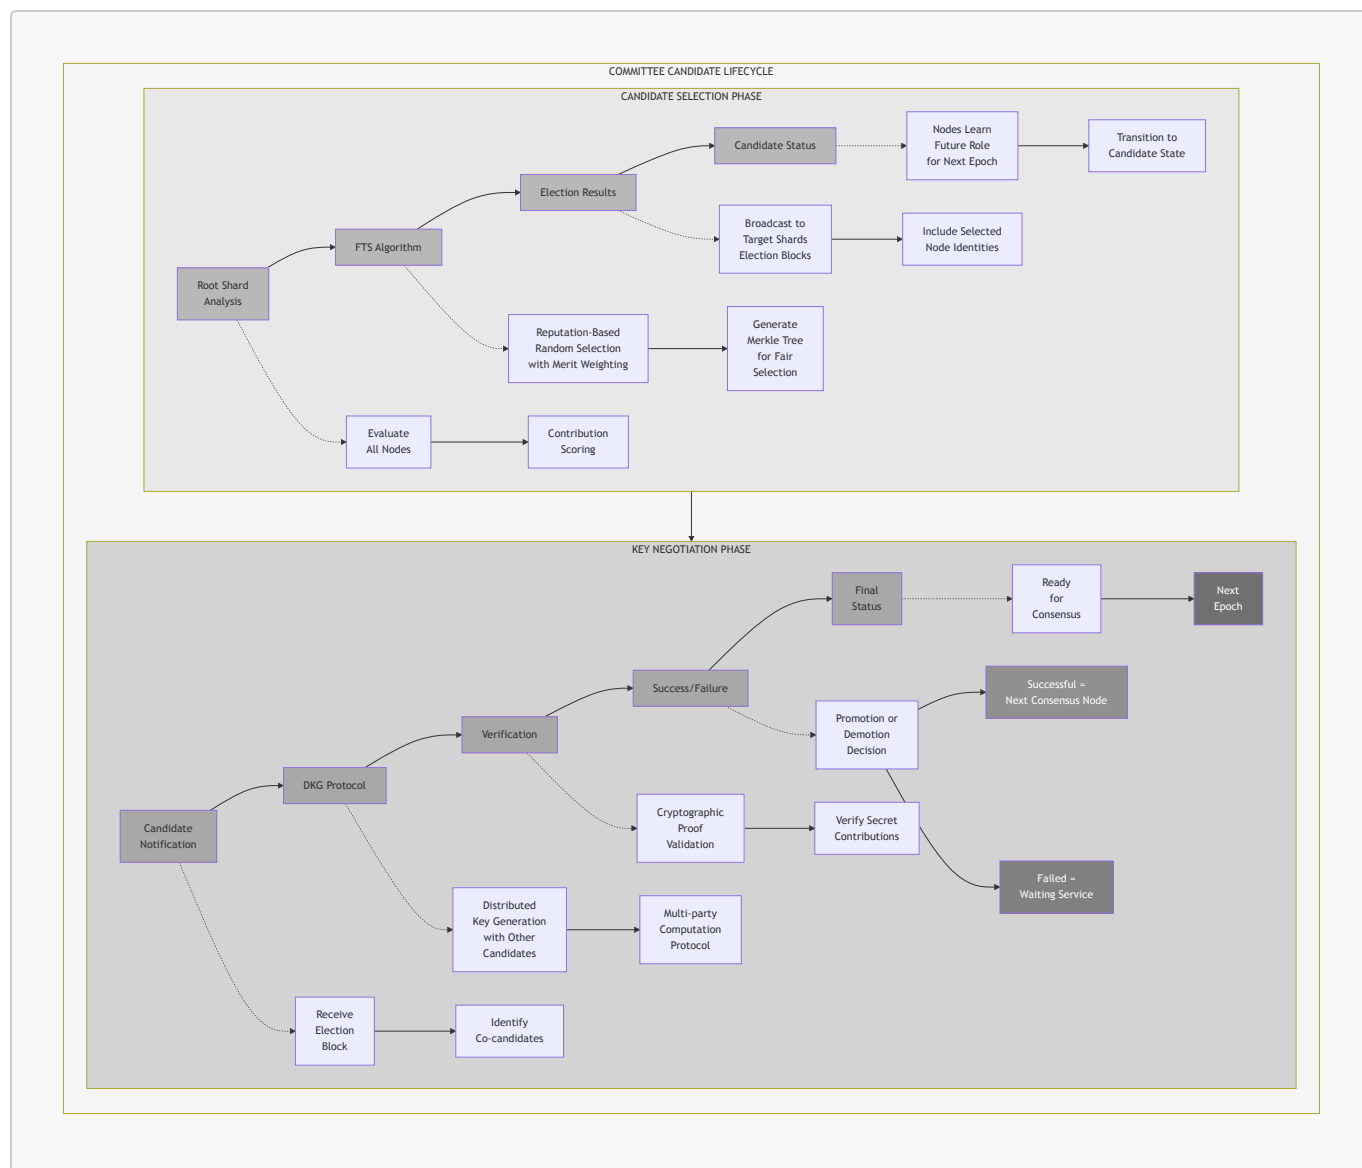


Execution Timeline:

- Current Epoch Operations: Consensus committee processes transactions using keys negotiated in the previous epoch
- Parallel Key Preparation: Candidate nodes for next epoch simultaneously negotiate cryptographic keys
- Seamless Transition: When epoch boundary arrives, pre-negotiated keys enable immediate consensus participation
- Continuous Service: No interruption to transaction processing throughout the reconfiguration cycle

3.2.1.2. Committee Candidate System

LoraShard implements a **sophisticated candidate system** that identifies and prepares future consensus participants well in advance of their actual service period, enabling comprehensive key preparation without impacting current operations.

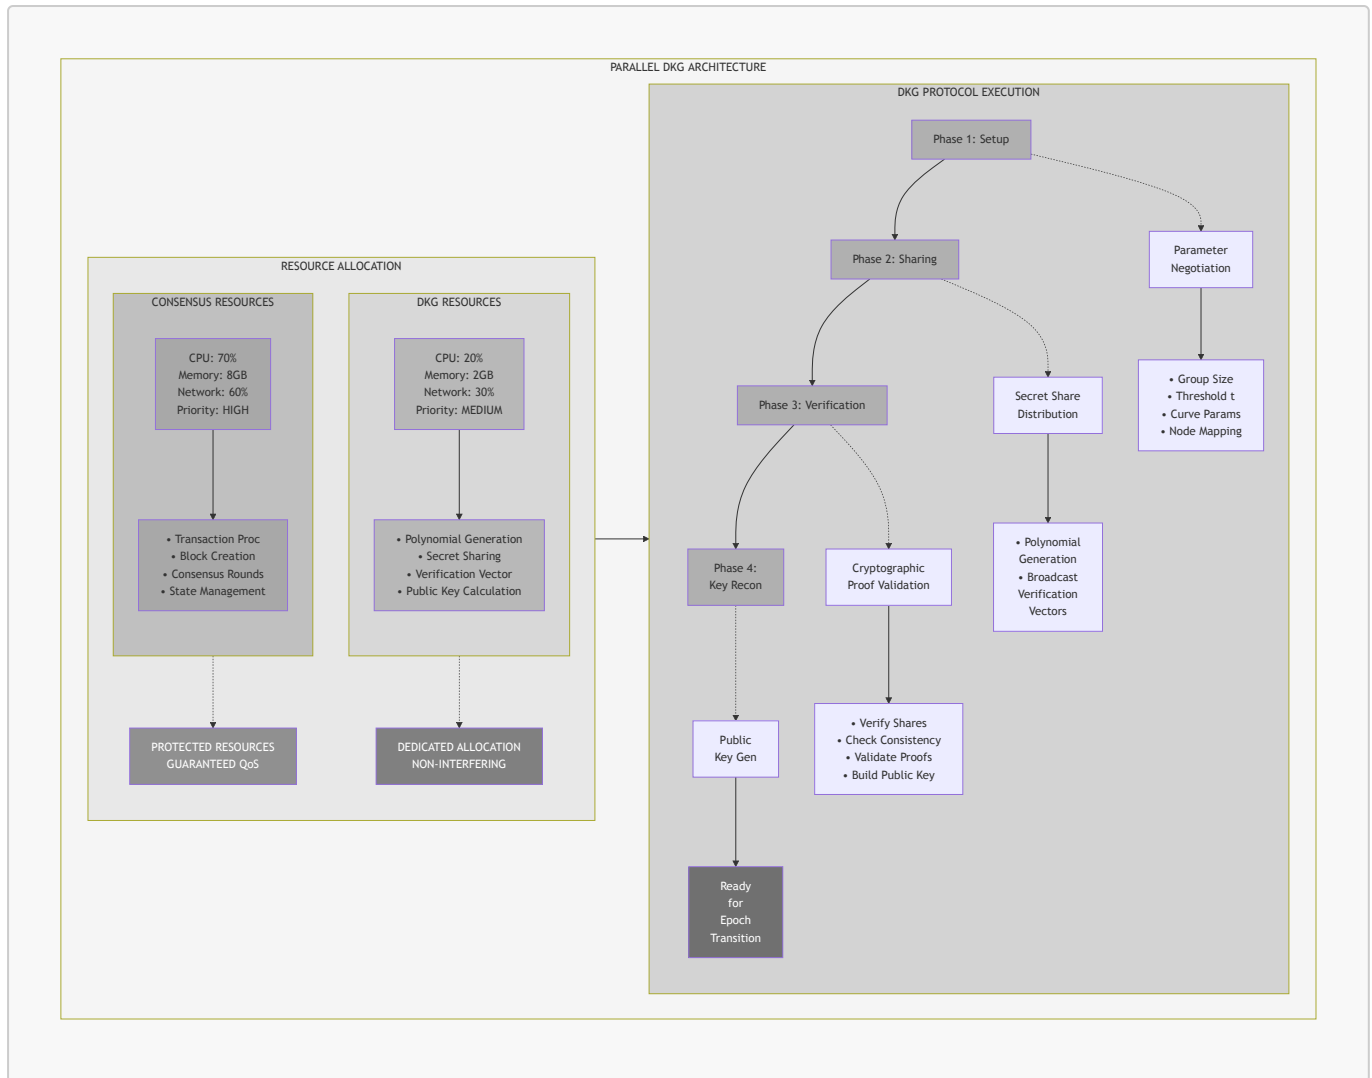


Candidate Node Responsibilities:

- **Dual Participation:** Continue current role while preparing for potential promotion
- **Resource Management:** Allocate computational resources between current duties and key negotiation
- **Cryptographic Preparation:** Generate and verify secret shares for threshold signatures
- **Readiness Assessment:** Validate successful key negotiation before epoch transition

3.2.1.3. Distributed Key Generation Integration

The pre-negotiation mechanism integrates **sophisticated DKG protocols** that operate in parallel with ongoing consensus operations through careful resource management and process isolation.



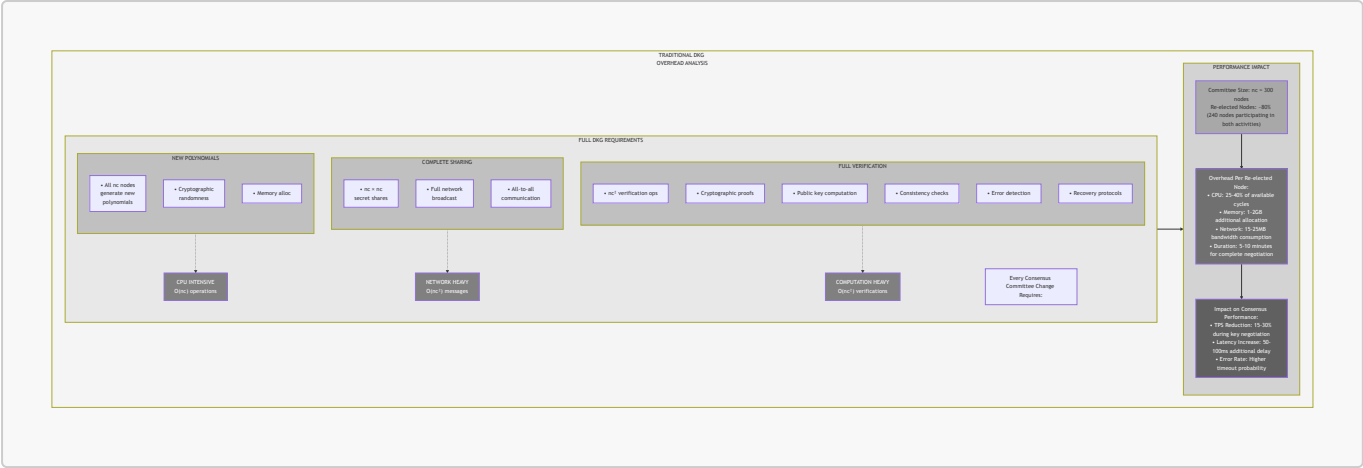
DKG Integration Features:

- **Resource Isolation:** DKG operations allocated dedicated resources to prevent consensus interference
- **Asynchronous Execution:** Key negotiation proceeds independently of consensus timing
- **Failure Tolerance:** Failed key negotiations result in candidate demotion without impacting current operations
- **Cryptographic Soundness:** Full DKG security properties maintained despite parallel execution

3.2.2. Secret-Reuse Strategy

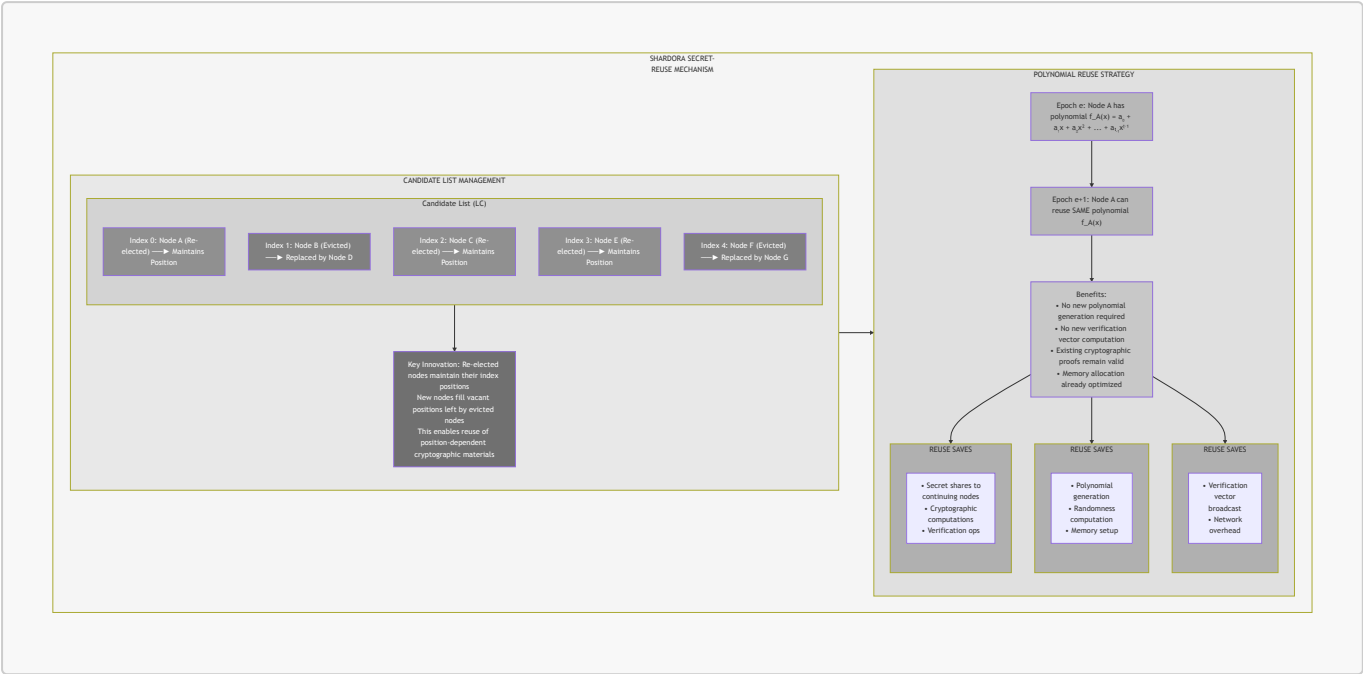
3.2.2.1. Lightweight Key Negotiation Problem

Even with parallel execution, traditional DKG protocols impose **significant computational and communication overhead** that can impact consensus performance when re-elected nodes must participate in both current consensus and future key negotiation simultaneously.



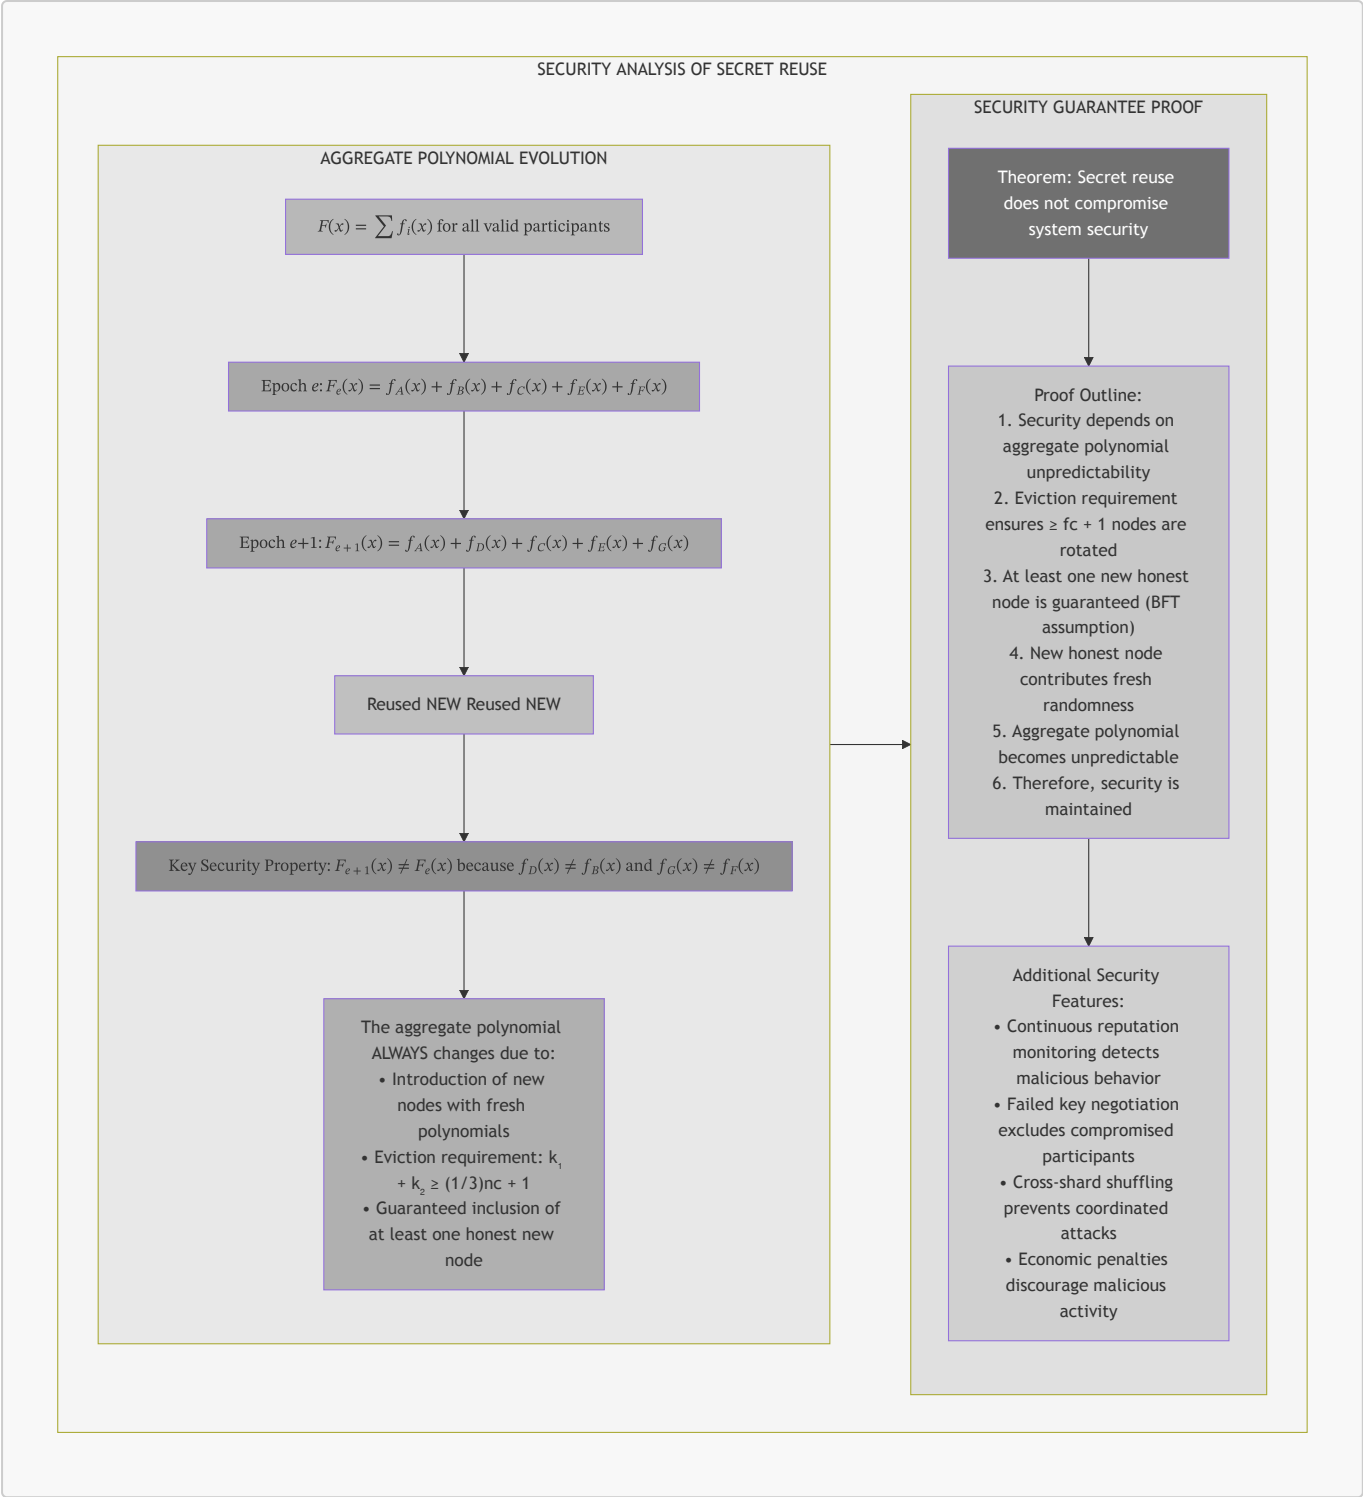
3.2.2.2. Secret Reuse Innovation

LoraShard introduces a **revolutionary secret-reuse strategy** that enables re-elected consensus nodes to leverage previously computed cryptographic materials, dramatically reducing the computational and communication overhead of repeated key negotiations.



3.2.2.3. Cryptographic Security Analysis

The secret-reuse strategy maintains **full cryptographic security** despite reusing polynomials because the overall system polynomial changes through the introduction of new nodes with fresh cryptographic material.



4. System Model

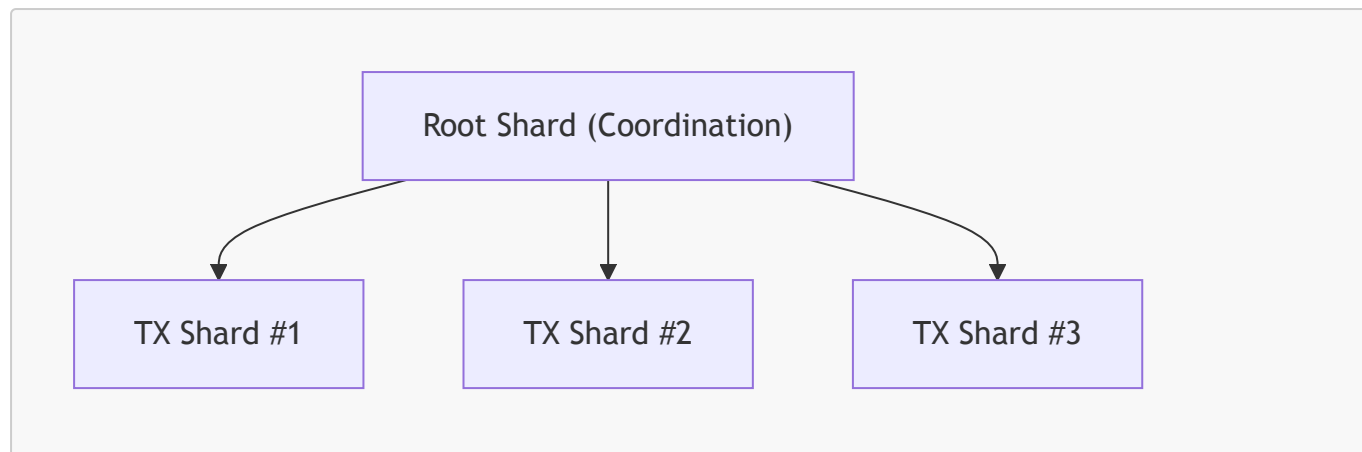
4.1. Network Model

LoraShard operates within a **peer-to-peer distributed network** comprising heterogeneous nodes with varying computational power, bandwidth capabilities, and storage resources. The network model is designed to accommodate the dynamic nature of blockchain environments while maintaining high performance and security standards through sophisticated coordination mechanisms.

The network architecture follows a **multi-shard topology** where nodes are strategically distributed across different shards to enable parallel transaction processing. Unlike traditional single-chain

architectures, LoraShard network model supports concurrent operations across multiple blockchain segments, each maintaining its own state and processing capabilities.

4.1.1. Network Architecture



Root Shard: Acts as the central coordination hub responsible for:

- **Epoch Management:** Broadcasting time blocks to synchronize all transaction shards
- **Reputation Tracking:** Maintaining and updating reputation scores for all network participants
- **Node Assignment:** Executing strategic node shuffling algorithms across shards
- **Global State Coordination:** Ensuring consistency and security across the entire network

Transaction Shards: Independent processing units that handle:

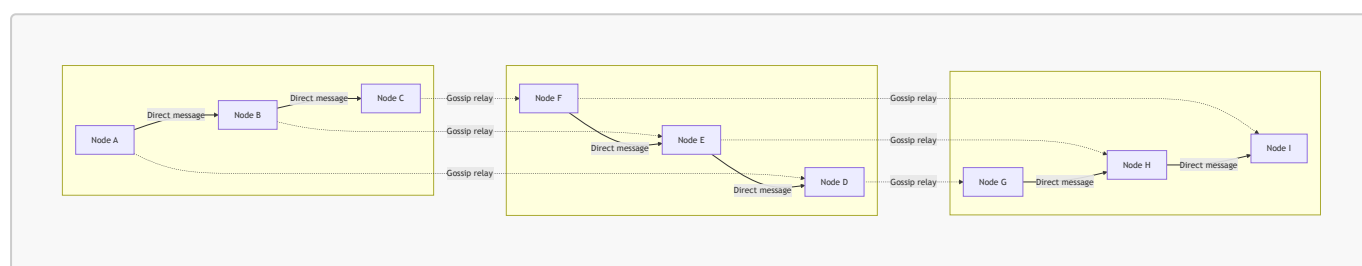
- **Parallel Transaction Processing:** Each shard processes its own subset of transactions
- **State Management:** Maintaining local ledger state for assigned accounts/contracts
- **Intra-shard Consensus:** Running BFT consensus protocols with threshold signatures
- **Cross-shard Communication:** Coordinating with other shards for cross-shard transactions

4.1.2. Network Topology and Communication

LoraShard assumes a **well-connected network** where honest nodes maintain robust communication channels with each other. The network operates under the following connectivity principles:

- **Full Mesh Connectivity:** Within each shard, all nodes maintain direct communication channels
- **Inter-shard Bridges:** Dedicated communication pathways exist between different shards
- **Gossip Protocol Integration:** Network-wide information dissemination through efficient gossip mechanisms
- **Adaptive Routing:** Dynamic path selection based on network conditions and node availability

Message Propagation Mechanism



The network employs a hybrid message propagation strategy:

- **Direct Communication:** Critical consensus messages are sent directly between consensus nodes
- **Gossip Propagation:** General network information spreads through gossip protocols for efficiency
- **Selective Broadcasting:** Different message types use appropriate propagation strategies based on urgency and importance

4.1.3. Synchronization and Timing Model

LoraShard operates under a partially synchronous network model with the following characteristics:

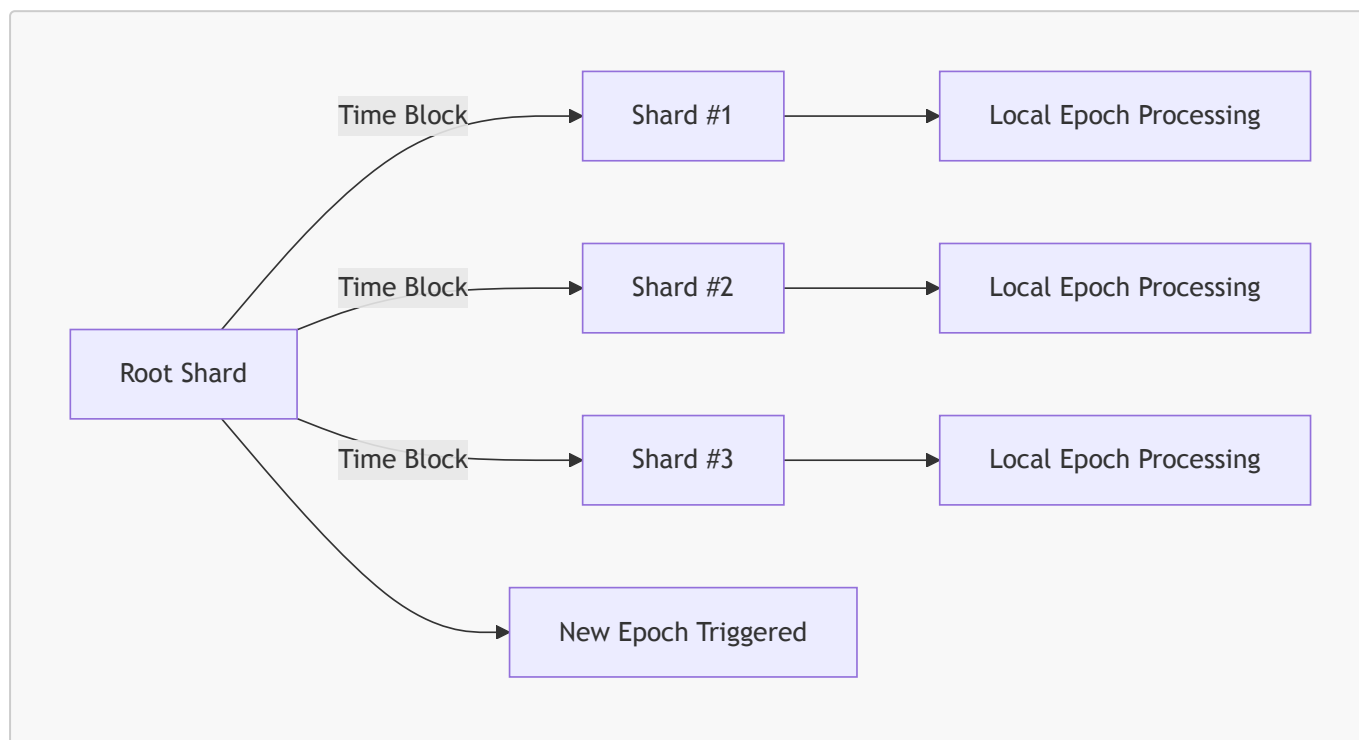
Known Bounded Delay (Δ): For critical operations such as identity creation and shard assignment, the network assumes messages sent by honest nodes will reach all other honest nodes within a fixed maximum delay Δ . This bounded delay ensures:

- **Deterministic Timeouts:** Nodes can set appropriate timeout values for critical operations
- **Liveness Guarantees:** The system maintains progress even under network stress
- **Security Assurance:** Prevents timing-based attacks on consensus mechanisms

Optimistic Timeouts: For regular consensus operations, LoraShard employs:

- **Exponentially Increasing Timeouts:** Starting with optimistic low values and increasing upon timeout
- **Adaptive Adjustment:** Timeout values adjust based on observed network performance
- **Fallback Mechanisms:** Graceful degradation when network conditions deteriorate

Clock Synchronization Mechanism



Time Block Broadcasting: The root shard periodically generates and broadcasts time blocks to all transaction shards, serving as:

- **Epoch Synchronization:** Coordinating the start of new epochs across all shards
- **Global Clock Reference:** Providing a network-wide time reference for operations
- **Reconfiguration Trigger:** Initiating shard reconfiguration processes at predetermined intervals

4.2. Node Types and Roles

LoraShard employs a sophisticated multi-tier node architecture that optimizes network performance through strategic role specialization. Unlike traditional blockchain systems where all nodes perform identical functions, LoraShard design recognizes that different network operations require different capabilities and responsibilities. This hierarchical approach enables efficient resource utilization while maintaining strong security guarantees through reputation-based node management.

The node classification system is dynamic and merit-based, allowing nodes to transition between roles based on their performance, reputation scores, and network contributions. This flexibility ensures that the most capable and reliable nodes handle critical consensus operations, while providing pathways for new or lower-performing nodes to contribute meaningfully to network security and operation.

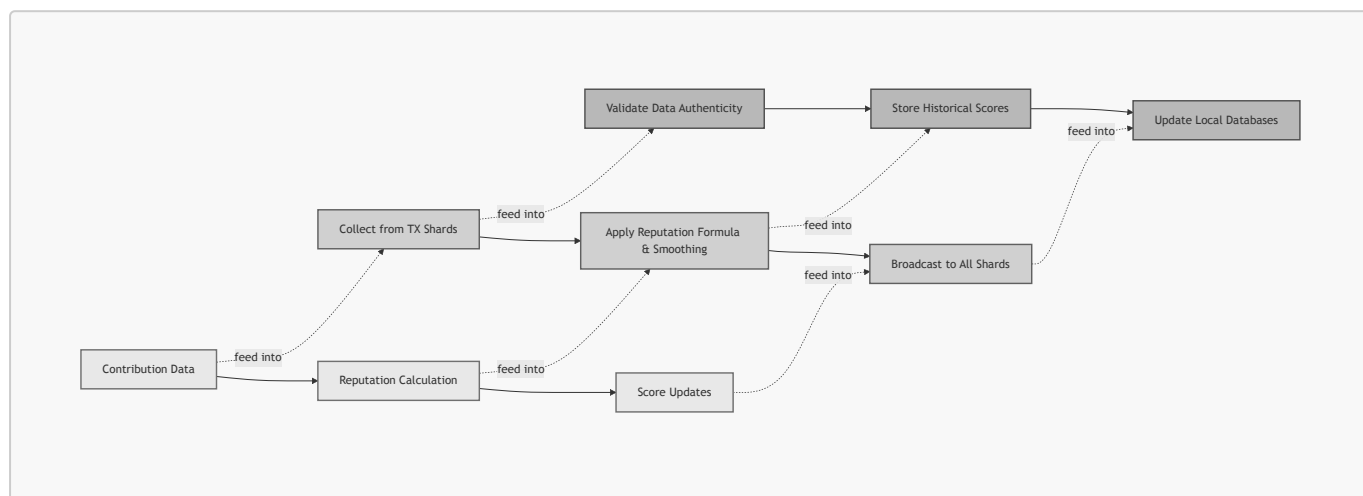
4.2.1. Root Shard Nodes

Root shard nodes serve as the central nervous system of the LoraShard network, orchestrating global coordination and maintaining network-wide consistency. These nodes operate within a specialized shard dedicated exclusively to network governance and coordination functions.

Epoch Management and Coordination:

- *Time Block Generation*: Create and broadcast time blocks that synchronize epoch transitions across all transaction shards
- *Global Clock Maintenance*: Maintain network-wide temporal consistency through **Verifiable Secret Sharing (VSS)** mechanisms
- *Epoch Randomness Generation*: Generate cryptographically secure random values using **Feldman VSS protocol** for unbiased node selection

Reputation System Management:



- *Reputation Score Calculation*: Process contribution data from transaction shards and compute updated reputation scores
- *Historical Tracking*: Maintain comprehensive records of node performance and behavior patterns
- *Score Distribution*: Broadcast updated reputation scores to all transaction shards for local decision-making

Strategic Node Orchestration:

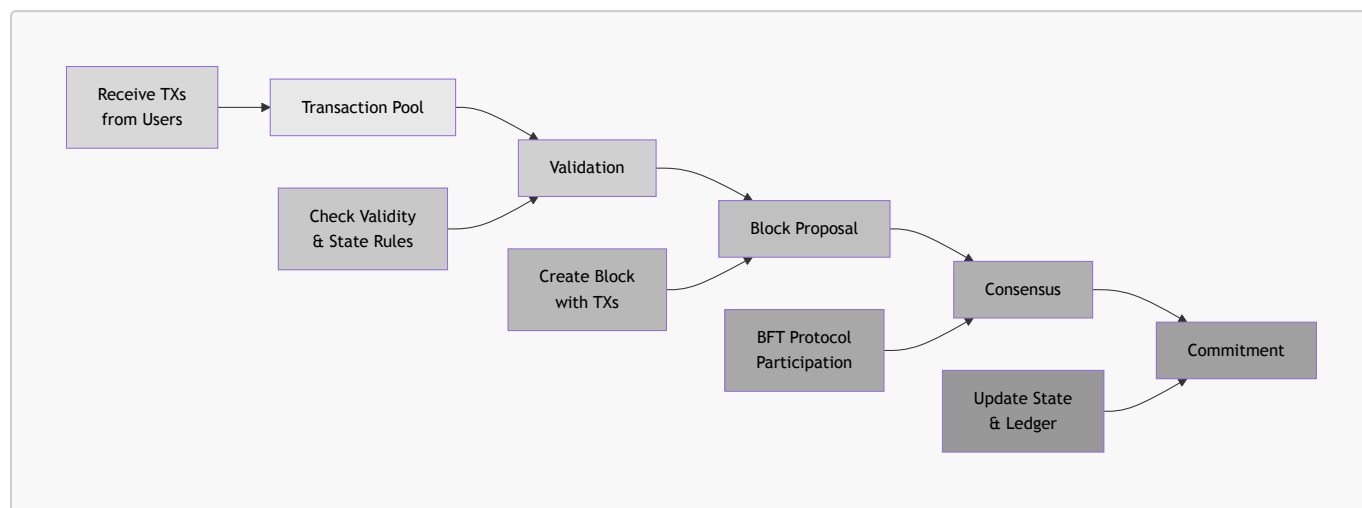
- Follow-the-Satoshi (FTS) Algorithm Execution: Run sophisticated node selection algorithms for fair and secure committee formation
- Cross-Shard Balancing: Ensure balanced reputation distribution across all transaction shards
- Election Block Generation: Create and disseminate election results to trigger node shuffling processes

4.2.2. Transaction Shard Nodes

Transaction shard nodes form the backbone of LoraShard parallel processing capabilities, handling the majority of user transactions and maintaining distributed ledger state. These nodes are organized into dual committee structures that enable continuous operation during reconfiguration periods. 4**2.2.1. Consensus Nodes**

Consensus nodes represent the active processing tier within each transaction shard, directly responsible for transaction validation, block creation, and consensus participation. These nodes are selected based on superior reputation scores and proven capability to handle high-throughput operations.

Transaction Processing Pipeline:



- *Transaction Validation*: Verify transaction authenticity, signature validity, and state consistency
- *Block Proposal*: Leaders create new blocks containing batches of valid transactions
- *Consensus Participation*: Engage in BFT consensus protocols with BLS threshold signatures
- *State Management*: Maintain current shard state and process state transitions

Leadership Rotation Mechanism:

- Dynamic Leader Selection: Leaders rotate after each successful block to ensure fair participation
- Reputation-Based Priority: Higher reputation nodes receive leadership opportunities more frequently
- Performance Monitoring: Leader performance is continuously evaluated and factored into reputation updates

Cross-Shard Coordination:

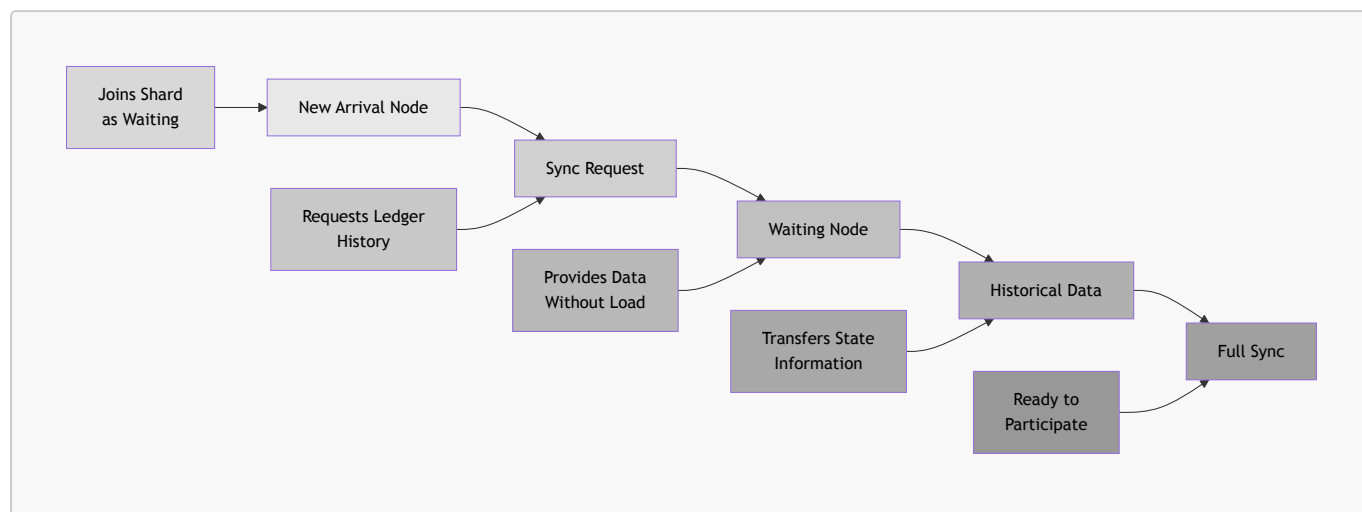
- Cross-Shard Transaction Processing: Handle transactions that span multiple shards through atomic protocols
- Proof Generation: Create cryptographic proofs for cross-shard transaction validation

- **Inter-Shard Communication:** Coordinate with consensus nodes in other shards for complex operations

4.2.2.2. Waiting Nodes

Waiting nodes serve as the strategic reserve within each shard, maintaining network readiness while providing essential support services. These nodes represent a pool of qualified participants ready to assume consensus responsibilities when needed.

Ledger Synchronization Services:



- *Historical Data Provision:* Serve complete ledger history to newly arrived nodes without impacting consensus operations
- *Parallel Synchronization:* Enable new nodes to sync historical state while consensus nodes continue processing transactions
- *Bandwidth Optimization:* Distribute synchronization load across multiple waiting nodes to optimize network resource usage

Continuous State Monitoring:

- **Real-Time State Tracking:** Maintain up-to-date shard state by monitoring all committed transactions
- **Verification Services:** Independently verify consensus node decisions to detect potential inconsistencies
- **Backup State Maintenance:** Serve as backup state repositories in case of consensus node failures

Election Readiness:

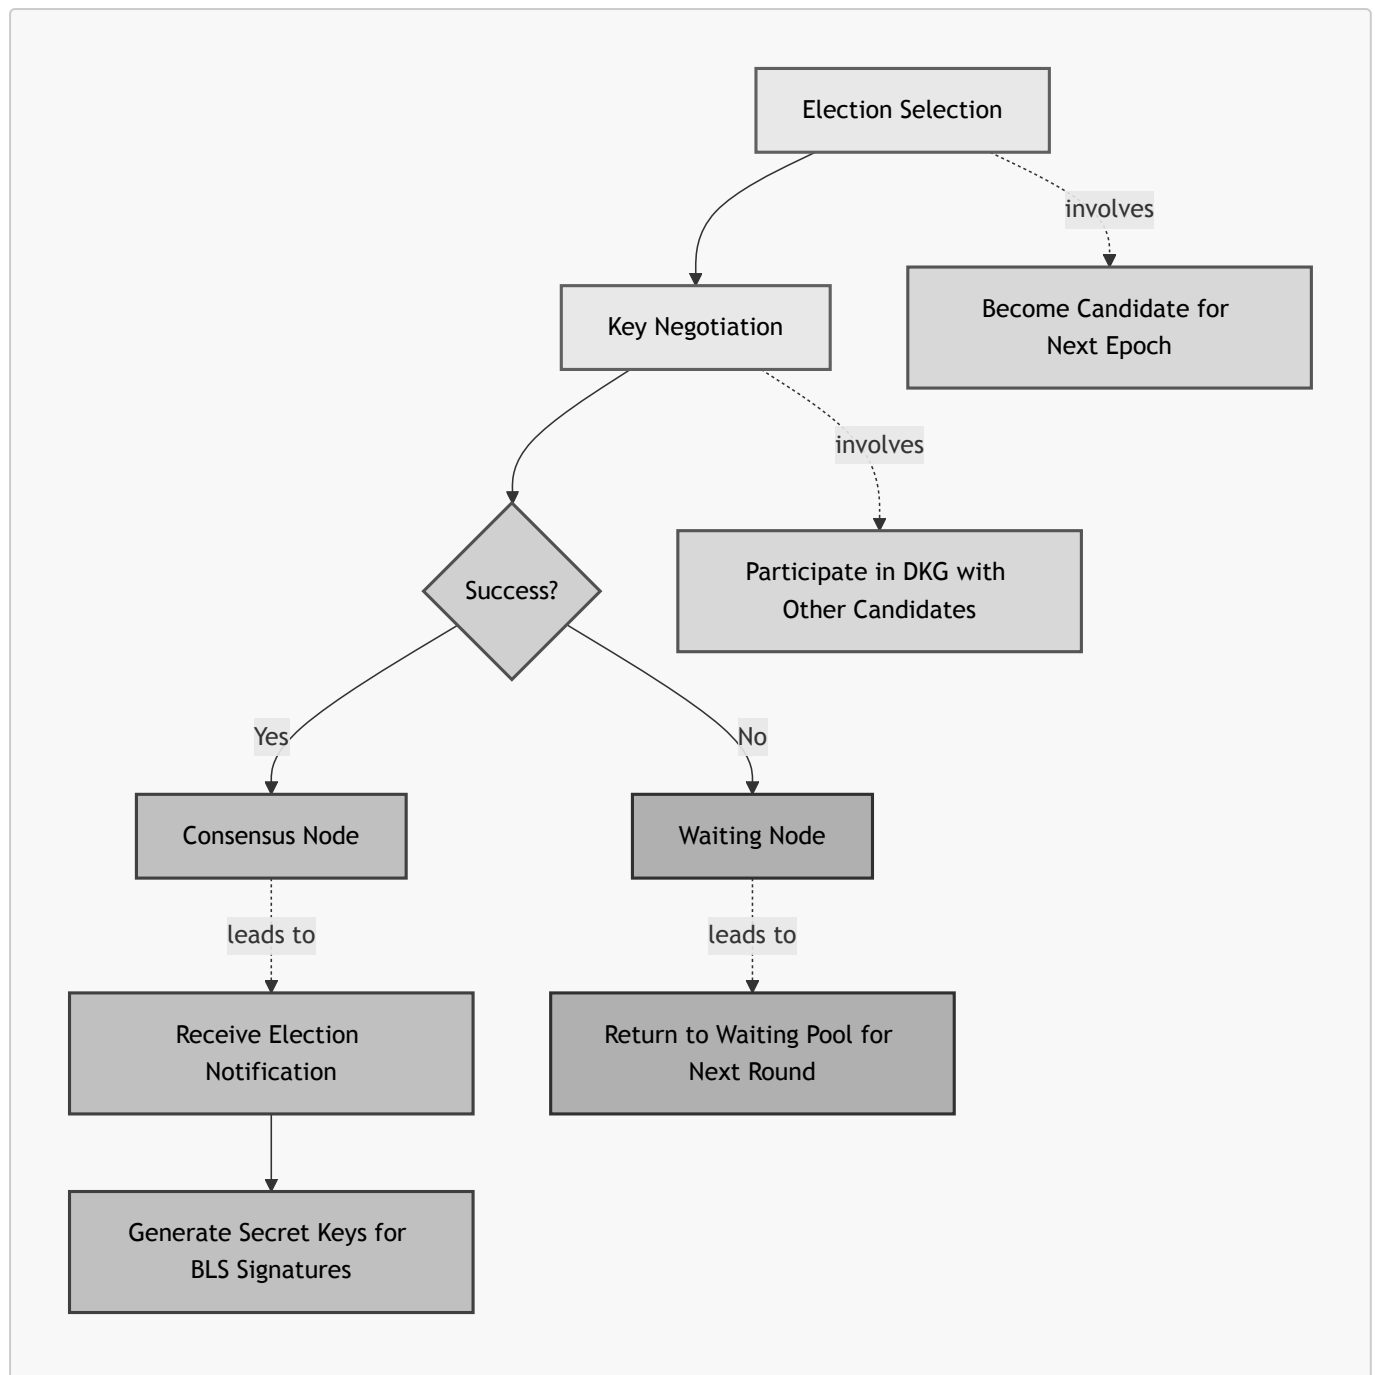
- **Performance Demonstration:** Continuously demonstrate readiness and capability for promotion to consensus roles
- **Reputation Building:** Accumulate reputation through consistent availability and reliable service provision
- **Key Preparation:** Participate in key negotiation processes when selected as consensus candidates

4.2.2.3. Candidate Nodes

Candidate nodes occupy a critical transitional state between waiting and consensus roles, representing nodes that have been selected for potential promotion but must complete additional requirements before

assuming full consensus responsibilities.

Key Negotiation Process:



Distributed Key Generation (DKG) Participation:

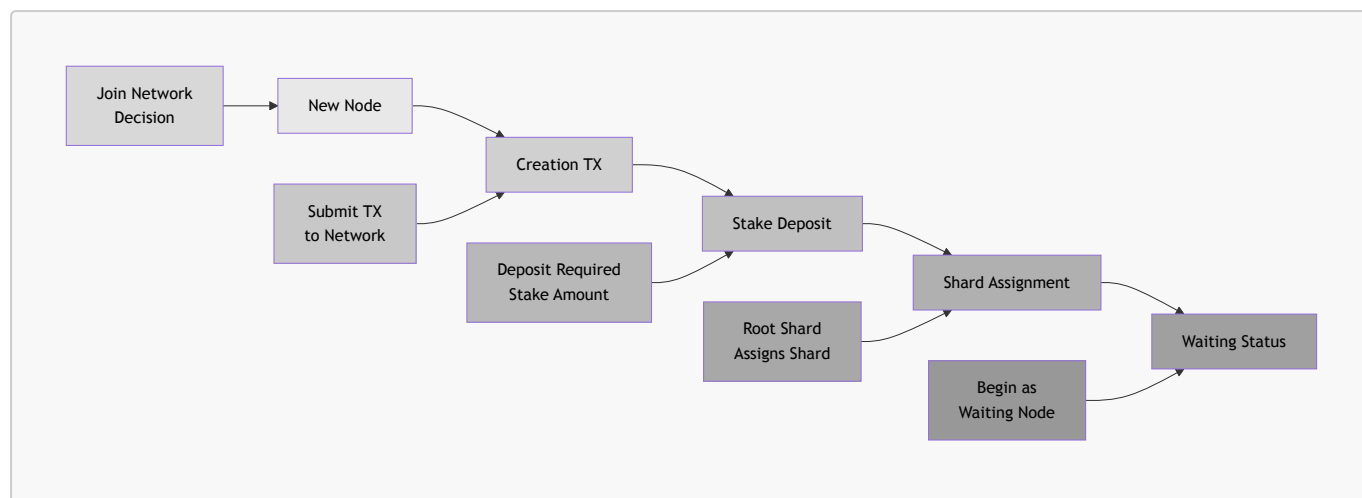
- Secret Share Generation: Create and distribute secret key contributions using polynomial-based cryptographic schemes
- Verification Vector Broadcasting: Share verification vectors to enable other candidates to validate secret contributions
- Threshold Signature Preparation: Collaborate to establish threshold signature capabilities for consensus participation

Promotion or Demotion Process:

- Success Path: Successful key negotiation completion leads to immediate promotion to consensus node status
- Failure Path: Key negotiation failure results in demotion back to waiting node status with opportunity for future selection
- Performance Evaluation: Success rates in key negotiation processes contribute to reputation score calculations

4.2.3. Node Lifecycle and Transitions

4.2.3.1. Node Entry Process



Creation Transaction Requirements:

- Stake Commitment: New nodes must deposit a minimum stake amount as economic security
- Identity Establishment: Cryptographic identity creation and registration with the network
- Preference Declaration: Nodes specify preference for root shard or transaction shard participation
- Resource Declaration: Nodes declare their computational, bandwidth, and storage capabilities

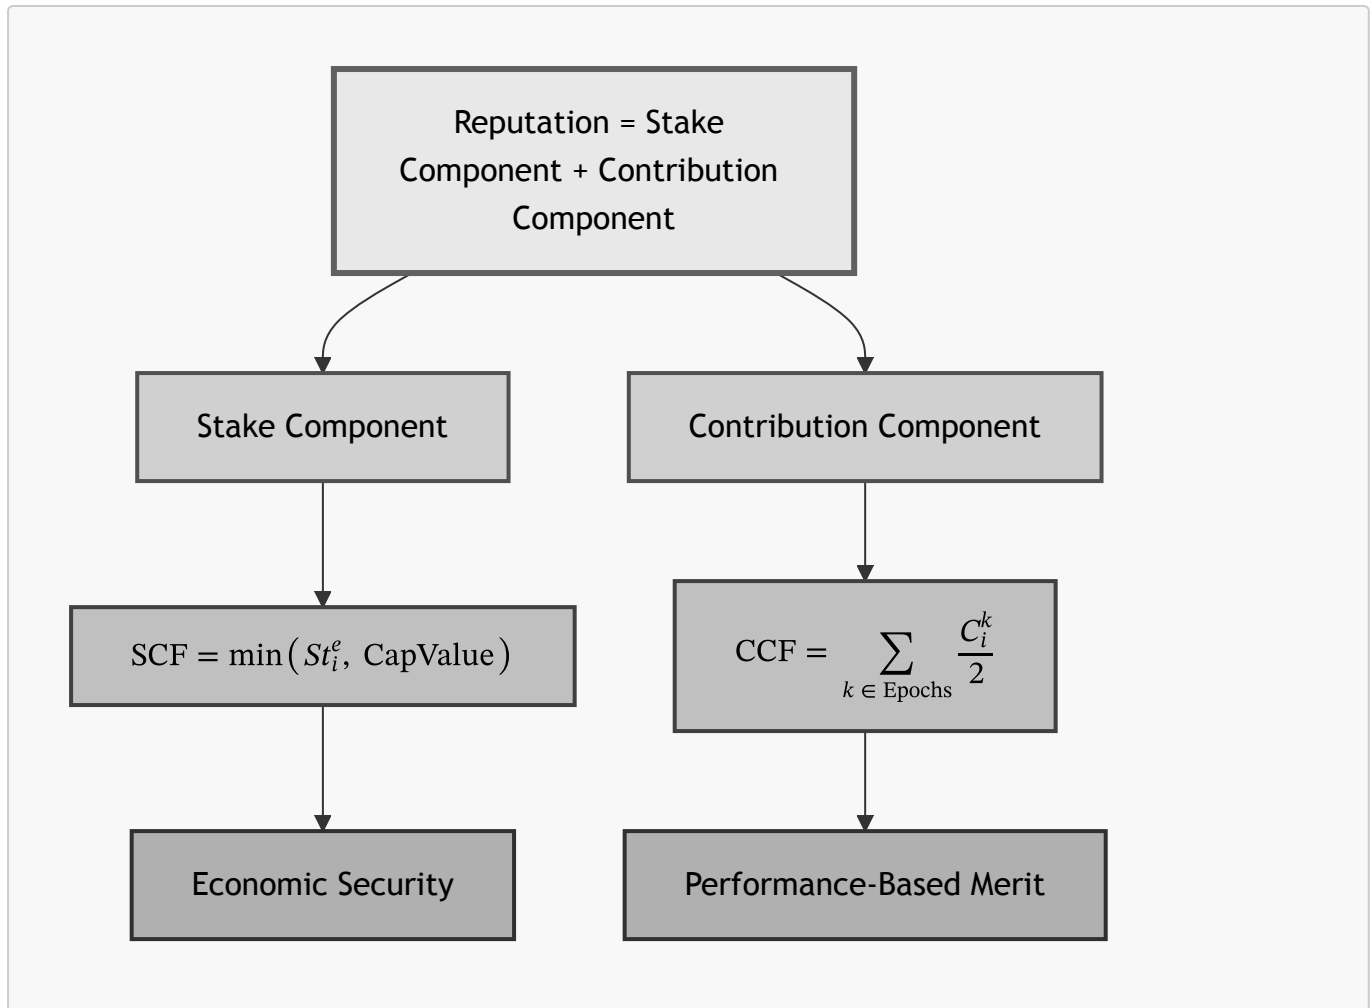
Shard Assignment Algorithm:

- Balanced Distribution: Root shard assigns nodes to transaction shards with lowest overall reputation to maintain balance
- Preference Consideration: Node preferences are considered within network balance constraints
- Capacity Assessment: Assignment considers shard capacity and node capabilities

4.2.3.2. Reputation-Based Promotion System

The promotion system operates on continuous performance evaluation rather than arbitrary selection, ensuring that the most capable and reliable nodes assume critical responsibilities.

Reputation Score Calculation:



Stake Component (St_i^e):

- Initial Stake: Required minimum deposit for network participation
- Accumulated Rewards: Block rewards automatically added to stake
- Additional Deposits: Voluntary stake increases to improve reputation
- Smoothing Function: Stake values are capped and normalized to prevent excessive wealth concentration

Contribution Component (C_i^k):

- Block Creation Success: Rewards for successfully proposing and confirming blocks
- Transaction Processing: Credit for processing valid transactions as consensus leader
- Gas Fee Optimization: Bonus for efficient transaction batching and gas utilization
- Penalty System: Negative scores for malicious behavior or poor performance

4.2.3.3. Demotion and Security Mechanisms

Direct Eviction Criteria:

- Lowest Performers: Bottom k_1 nodes in each epoch face automatic demotion to other shards
- Reputation Thresholds: Nodes falling below minimum reputation thresholds face immediate review
- Malicious Behavior: Detected Byzantine behavior results in immediate eviction and stake penalties

Random Eviction for Security:

- Stochastic Selection: k_2 additional nodes are randomly selected for eviction to prevent predictable patterns
- FTS Algorithm Application: Random selection weighted by inverse reputation scores to maintain fairness
- Security Enhancement: Prevents adversaries from gaming the system through reputation manipulation

Economic Penalties:

- Stake Slashing: Malicious behavior results in partial or complete stake forfeiture
- Reputation Damage: Poor performance permanently impacts historical reputation scores
- Recovery Pathways: Honest nodes can recover reputation through consistent good performance

4.2.4. Committee Formation and Coordination

4.2.4.1. Dual Committee Architecture

The dual committee structure represents a fundamental innovation in blockchain sharding, enabling continuous operation during reconfiguration by separating transaction processing from auxiliary functions.

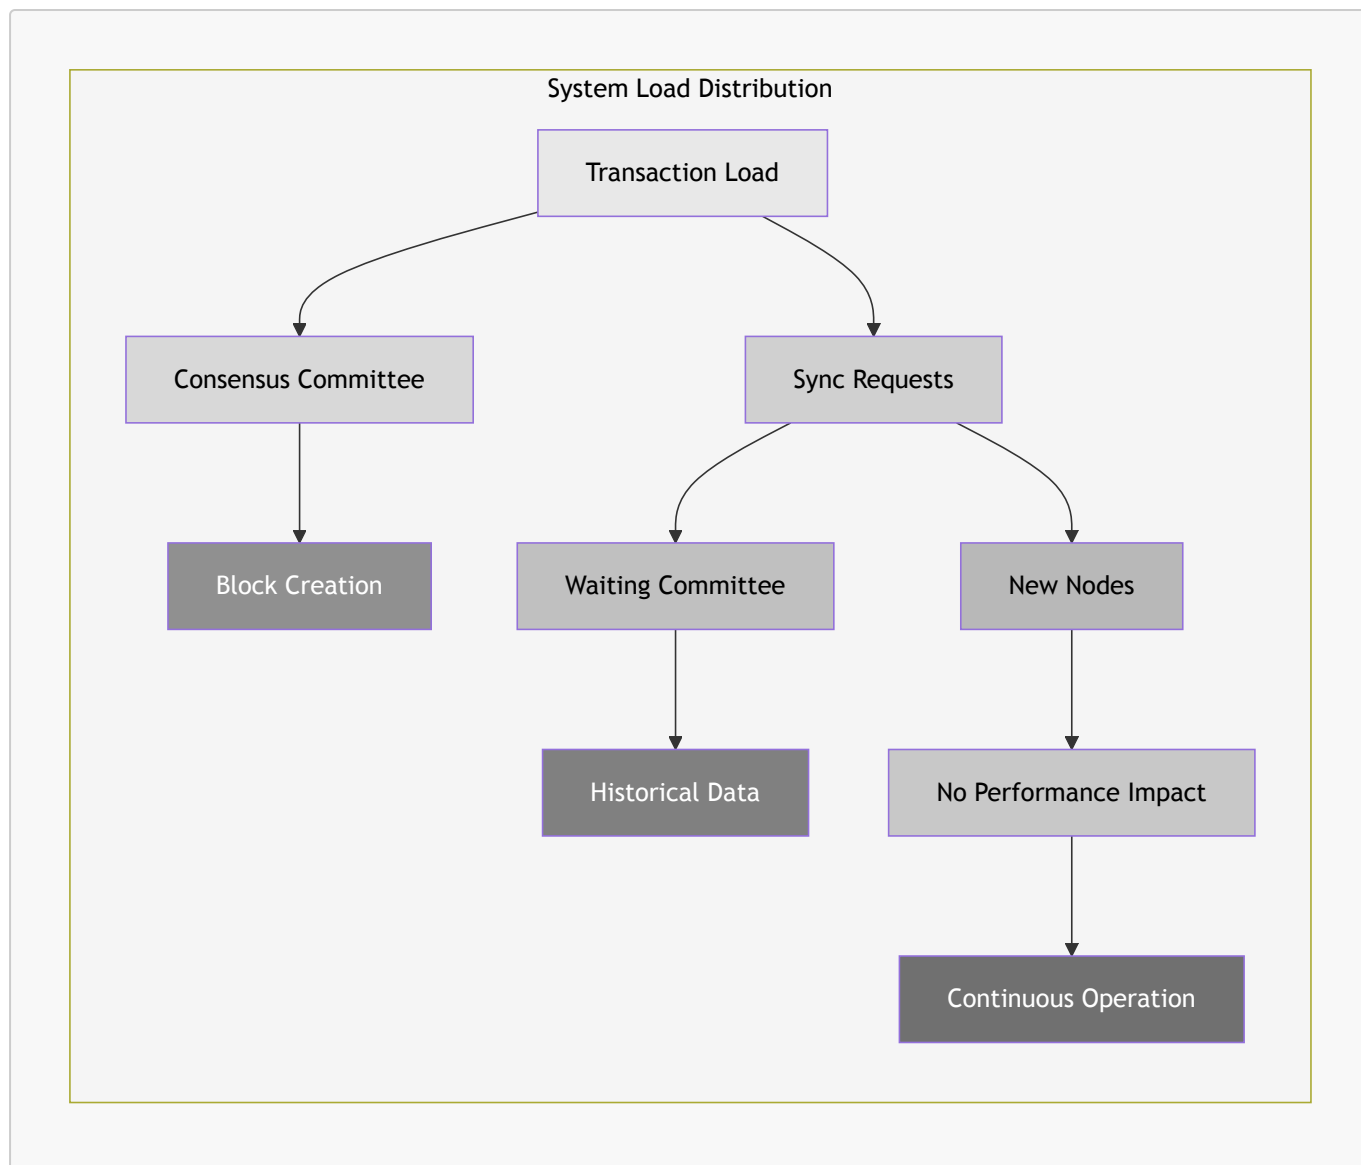
Consensus Committee Characteristics:

- Size: Typically n_c nodes selected from highest reputation participants
- Threshold: Requires $t \geq (2/3)n_c + 1$ nodes for BFT consensus
- Rotation: Partial rotation every epoch to maintain security while preserving operational continuity
- Performance Focus: Optimized for high-throughput transaction processing

Waiting Committee Characteristics:

- Size: Remaining shard nodes not in consensus committee
- Support Role: Provides ledger synchronization and backup services
- Election Pool: Source of candidates for next epoch's consensus committee
- Continuous Operation: Maintains readiness for immediate promotion when needed

4.2.4.2. Inter-Committee Coordination Mechanisms



- Load Isolation: Consensus nodes focus exclusively on transaction processing during reconfiguration
- Parallel Operations: Waiting nodes handle synchronization without impacting consensus performance
- Dynamic Adjustment: Committee sizes adjust based on network load and node availability

4.3. Transaction Model

LoraShard employs a sophisticated transaction architecture built upon the account/balance model, optimized for high-throughput parallel processing across multiple shards. Unlike traditional single-chain systems, LoraShard transaction model is designed to handle both intra-shard and cross-shard operations efficiently while maintaining strong consistency guarantees and atomic transaction properties.

The transaction model incorporates specialized transaction types that serve different network functions, from user-initiated transfers to system-level coordination messages. This multi-tiered approach enables LoraShard to maintain network synchronization, manage node participation, and coordinate complex multi-shard operations while preserving the simplicity and security of the underlying account/balance paradigm.

The system ensures transaction atomicity across shard boundaries through sophisticated two-phase protocols, while maintaining high performance through parallel execution within individual shards. This

design allows LoraShard to scale transaction throughput linearly with the number of shards while preserving the security and consistency properties essential for a production blockchain system.

4.3.1. Account/Balance Model Foundation

LoraShard adopts the **account/balance model** as its fundamental transaction paradigm, where each user maintains an account with an associated balance that is modified through transaction execution. This model provides intuitive transaction semantics and efficient state management compared to UTXO-based alternatives.

Account State Structure:

Account ID: 0x1a2b3c...
Balance: 1000.0 LRS
Nonce: 42
Shard Assignment: Shard #2
Contract Code: [if applicable]
Storage State: {key: value pairs}

Balance Verification Protocol:

- Pre-execution Validation: Every transaction undergoes balance sufficiency checks before execution
- State Consistency: Account states are maintained consistently across transaction executions
- Atomic Updates: Balance modifications occur atomically to prevent double-spending and inconsistent states

Account Distribution Strategy:

- Deterministic Shard Assignment: Accounts are assigned to shards using consistent hashing algorithms
- Load Balancing: Account distribution aims to balance transaction loads across shards
- Migration Support: Accounts can be reassigned to different shards during network rebalancing operations

4.3.2. Transaction Type Classification

4.3.2.1. Time Transactions

Time transactions serve as the global coordination mechanism for the LoraShard network, enabling synchronized epoch transitions and network-wide operational coordination. These transactions are exclusively generated and processed by the root shard.

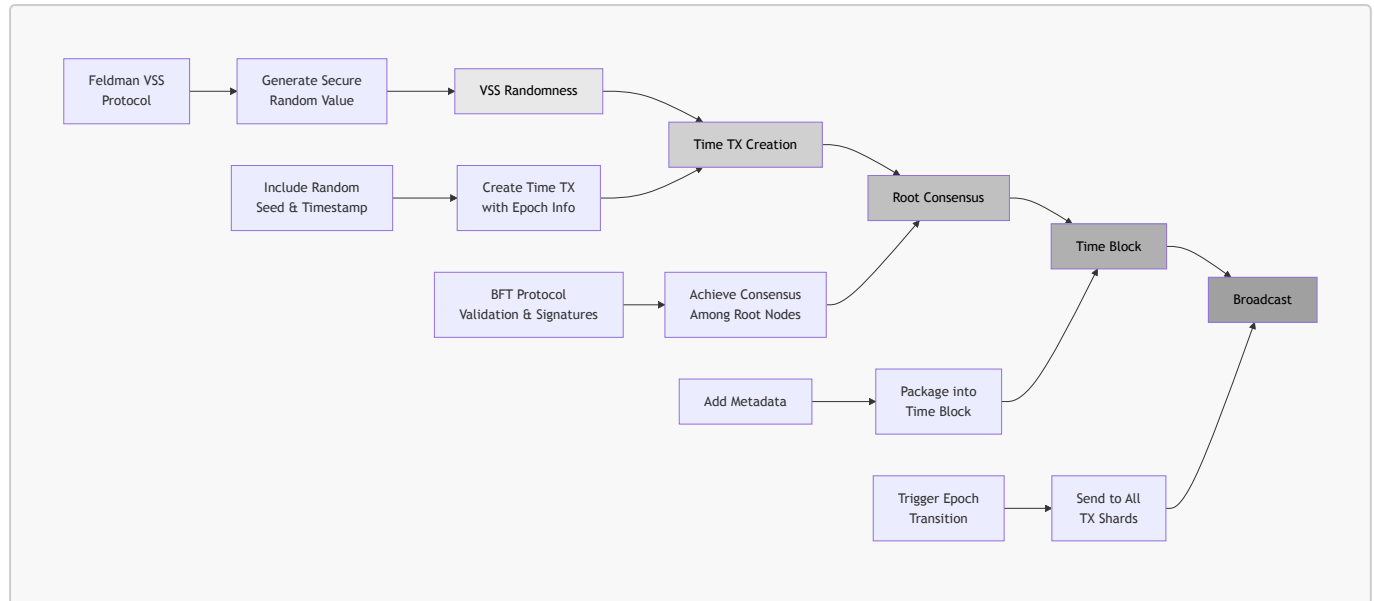
Time Transaction Structure:

Time Transaction Components:

Type: TIME_TRANSACTION
Epoch Number: 1547

Timestamp: 2025-06-23T10:30:00Z
Random Seed: 0xa1b2c3d4...
Root Signature: BLS_Signature
Verification Data: Merkle_Proof

Time Block Generation Process:



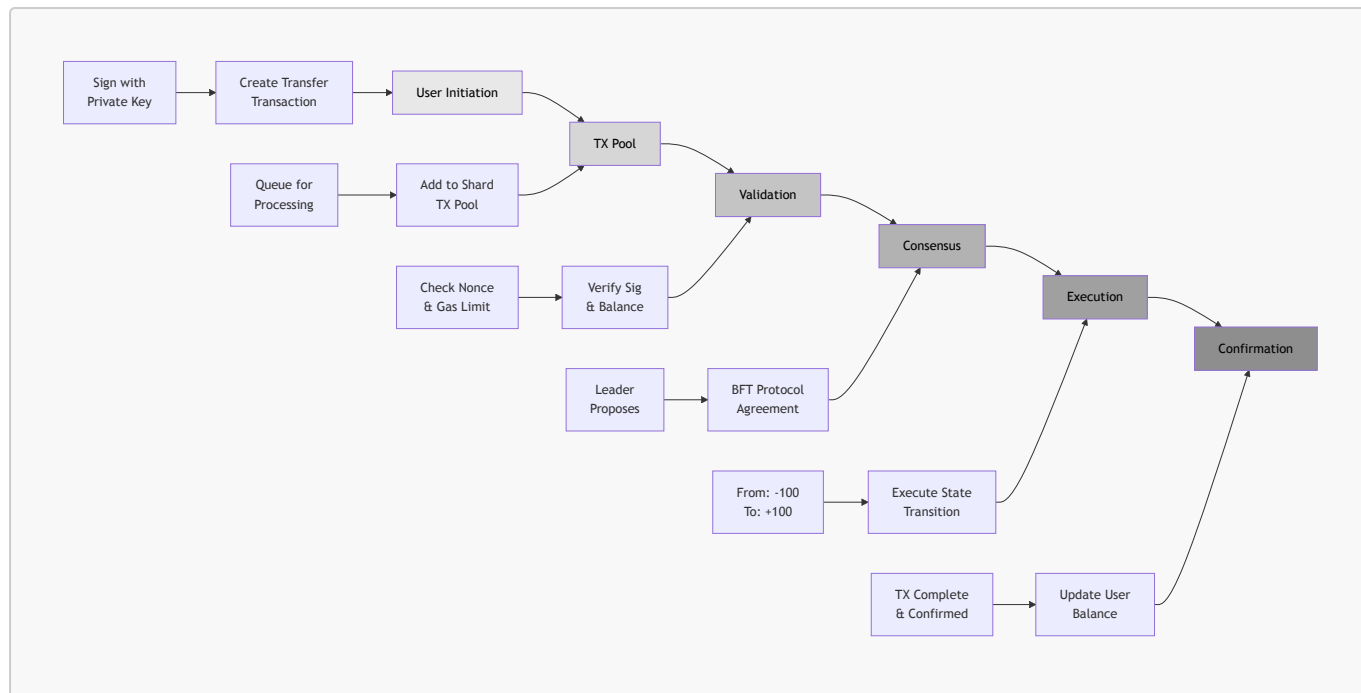
Operational Mechanics:

- Periodic Generation: Time transactions are created at predetermined intervals to maintain network rhythm
- Epoch Synchronization: Receipt of time blocks triggers new epoch initiation across all transaction shards
- Randomness Distribution: Embedded random values provide entropy for node selection algorithms
- Global Clock Reference: Time blocks establish network-wide temporal consistency

4.3.2.2. Transfer Transactions

Transfer transactions represent the primary user-facing transaction type, enabling value transfers between accounts within the LoraShard ecosystem. These transactions are optimized for high-throughput processing while maintaining security and consistency.

Intra-Shard Transfer Flow:

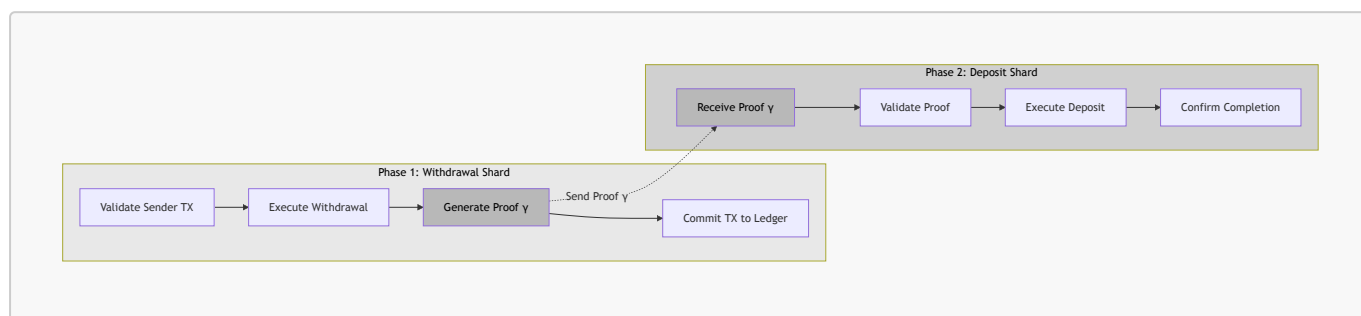


Atomic Operation Components:

- Withdrawal Phase: Sender account balance is decremented by transfer amount plus fees
- Deposit Phase: Receiver account balance is incremented by transfer amount
- Fee Distribution: Transaction fees are distributed to consensus nodes and network treasury
- State Consistency: All balance changes occur atomically within a single transaction execution

Cross-Shard Transfer Protocol:

Cross-shard transfers require sophisticated coordination between multiple shards to maintain atomicity and consistency across shard boundaries.



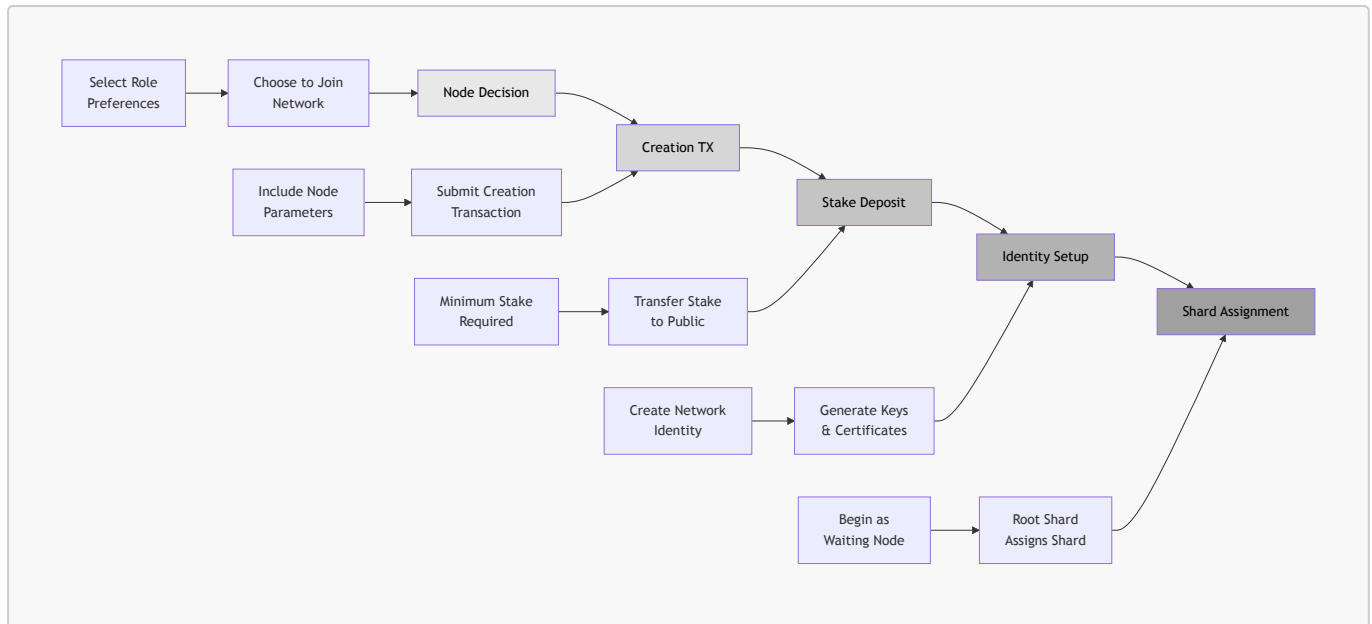
Proof of Acceptance Generation:

- Aggregated Signature: BLS threshold signature from withdrawal shard consensus committee
- Transaction Inclusion Proof: Merkle proof demonstrating transaction inclusion in committed block
- Header Hash: Block header hash providing additional verification context
- Account State Proof: Merkle proof of sender account state changes

4.3.2.3. Creation Transactions

Creation transactions provide the **entry mechanism** for new nodes to join the LoraShard network, establishing their identity, staking requirements, and initial role assignments.

Node Registration Process:



Creation Transaction Components:

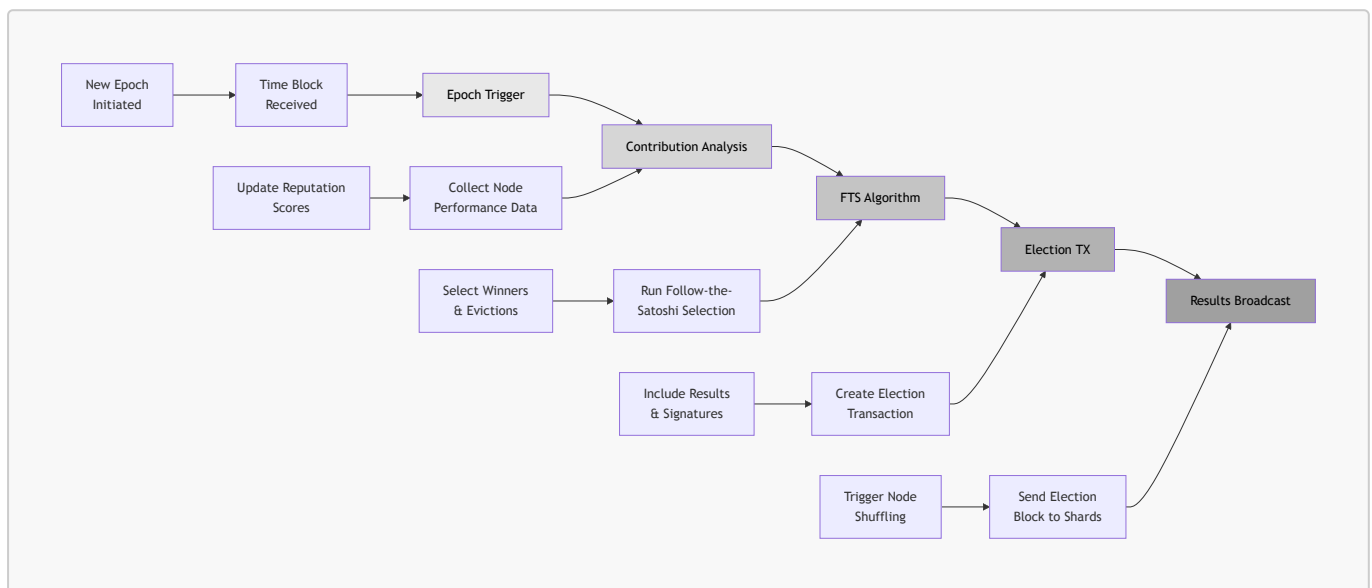
- Node Identity: Public key and cryptographic identity establishment
- Stake Commitment: Economic security deposit transferred to network treasury
- Role Preferences: Declaration of desired participation role (root vs transaction shard)
- Capability Declaration: Hardware and network capacity specifications
- Geographic Information: Optional location data for network optimization

Validation and Processing:

- Stake Verification: Confirmation that deposit meets minimum network requirements
- Identity Uniqueness: Prevention of duplicate node registrations
- Capacity Assessment: Evaluation of declared capabilities against network standards
- Shard Assignment Algorithm: Root shard determines optimal placement based on network balance

4.3.2.4. Election Transactions

Election transactions represent critical governance operations that coordinate node shuffling, committee formation, and network reconfiguration across the LoraShard ecosystem.



Election Transaction Structure:

- Shard Identifier: Target shard for election results
- Elected Nodes: List of nodes selected for consensus committee promotion
- Evicted Nodes: List of nodes selected for committee eviction
- Assignment Changes: Cross-shard reassignments for network balancing
- Randomness Proof: Cryptographic proof of fair selection process

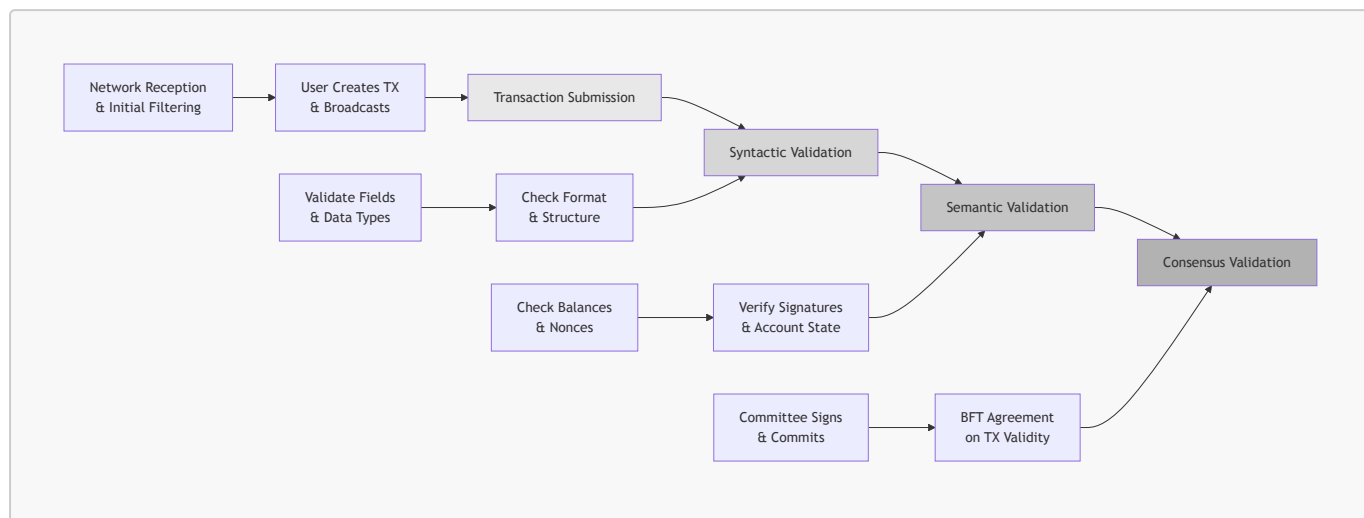
Follow-the-Satoshi Algorithm Integration:

- Reputation Weighting: Node selection probabilities based on accumulated reputation scores
- Merkle Tree Construction: Efficient data structure for weighted random selection
- Verifiable Randomness: Use of epoch random seeds for transparent and unbiased selection
- Security Guarantees: Cryptographic proofs prevent manipulation of selection process

4.3.3. Transaction Processing Pipeline

4.3.3.1. Validation and Verification Framework

LoraShard employs a comprehensive validation framework that ensures transaction integrity, prevents double-spending, and maintains network security through multiple verification layers.



Syntactic Validation Layer:

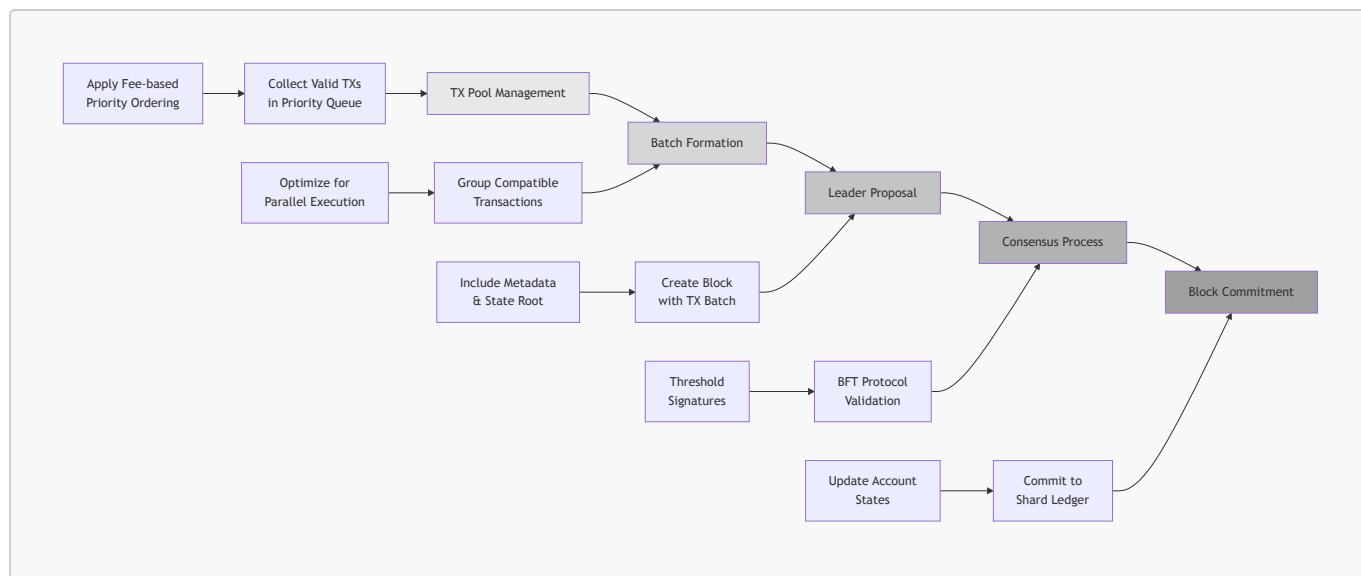
- Format Verification: Transaction structure compliance with protocol specifications
- Field Validation: Data type checking and range validation for all transaction fields
- Size Limitations: Enforcement of maximum transaction size and gas limit constraints
- Encoding Standards: Verification of proper data encoding and serialization

Semantic Validation Layer:

- Cryptographic Verification: Digital signature validation using sender's public key
- Account State Consistency: Verification of sender account existence and current state
- Balance Sufficiency: Confirmation of adequate account balance for transaction execution
- Nonce Sequence Validation: Prevention of replay attacks through nonce ordering checks

4.3.3.2. Consensus Integration and Block Formation

LoraShard optimizes throughput through **intelligent transaction batching** that groups compatible transactions for efficient parallel processing while maintaining ordering requirements.



Batching Optimization Criteria:

- Gas Efficiency: Maximizing transaction throughput within block gas limits
- Dependency Analysis: Grouping transactions with minimal state dependencies
- Fee Prioritization: Higher fee transactions receive processing priority
- Load Balancing: Even distribution of computational load across consensus nodes

Block Formation Process:

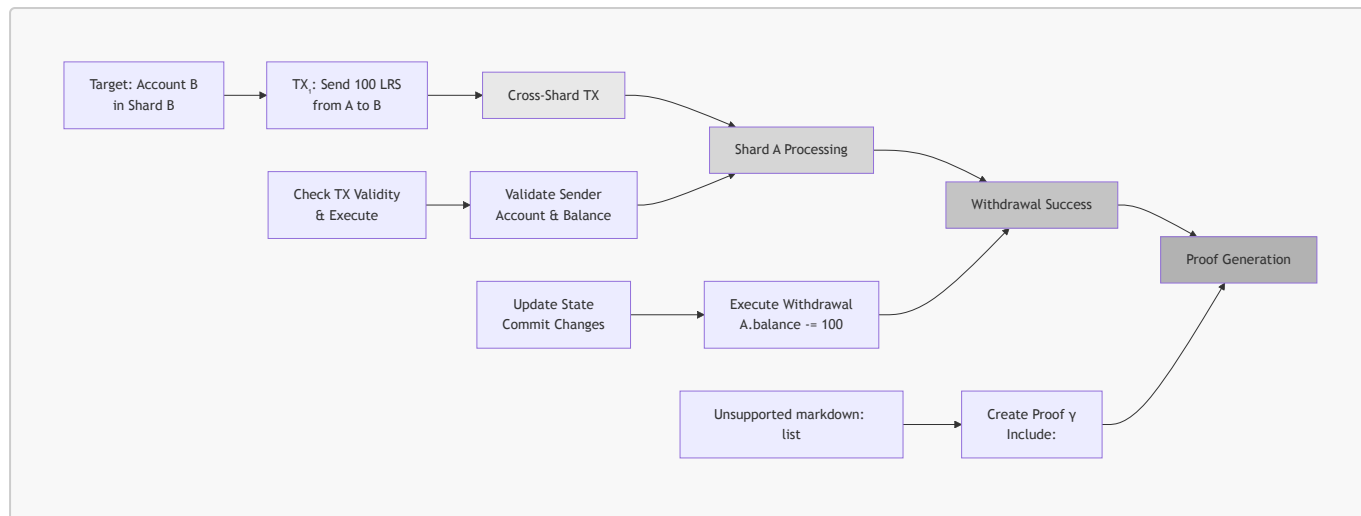
- Transaction Selection: Leader selects optimal transaction batch from pool
- State Root Calculation: Computation of Merkle root for resulting account states
- Block Header Construction: Assembly of block metadata, timestamps, and cryptographic proofs
- Consensus Proposal: Distribution of proposed block to consensus committee for validation

4.3.4. Cross-Shard Transaction Coordination

4.3.4.1. Two-Phase Atomic Protocol

Cross-shard transactions require sophisticated coordination protocols to ensure atomicity across multiple independent shards while maintaining high performance and security standards.

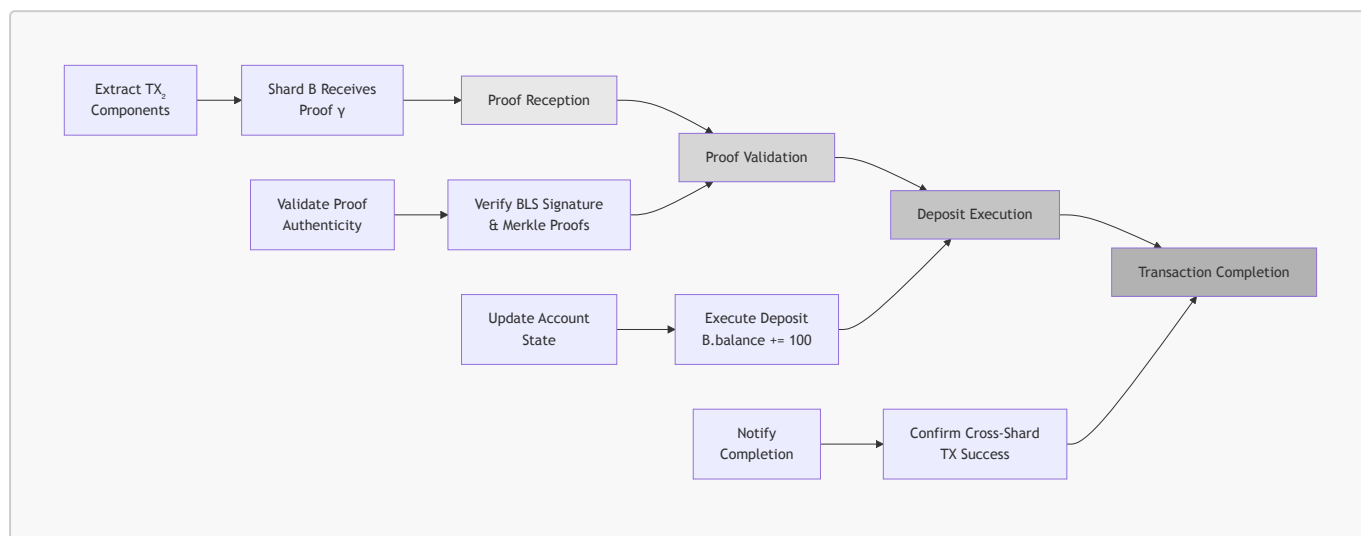
Phase 1: Withdrawal Execution:



Proof of Acceptance Components:

- Aggregated BLS Signature (σ_c): Threshold signature from consensus committee confirming transaction commitment
- Public Key Information (PK): Verification keys for signature validation with epoch identification
- Transaction Details: Complete withdrawal transaction data for verification
- Block Header Hash (Hheader): Hash of block containing the withdrawal transaction
- Merkle Proof (HTXs): Cryptographic proof of transaction inclusion in committed block

Phase 2: Deposit Execution:



- Atomic Execution: Either both withdrawal and deposit succeed, or neither occurs
- Cryptographic Verification: All state changes are cryptographically provable and verifiable
- Consensus Agreement: Both phases require consensus committee agreement for execution
- Rollback Protection: Failed cross-shard transactions do not result in partial state changes

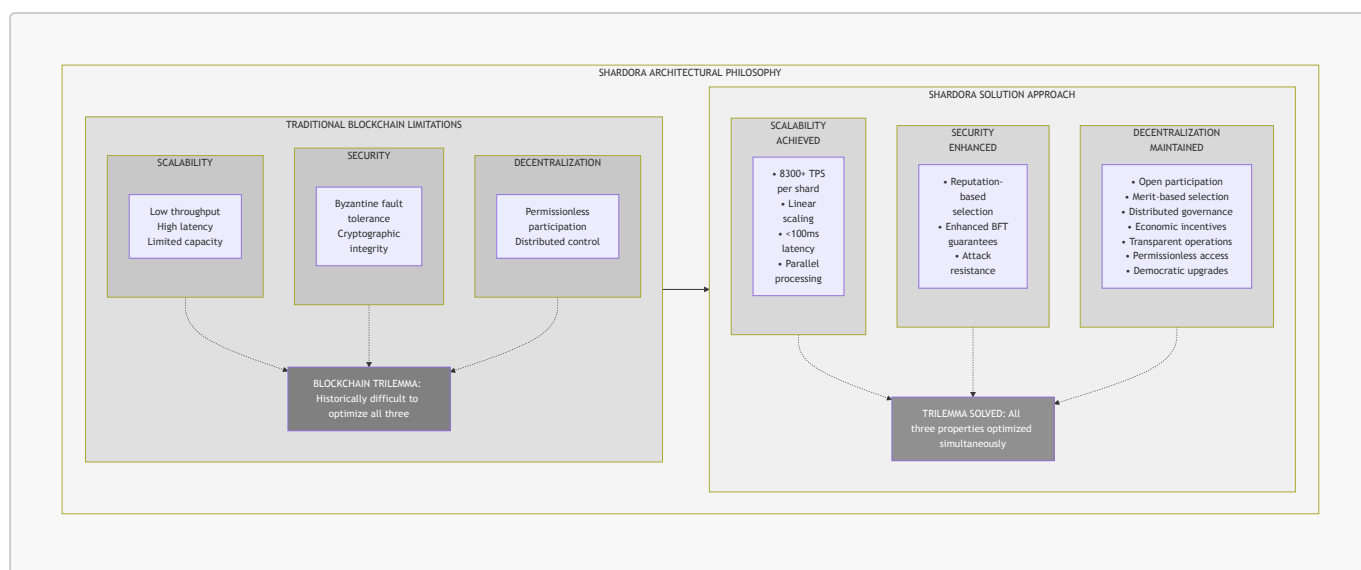
5. Deep Dive into Architecture

To fully appreciate LogaShard's innovations, it is beneficial to understand the foundational architectural principles common to advanced sharded blockchains, particularly those that are EVM-compatible and aim for linear scalability. While specific architectural diagrams for LogaShard itself are not provided, we can infer its likely structure and operational flow by examining the principles of similar systems like Sharding, which shares many core concepts of dynamic state sharding and transaction-level consensus.

5.1. High-Level Architecture (Referencing Shardeum as a Model)

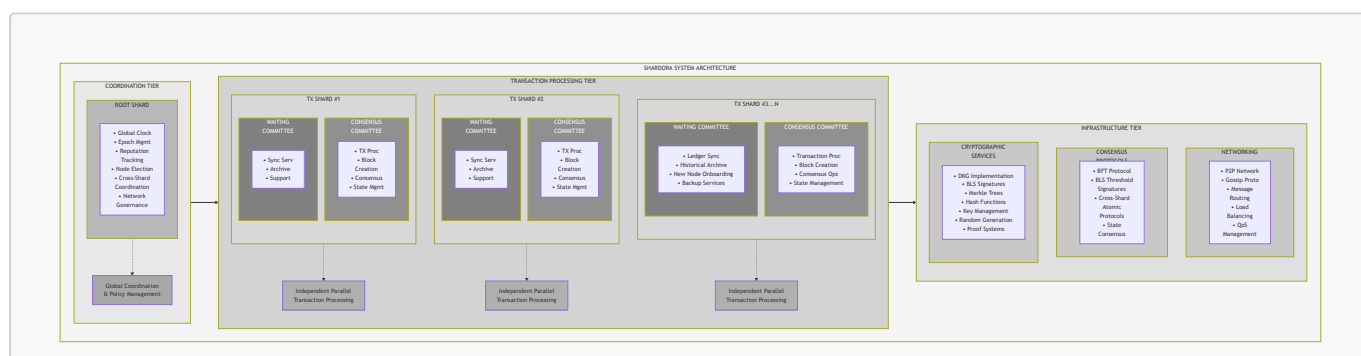
LoraShard represents a **paradigmatic advancement** in blockchain sharding architecture, built upon the foundational principle that **high performance and strong security are not mutually exclusive** when intelligent architectural decisions are made. The system demonstrates how sophisticated coordination mechanisms, parallel execution models, and reputation-based governance can work together to create a blockchain platform that scales linearly while maintaining the security guarantees essential for production deployment.

The architectural philosophy centers on **separation of concerns at every level** - from the temporal separation of operations through epoch structures to the functional separation of responsibilities through dual committee frameworks. This approach enables LoraShard to address the fundamental bottlenecks that have prevented previous sharding solutions from achieving practical scalability without compromising security or decentralization.



- Separation of Concerns: Different network functions handled by specialized components
- Parallel Execution: Multiple critical processes operate simultaneously without interference
- Merit-Based Governance: Performance and reliability drive participation rather than pure randomness
- Temporal Coordination: Structured time windows enable complex coordination without service disruption
- Graceful Degradation: System maintains core functionality under various failure scenarios

LoraShard implements a sophisticated multi-tier architecture that organizes network components according to their specialized functions while maintaining efficient communication and coordination pathways between all system elements.



5.2. Sharding Mechanism

Sharding is the cornerstone of scalability in networks like LogaShard. It involves breaking the network into smaller, independent pieces called shards, each processing a portion of the network's transactions in parallel.³

5.2.1. Dynamic State Sharding

Shardeum, and by extension, LogaShard, implements "dynamic state sharding". This is considered the most advanced and complex form of sharding because it dynamically partitions the network's

state, network, and transactions. Unlike static sharding, where shard structures remain constant, dynamic state sharding allows shards to adapt to real-time changes in network conditions, such as workload fluctuations.

This adaptability involves the ability to split or merge shards, add or remove nodes from shards, or reassign nodes to different shards as needed. This reorganization, driven by various algorithms that respond to real-time changes, helps maintain an optimal balance between the system's load and available resources, ensuring efficient and secure network operation. Key benefits of dynamic state sharding include:

- **Efficient Resource Utilization:** Resources are allocated dynamically based on demand, preventing bottlenecks and underutilization.
- **Adaptive Scalability:** The network's capacity scales automatically with increasing demand, adjusting the number and size of shards.
- **Atomic Composability:** This mechanism ensures that all transactions, including complex cross-shard ones, maintain atomic properties, meaning they either fully succeed or completely fail, preserving data integrity across the entire sharded network.
- **Synchronous Processing of Multi-Shard Transactions:** Despite being distributed, the system aims to process multi-shard transactions synchronously, reducing latency and complexity often associated with asynchronous cross-shard communication.

5.2.2. Types of Sharding (Contextual)

To understand dynamic state sharding, it's useful to differentiate it from other sharding types:

- **Network Sharding:** Partitions nodes into smaller, interconnected groups, each processing a subset of transactions and storing a portion of the state. Each shard operates independently with its own consensus algorithm, ledger, and transaction pool, communicating to maintain overall integrity.
- **Compute Sharding:** Nodes still store the entire blockchain, but only process transactions assigned to their shard. Cross-shard transactions require additional steps like a two-phase commit protocol.¹
- **Data Sharding:** Divides the blockchain's data into smaller pieces, with each shard responsible for data storage and maintenance, reducing storage burden and enabling faster data access.
- **State Sharding:** Divides the blockchain's state (account balances, smart contract data) into separate segments, reducing storage and computational requirements for individual nodes, making

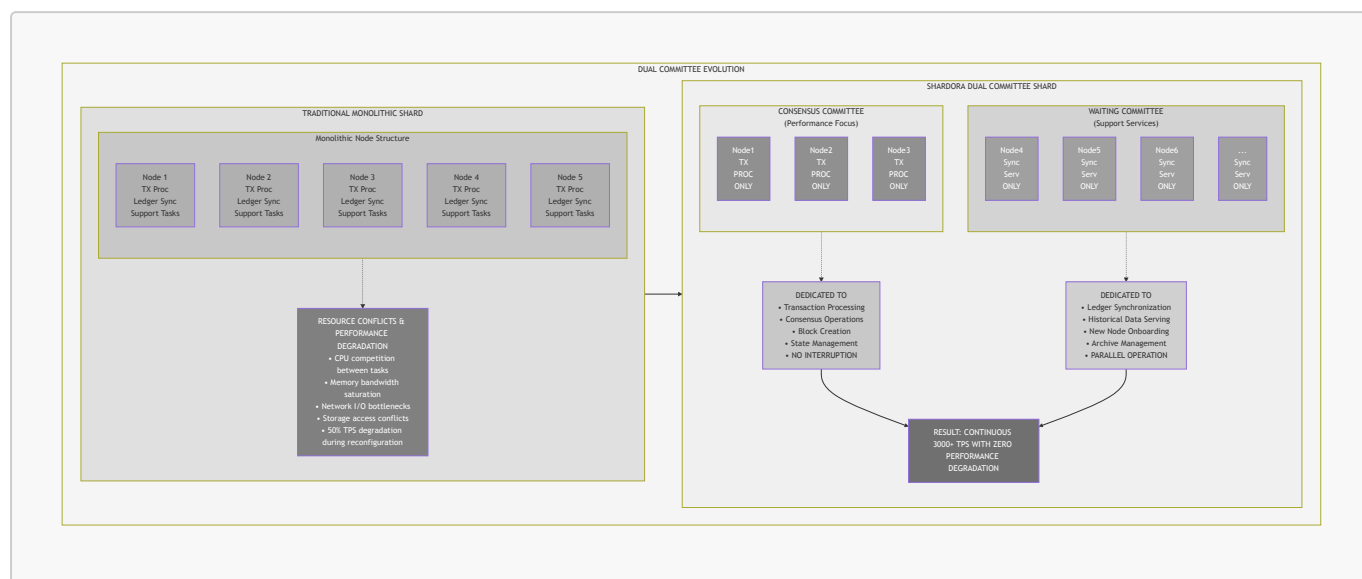
participation easier.

LogaShard's dynamic state sharding combines elements of these, dynamically adapting to optimize performance and security.

5.3. Dual Committee Structure

The Dual Committee Structure represents one of LoraShard's most innovative architectural contributions, fundamentally reimagining how blockchain shards organize and manage their node resources. This sophisticated approach addresses the core performance bottlenecks that have limited previous sharding systems by separating critical functions into specialized committees that operate in parallel without interfering with each other's performance characteristics.

Traditional blockchain shards suffer from the monolithic node problem where every node must perform identical functions, leading to resource conflicts when nodes simultaneously handle transaction processing and auxiliary services like ledger synchronization. LoraShard's dual committee approach eliminates these conflicts through intelligent functional decomposition, enabling each committee to optimize for its specific responsibilities while maintaining seamless coordination.



The dual committee structure enables **unprecedented performance consistency** by allowing consensus operations to proceed uninterrupted while synchronization and support services operate in parallel. This architectural innovation is the foundation that makes LoraShard's elimination of TPS-degradation and Zero-TPS issues possible.

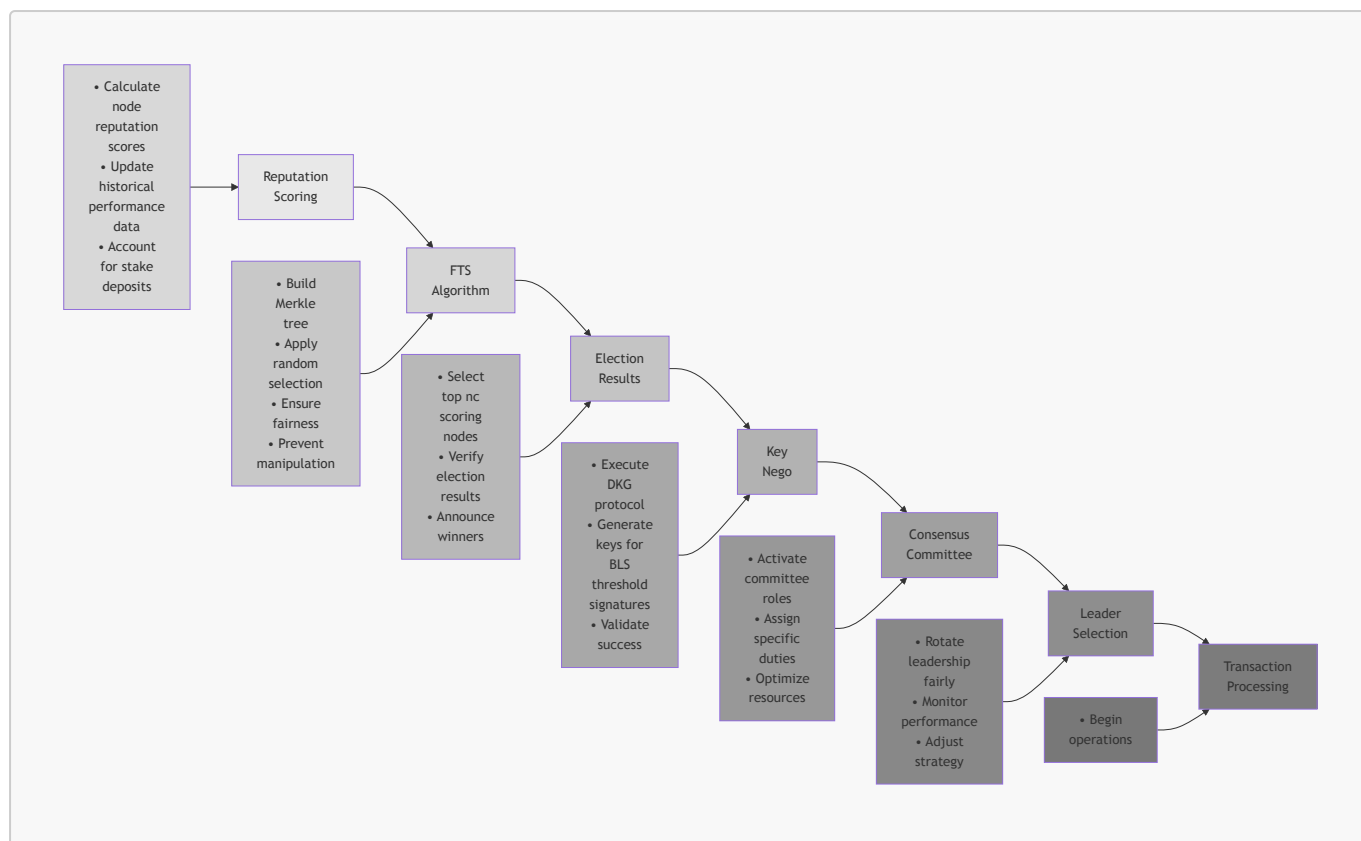
5.3.1. Consensus Committee

The Consensus Committee represents the performance-critical core of each LoraShard shard, comprised of the highest-reputation nodes optimized exclusively for transaction processing, block creation, and consensus operations. This committee operates as a high-performance transaction processing engine without any auxiliary responsibilities that could impact performance.

Committee Composition and Selection:

- **Size:** Configurable between 100-300 nodes based on shard requirements and network conditions
- **Selection Criteria:** Top-reputation nodes within the shard, chosen through the Follow-the-Satoshi (FTS) algorithm

- Term Duration: One epoch (typically 10 minutes) with partial rotation for security
- Rotation Rate: Approximately 20-30% of nodes are rotated each epoch to maintain security while preserving performance continuity



5.3.1.1. Transaction Processing Pipeline

The consensus committee operates a sophisticated multi-stage processing pipeline that maximizes throughput through intelligent resource management, parallel execution, and optimized workflow coordination.

Stage 1: Transaction Ingestion and Validation

The first stage focuses on efficient transaction reception and initial processing to ensure only valid transactions enter the processing pipeline:

- Network Reception: P2P message routing with rate limiting and priority queue management
- Format Validation: Transaction format verification, duplicate detection, and size checking
- Cryptographic Verification: Signature validation, balance checking, and nonce verification
- Priority Ranking: Gas price sorting, account activity analysis, and smart contract requirements assessment
- Batch Formation: Organizing transactions by priority, creating optimal batch sizes, and load balancing preparation

Stage 2: Parallel Batch Processing

The core processing stage implements advanced parallel execution with conflict resolution and state management:


- Intelligent Batching: Optimal batch sizing based on current network conditions and resource availability

- Conflict Resolution: Dependency analysis and parallel execution path determination
- Multi-threaded Execution: Parallel state access and validation across multiple processing threads
- BFT Protocol Rounds: Leader proposal, pre-vote collection, pre-commit processing, and threshold signature aggregation
- State Updates: Account balance modifications, nonce increments, and smart contract execution with rollback preparation

Stage 3: Consensus and Finalization

The final stage ensures secure consensus achievement and efficient block finalization:

- Consensus Agreement: Final validation with BFT protocol completion and signature aggregation
- Block Creation: Assembly of block structure, Merkle root calculation, and cryptographic proof generation
- State Commitment: Application of state changes, database persistence, and finalization confirmation
- Network Broadcast: Block propagation, transaction confirmations, and cross-shard notifications
- Committee Notification: Informing waiting committee of new block data and preparing for next processing cycle



TRANSACTION PROCESSING PERFORMANCE METRICS

TYPICAL PROCESSING TIMES

Transaction Ingestion
5-10ms per transaction



Validation and Sorting
15-25ms per batch



Parallel Execution
30-60ms depending on
complexity

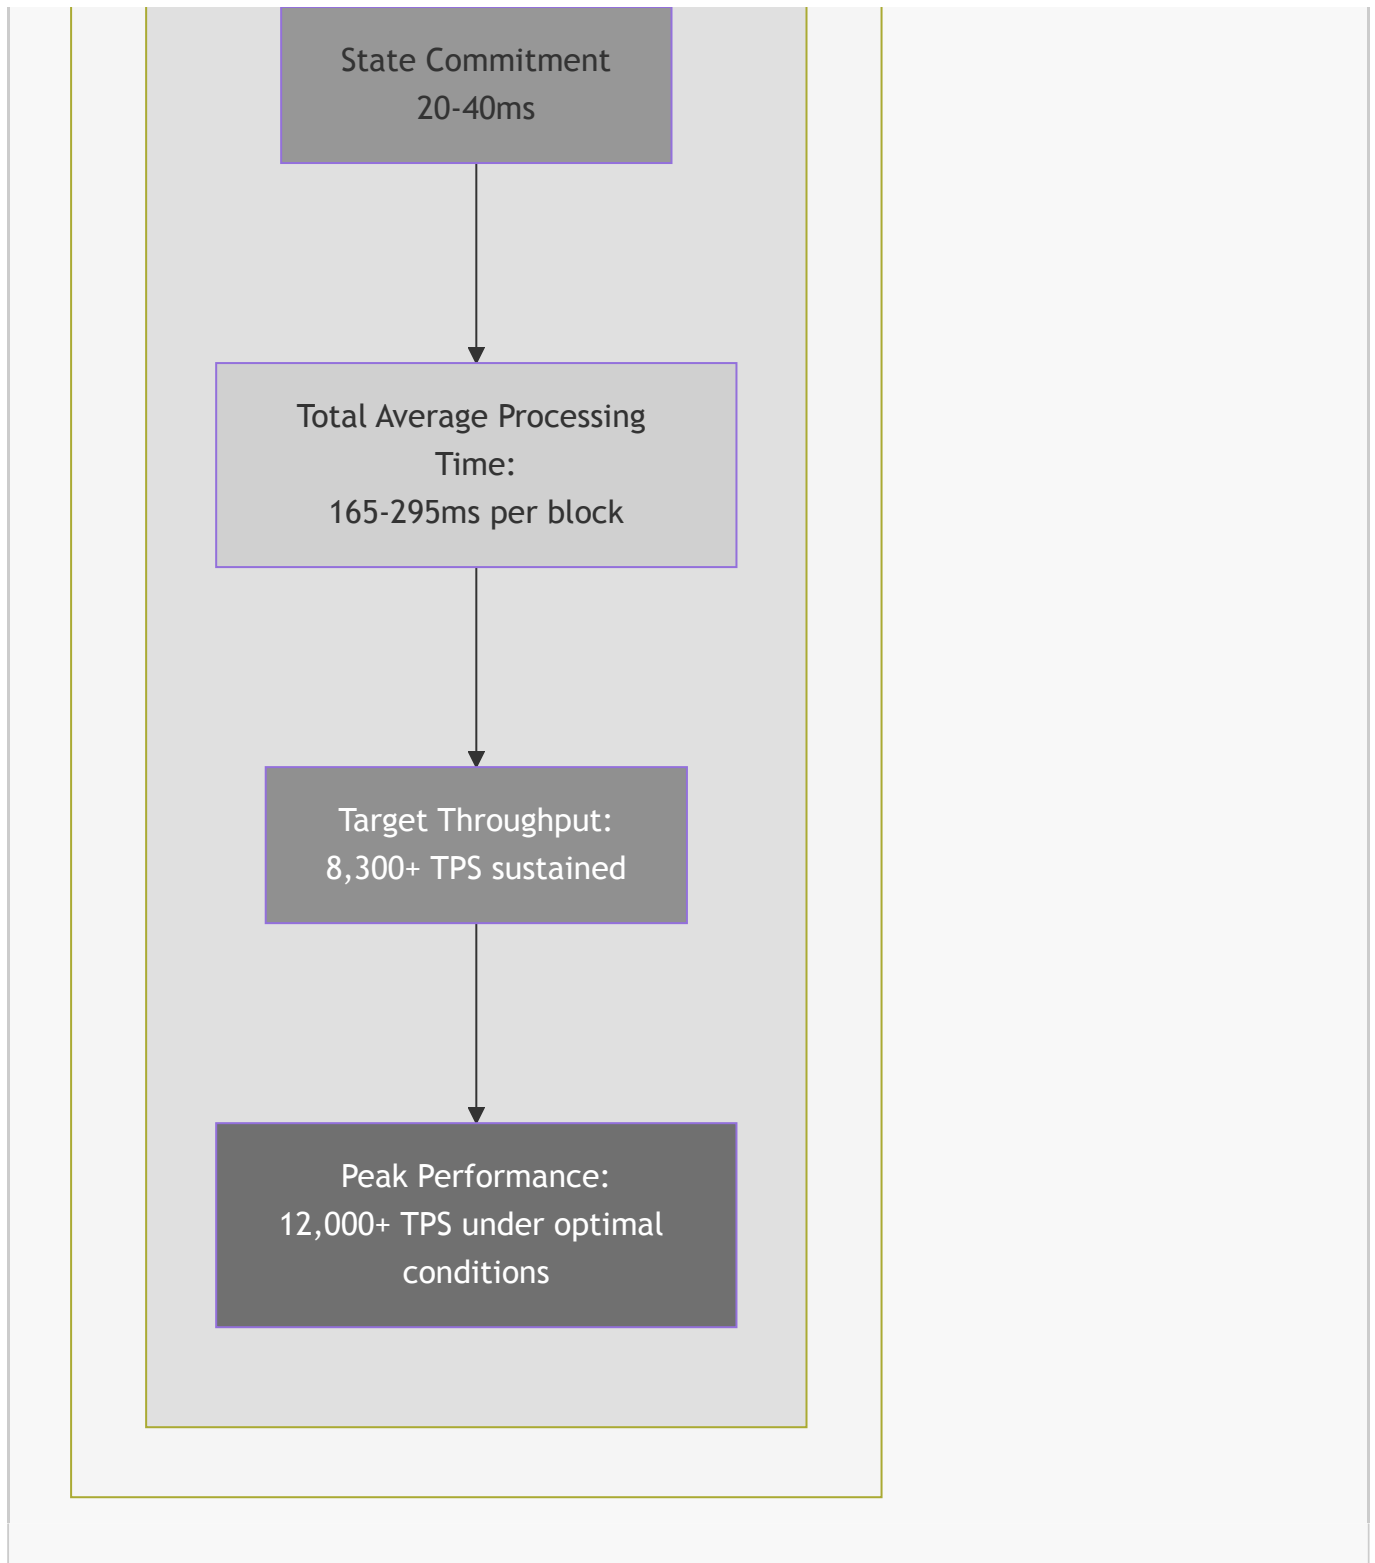


Consensus Rounds
85-140ms total (all BFT
phases)



Block Finalization
10-20ms





5.3.1.2. BFT Consensus with Threshold Signatures

The consensus committee implements an enhanced BFT protocol that combines traditional Byzantine fault tolerance with BLS threshold signatures to achieve both security and efficiency.

Consensus Round Structure:

Each consensus round follows a six-phase protocol designed to ensure safety, liveness, and optimal performance:

1. Leader Proposal Phase (10-20ms): Current leader creates block proposal with transaction batch, calculates state root, generates signature, and broadcasts to all committee members

2. Pre-Vote Processing (30-50ms): Committee members receive proposal, validate block header and transactions, sign with BLS secret share, and broadcast pre-vote messages
3. Pre-Commit Phase (20-30ms): Leader collects pre-vote signatures, verifies threshold reached, aggregates signatures using Lagrange interpolation, and broadcasts pre-commit message
4. Commit Processing (15-25ms): Committee members verify aggregated pre-vote signature, validate consensus threshold, generate commit signatures, and send to leader
5. Final Commit (10-15ms): Leader aggregates commit signatures, performs final validation, confirms consensus decision, and prepares block finalization
6. Block Completion (10-15ms): Archive block data, update local ledger state, notify waiting committee, and prepare for next round

Threshold Signature Implementation:

The BLS threshold signature mechanism provides linear communication complexity while maintaining strong security guarantees:

- Key Generation: Distributed Key Generation (DKG) protocol creates shared secret keys during committee formation
- Individual Signing: Each node signs messages with its secret share using BLS signature scheme
- Signature Collection: Leader collects individual signatures until threshold ($\lceil 2/3 \rceil n + 1$) is reached
- Aggregation Process: Lagrange interpolation combines individual signatures into single threshold signature
- Verification: Any node can verify aggregated signature using system public key with bilinear pairing

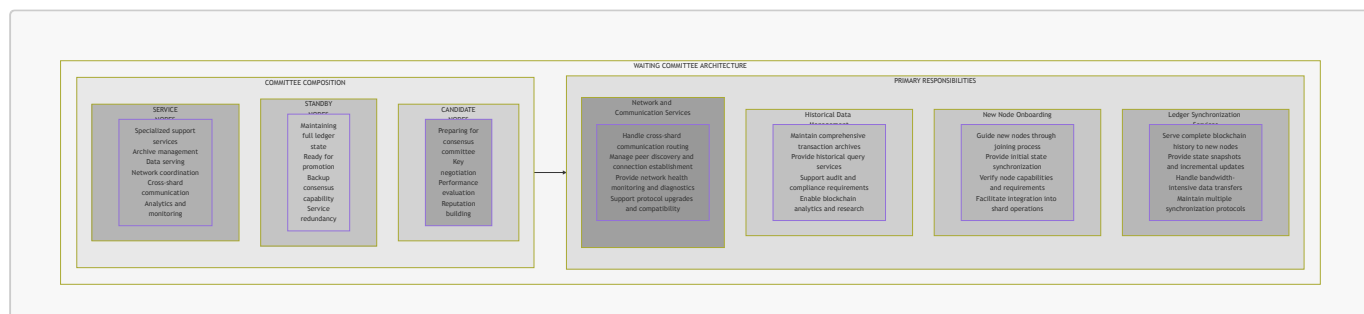
5.3.2. Waiting Committee

The Waiting Committee serves as the specialized support infrastructure of each LoraShard shard, handling all auxiliary services that traditionally burden consensus nodes in other blockchain systems. This committee operates in parallel with the consensus committee, ensuring that support services never interfere with transaction processing performance.

Committee Structure and Composition:

The waiting committee is designed as a flexible, scalable service layer that adapts to shard requirements:

- Variable Size: Typically 2-5 times larger than consensus committee, automatically scaling based on synchronization load
- Node Types: Mix of candidate nodes (preparing for consensus committee), standby nodes (maintaining readiness), and service nodes (specialized for specific support functions)
- Entry Mechanism: New nodes joining the shard automatically become waiting committee members
- Promotion Path: High-performing waiting nodes can be elected to consensus committee during reshuffling



6.3.2.1. Ledger Synchronization Services

The waiting committee's primary function is providing comprehensive ledger synchronization services without impacting consensus committee performance. This system implements multiple synchronization strategies optimized for different scenarios.

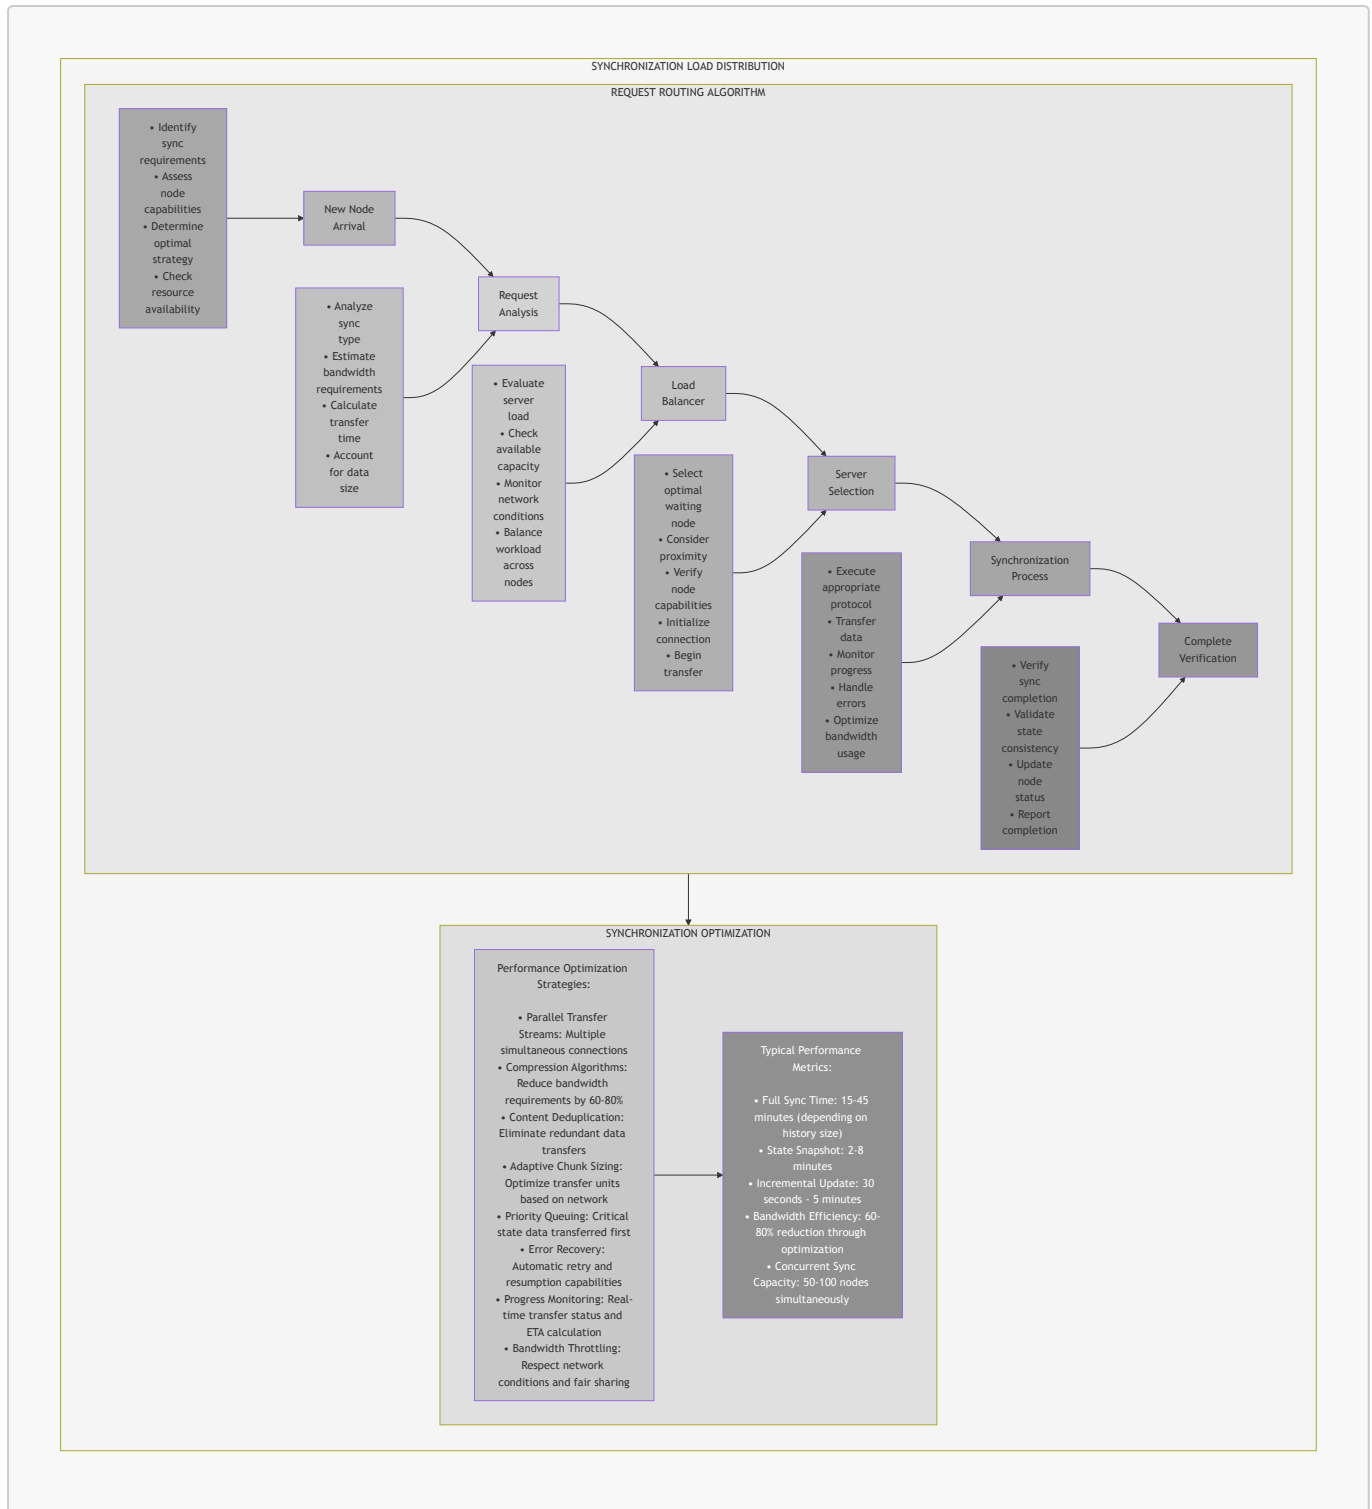
Multi-Protocol Synchronization Framework:

The waiting committee implements diverse synchronization protocols to handle various node requirements and network conditions:

- Full Synchronization: Complete blockchain history transfer for new nodes requiring full validation capability
- State Snapshot Sync: Current state transfer with merkle proofs for nodes needing quick integration
- Incremental Updates: Differential synchronization for nodes that have been offline temporarily
- Fast Sync Protocols: Optimized synchronization for nodes with limited bandwidth or storage
- Checkpoint Synchronization: Verified state checkpoints for rapid bootstrap with security guarantees

Intelligent Load Distribution:

The synchronization system implements sophisticated load balancing to ensure efficient resource utilization:



5.3.2.2. Historical Data and Archive Management

The waiting committee maintains comprehensive historical data services that support blockchain analytics, compliance requirements, and research activities without burdening consensus operations.

Archive Architecture:

- Tiered Storage System: Hot data (recent blocks), warm data (medium-term history), and cold data (long-term archives)
- Indexing Services: Advanced indexing for efficient historical queries and transaction lookup
- Compression and Optimization: Data compression algorithms reducing storage requirements by 70-85%
- Redundancy and Backup: Multiple copies across different waiting nodes ensuring data availability

- Query Optimization: Specialized query engines for historical analysis and research

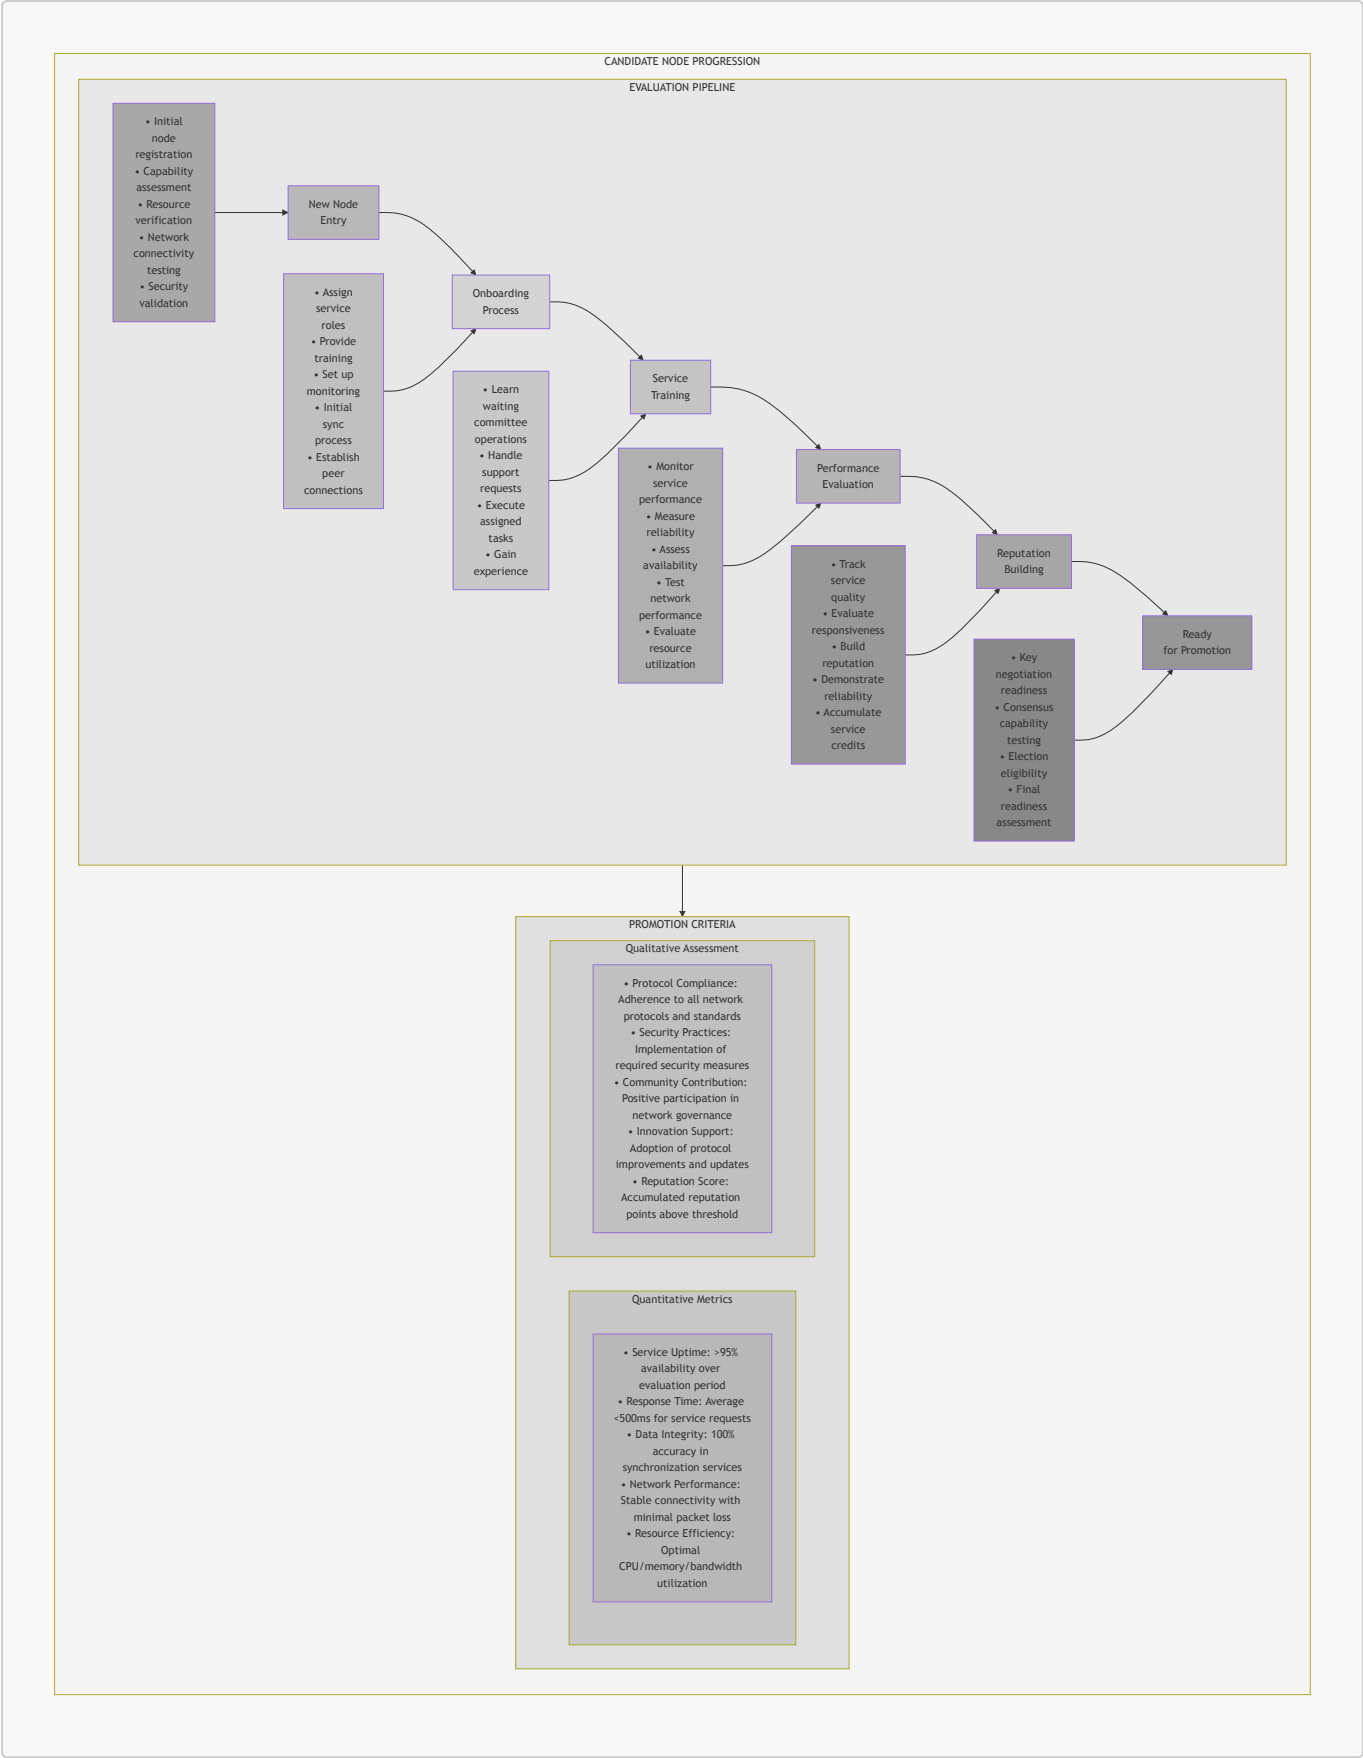
Data Serving Capabilities:

The waiting committee provides sophisticated data serving services that enable various blockchain applications and analysis tools:

- Transaction History Queries: Fast lookup of historical transactions by various criteria
- Account State History: Historical account balance and state information
- Block Range Requests: Efficient serving of block ranges for analysis applications
- Statistical Data: Pre-computed statistics and analytics for common queries
- API Services: RESTful and GraphQL APIs for external application integration

5.3.2.3. Candidate Node Management

The waiting committee includes a sophisticated candidate management system that prepares nodes for potential promotion to the consensus committee while maintaining their service responsibilities.



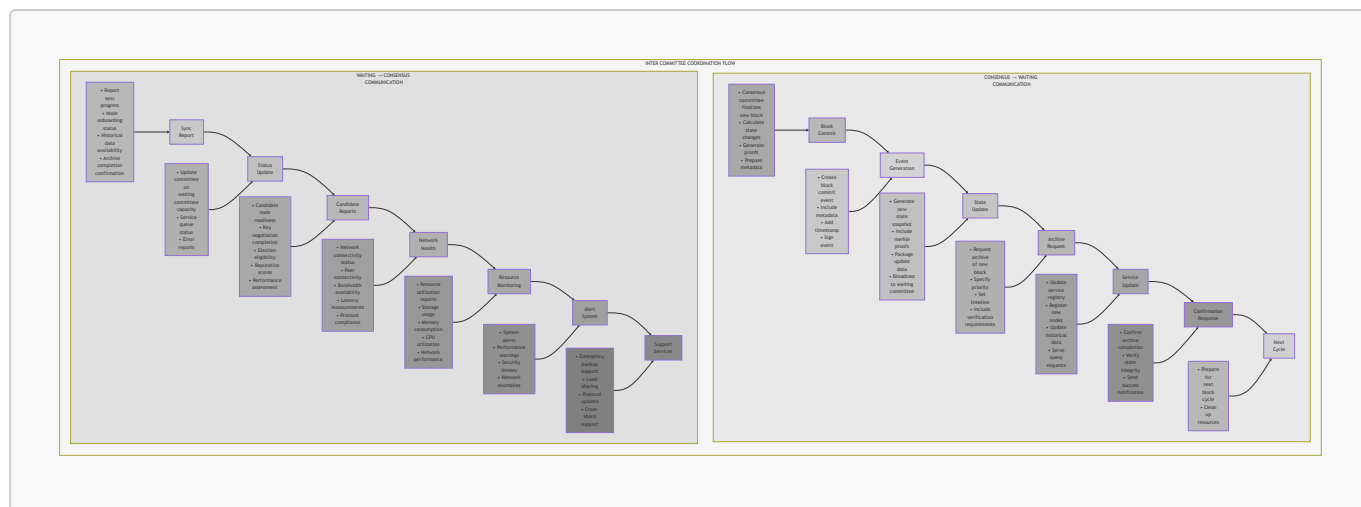
5.3.3. Committee Coordination

The coordination mechanism between consensus and waiting committees represents a critical architectural component that ensures seamless operation while maintaining the performance isolation that makes the dual committee structure effective.

Communication Architecture:

The coordination system implements asynchronous communication patterns that allow both committees to operate independently while sharing essential information:

- Event-Driven Messaging: Asynchronous event notifications that don't block committee operations
- Shared State Management: Efficient state sharing mechanisms with minimal overhead
- Priority Communication Channels: Different communication priorities for various types of coordination
- Fault-Tolerant Protocols: Robust communication that continues functioning during network issues
- Load-Balanced Distribution: Intelligent message routing to prevent communication bottlenecks



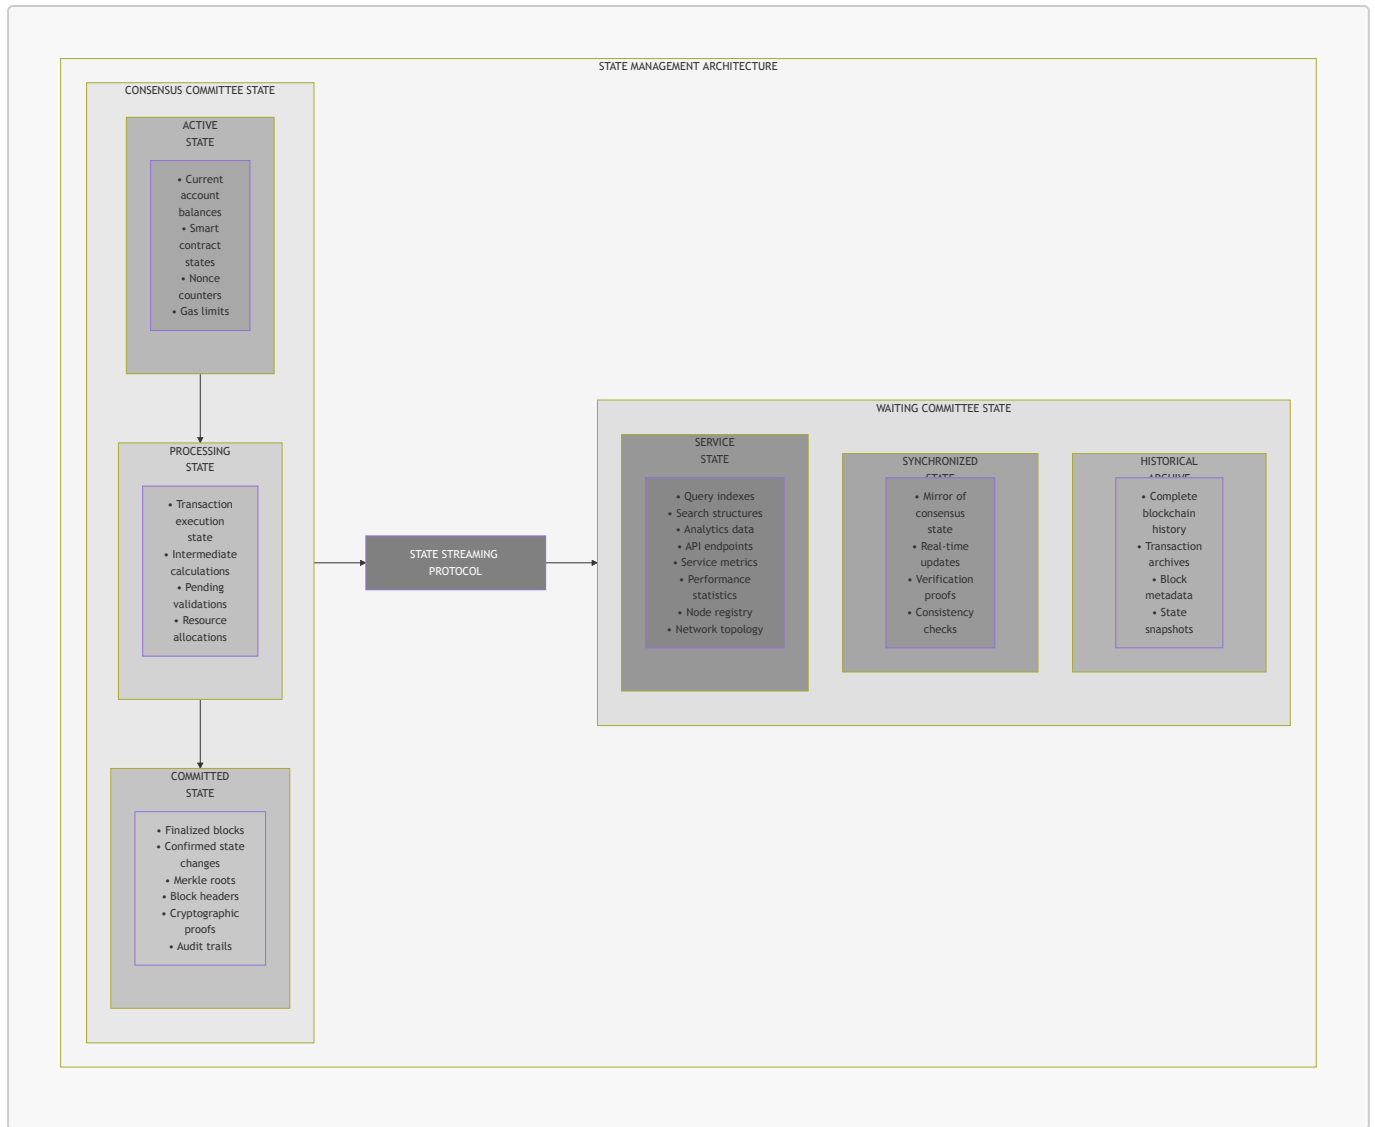
5.3.3.1. Coordinated State Management

The dual committee structure requires sophisticated state management to ensure both committees maintain consistent views of the blockchain state while allowing independent operation.

State Synchronization Strategy:


The coordination system implements multi-layered state management that balances consistency with performance:

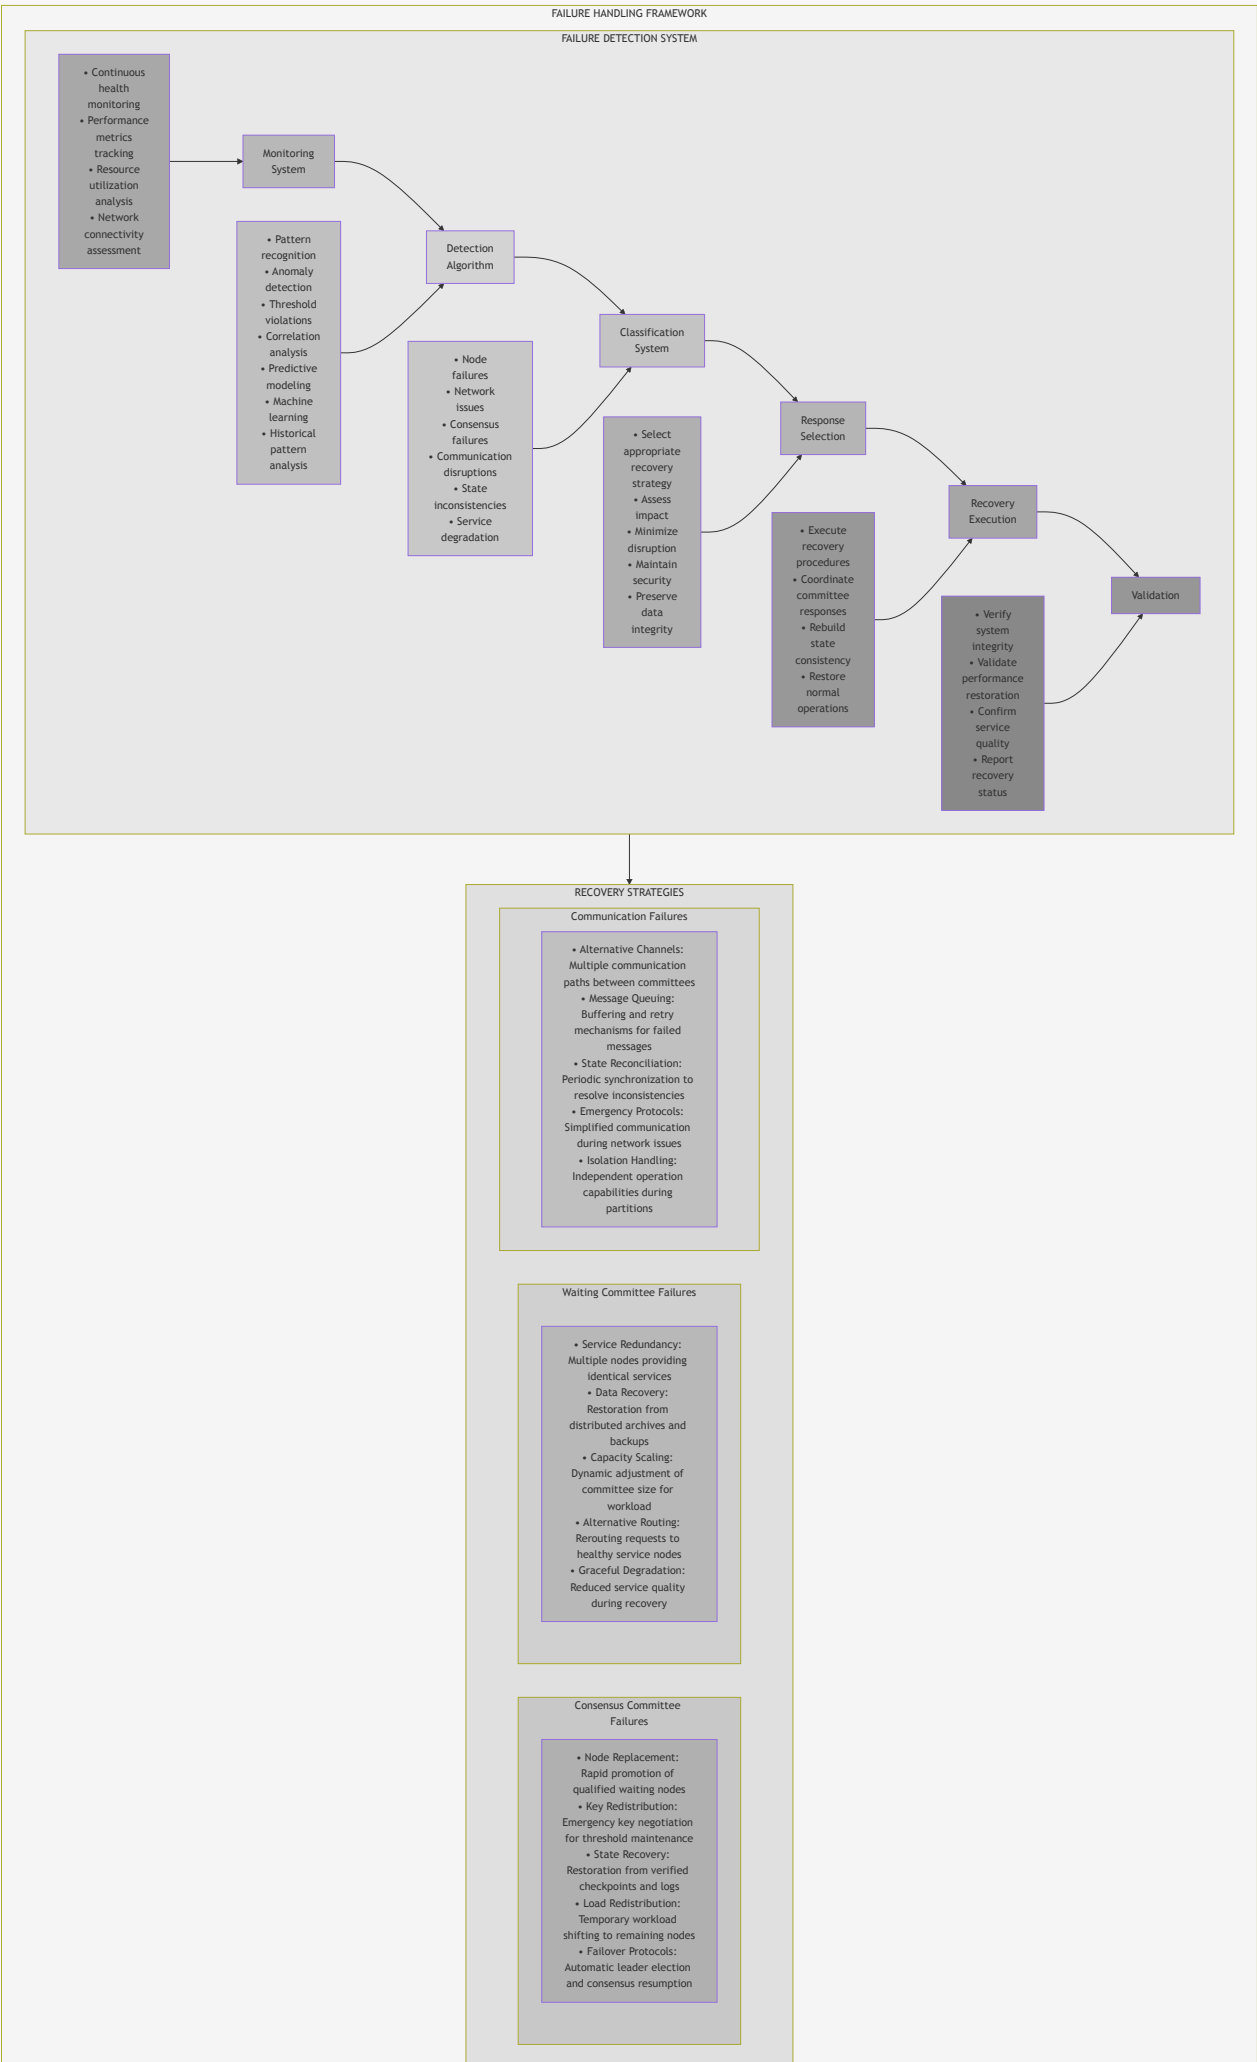
- Real-time State Streaming: Consensus committee streams state changes to waiting committee with minimal delay
- Checkpoint Synchronization: Periodic state checkpoints ensure long-term consistency across committees
- Differential Updates: Incremental state changes reduce synchronization overhead
- State Verification: Cryptographic verification ensures state integrity across committee boundaries
- Rollback Coordination: Coordinated rollback mechanisms handle consensus failures or reorganizations



5.3.3.2. Failure Handling and Recovery Coordination

The dual committee structure includes comprehensive failure handling mechanisms that ensure system resilience while maintaining the performance benefits of functional separation.





5.4. Consensus Mechanism

- Refer into [LoraBFT](#)

5.5. Transaction Lifecycle

The transaction lifecycle in a sharded blockchain like LogaShard (drawing from Shardeum's detailed process) involves a sophisticated series of steps from user initiation to final state update, designed for efficiency and integrity. The network primarily counts user transactions for its TPS metric, excluding internal network maintenance transactions.²⁴

5.5.1. Transaction Initiation and Submission

The process begins when a user initiates a transaction (e.g., a coin transfer or smart contract interaction) typically via a wallet like MetaMask. This request is sent to an RPC (Remote Procedure Call) server, which acts as the standard interface for submitting transactions and handling network queries, connecting the user's wallet to the network.

The RPC server then fetches the transaction's nonce (a sequential number to maintain order) and estimates the required gas. This gas estimation involves simulating the transaction on a local data RPC server with all relevant account data to predict computational resources. For smart contracts, this includes pre-running them to tally opcodes and their gas for a detailed estimate. A "nonce mode" feature enhances handling by queuing transactions with higher-than-expected nonces, processing them once preceding nonces are cleared, rather than immediate rejection. Once nonce and gas details are confirmed, the user signs the transaction with their private key, and the signed transaction is submitted from the wallet to the RPC server, which then injects it into the validator network.

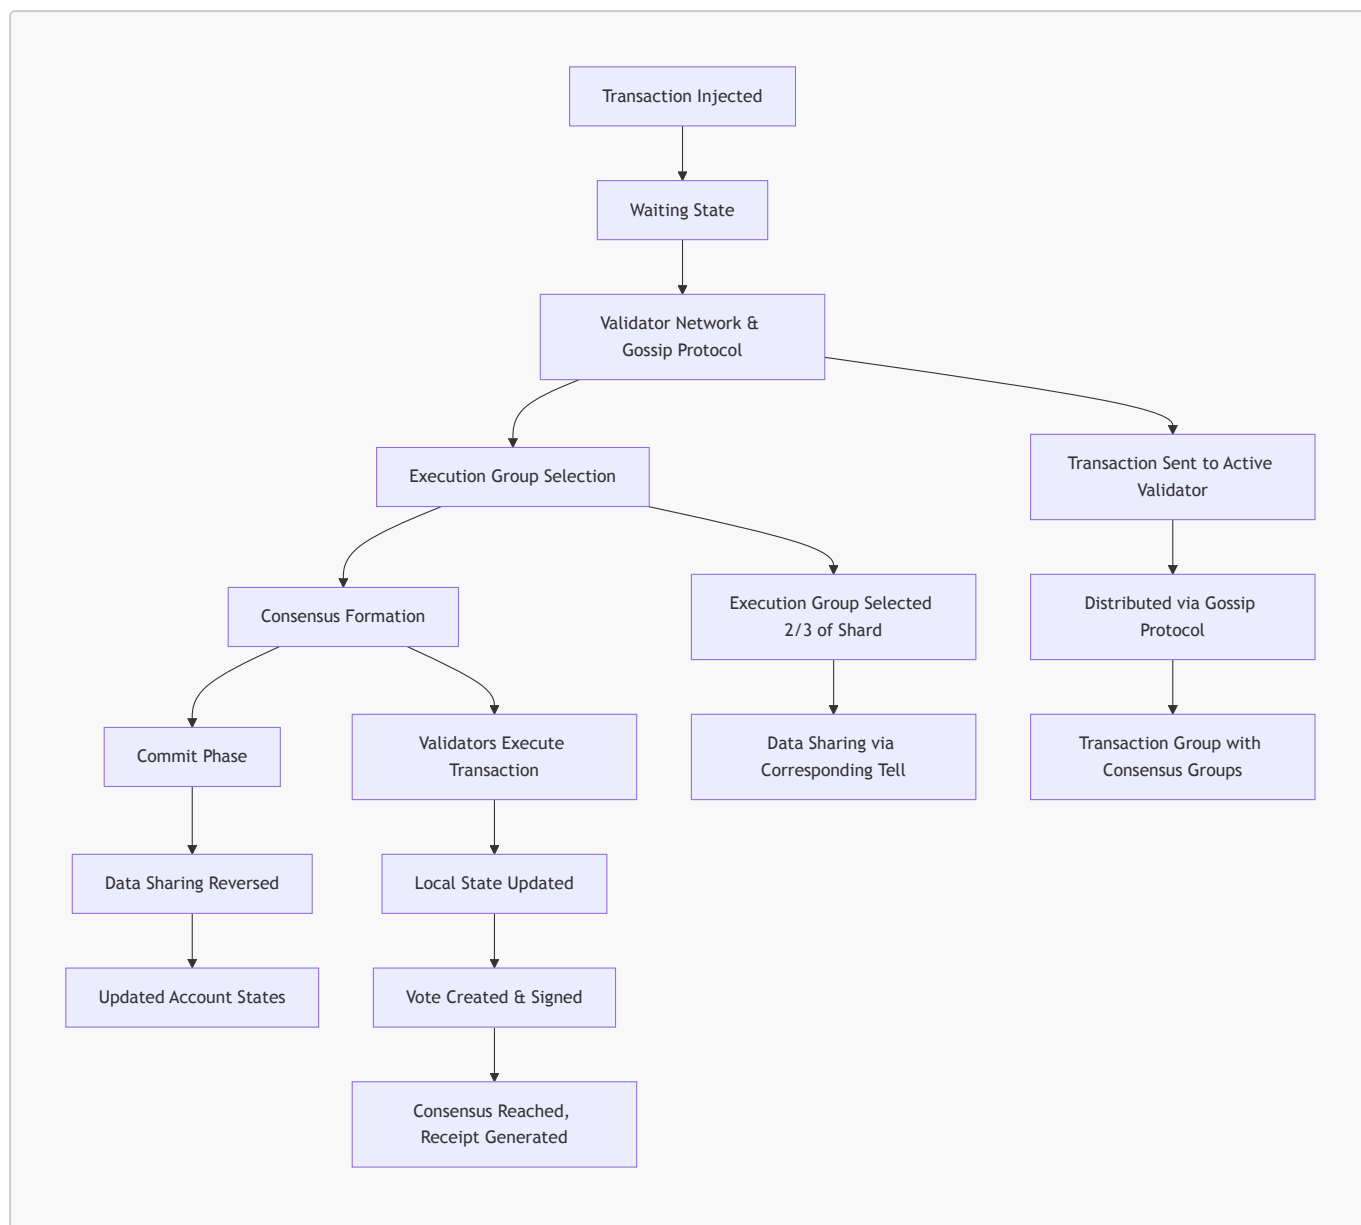
5.5.2. Transaction Processing

A transaction progresses through a state machine-like lifecycle, initially entering a waiting state for approximately six seconds after injection. This waiting period allows the transaction to mature and be queued by multiple validators, ensuring correct ordering based on timestamps.

- **Validator Network and Gossip Protocol:** The transaction is first sent to one active validator, then distributed to other validators within its "transaction group" using the gossip protocol. A transaction group comprises multiple "consensus groups," each responsible for a specific account or data involved in the transaction. For example, if a transaction involves accounts A, B, and C, there will be consensus groups for each, collectively forming the transaction group. Validators within these groups collaborate, each holding a part of the necessary transaction data. The gossip protocol optimizes communication by propagating messages efficiently to a few nodes, which then spread the message further, reducing the load on any single node.
- **Execution Group Selection:** Before queuing, an "execution group" is selected, typically comprising 2/3 of the shard size, deemed sufficient to determine a transaction's value. However, the executing shard may not have all necessary information for accounts on other shards, necessitating data sharing. Validators use a "corresponding tell" process to efficiently share

required data, ensuring data is forwarded only to specific nodes as defined by sharding logic, preventing undue influence.

- **Consensus Formation:** Validators execute the transaction, apply it to their local state (an in-memory copy), and form a consensus on the transaction's result. This process requires execution group validators to receive all necessary data, sometimes from non-execution group nodes, ensuring no upstream transactions could modify the state before processing. Validators create a vote by hashing the transaction proposal, cryptographically signing it, and sharing this signed hash with other nodes. If a majority of nodes agree and have the same hash, the transaction is confirmed, producing a provable receipt. Once a supermajority of votes is achieved, a valid signed receipt is produced, and any data owned by the node related to the transaction can be committed (saved).
- **Commit Phase:** After the receipt is created and confirmed, validators commit the receipt. This involves reversing the data sharing process, where nodes that received account data for execution send updated data back to the original nodes, ensuring all nodes in the transaction group have the correct and updated account states.



5.5.3. Data Distribution and State Updates

The confirmed transaction receipt is then sent to an "archiver," a specialized node responsible for long-term storage and data integrity.

- **Archiver Setup:** The archiver approves the receipt, writes it to a data log for inter-process communication, and stores it in a database optimized for storage (though not necessarily for queries). Ensuring valid receipts reach an Archive Server is crucial to prevent serious issues.
- **Data Forwarding and Syncing:** The archiver forwards the receipt to a "distributor," which disseminates the transaction data to a "collector" using WebSockets for efficient, real-time updates.²⁴ The collector then sends data back to the local data RPC server, which notifies the user's wallet (e.g., MetaMask) that the transaction is complete. The archiver also forwards data to an explorer collector, updating the blockchain explorer for public verification.
- **State Management:** Validators maintain only the latest state for optimized performance, while archivers store historical state data, ensuring efficient access to current data while preserving auditability. This separation allows for flexible topologies, where one archiver can connect to multiple collectors, or a collector to several distributors, ensuring effective scaling and addressing potential mainnet overloads.

This intricate transaction lifecycle, supported by dynamic sharding and a hybrid consensus, underpins the network's ability to handle high transaction volumes while maintaining decentralization and security. LogaShard's specific innovations are built upon these foundational principles, addressing the points of failure during dynamic reconfiguration.

6. Benchmark

LogaShard's performance needs to be understood in the context of existing blockchain throughput. Traditional blockchains like Bitcoin and early Ethereum have low TPS (7 and 15-30 respectively), while even with improvements, Ethereum reaches around 173-607 TPS. Sharded blockchains aim for much higher, with Polkadot achieving 1000 TPS per parachain and RapidChain 4220 TPS. This is still far from Visa's 24,000 TPS, highlighting the need for innovations like LogaShard to bridge this gap and enable mainstream Web3 adoption.

Shardeum, a sharded Layer 1, provides a benchmark. Its Alphanet reached 100 TPS, and Betanet (Sphinx) improved to 150-250 TPS. The underlying Shardus protocol demonstrated 5000 TPS with 1000 nodes, showcasing true linear scaling. Shardeum also achieved a record 171,000 testnet validators, emphasizing scalability through decentralized node operation, rather than vertical scaling. This linear scalability, where adding more nodes increases throughput, aligns with decentralization principles and sets a strong precedent for evaluating LogaShard's claims of high TPS and parallelism.

LogaShard's Achieved Throughput and Latency Benchmarks

LogaShard achieves **8300 TPS** on a **single shard (600 nodes, LAN conditions)**, significantly reducing **ledger synchronization and key negotiation overhead** by **90%**. This innovation tackles **TPS-Degradation** and **Zero-TPS** issues, ensuring **continuous transaction processing**.

While highly optimized, its **LAN-based benchmark** may not fully reflect real-world **WAN challenges** like latency and bandwidth constraints. Further validation is needed to assess **performance in decentralized networks**.

Table 1: Comparative TPS Benchmarks

Blockchain/System	TPS (Transactions Per Second)	Context/Conditions	Key Features/Mechanism
Bitcoin	7	Mainnet	Proof of Work (PoW)
Ethereum (Pre-Merge)	15-30	Mainnet	Proof of Work (PoW)
Ethereum (Post-Merge/Blobs)	173.6 - 607	Mainnet (using blobs/calldata)	Proof of Stake (PoS)
Polkadot (Parachain)	~1000	Single Parachain	Nominated Proof of Stake (NPoS), Sharding
RapidChain	4220	Sharded blockchain	Sharding
Shardeum (Alphanet 2.1)	100	Testnet	Dynamic State Sharding, PoS + PoQ
Shardeum (Betanet Sphinx 1.1)	250	Testnet (1280 Validators)	Dynamic State Sharding, PoS + PoQ
Shardeum (Linear Scaling)	5000	Testnet (1000 Nodes)	Dynamic State Sharding, PoS + PoQ, Linear Scaling
LogaShard	8300	Single Shard, 600 Nodes, LAN Conditions	Parallelized Dual Committee, Key Pre-negotiation, Sharding

7. Current Challenges

Despite the innovative solutions proposed by LogaShard and similar sharding architectures, several challenges persist in the practical implementation and optimization of sharded blockchains. These challenges are often inherent to the complexity of distributed systems and the dynamic nature of sharding.

7.1. Scalability Challenges of Cross-Shard Transactions

Cross-shard transactions (CTXs) continue to be a primary bottleneck for sharded blockchains. While sharding aims to parallelize transaction processing, CTXs inherently require coordination across multiple shards, introducing complexity and potential delays.

The overhead associated with CTXs is significant. Protocols like two-phase commit (2PC), often used to ensure atomicity, involve multiple rounds of communication and consensus across shards, leading to increased latency and reduced throughput. Some sharding schemes may also record CTXs redundantly across involved shards, increasing storage and communication burdens. As a network scales, the proportion of CTXs can become very high, potentially dominating the overall transaction volume. If not efficiently managed, this can negate the scalability benefits of sharding, causing the network to experience congestion and higher fees, similar to unsharded blockchains.

Ensuring atomicity and consistency across shards for CTXs remains difficult, especially when transactions depend on data from multiple shards. Conflicts may arise if two shards attempt to update the same data simultaneously. Specialized protocols are required to manage these transactions, which complicates the overall architecture and necessitates robust mechanisms to handle failures and retries in inter-shard communications. The challenge lies in balancing the efficiency of parallel processing with the need for strong consistency guarantees across a fragmented state.

7.2. Data Availability Challenges in Sharded Blockchains

Ensuring "data availability" is paramount for the integrity and security of sharded blockchains. In a sharded system, where individual nodes typically store only a portion of the total state, guaranteeing that all necessary transaction data is accessible for verification by any participant becomes complex.

One significant challenge is the potential for "data withholding attacks". Malicious actors might intentionally fail to share critical data, leading to inconsistencies across the network and compromising its trustworthiness. This is particularly problematic in sharded environments where a node might need data from another shard to validate a transaction.

Furthermore, as transaction volumes and historical data accumulate, "storage bloat" becomes a concern. Requiring every node to store all data becomes expensive and burdensome, especially for less powerful devices, which can inadvertently lead to centralization if only nodes with significant resources can keep up. This also increases "verification overhead," as processing every transaction becomes slower with larger datasets, impacting network speed and throughput.

Addressing these issues often involves implementing "Data Availability Layers (DALs)" or "Data Availability Sampling (DAS)". DALs separate data availability from other blockchain functions, storing data reliably off-chain while ensuring verifiability. DAS allows light nodes to verify data availability without downloading the entire dataset. While these solutions improve efficiency, they add architectural complexity and introduce new challenges related to their own security and decentralization.

7.3. Impact of Validator Set Size and Churn on Sharded PoS Performance

The dynamics of the validator set in a sharded Proof-of-Stake (PoS) network present several challenges to performance and decentralization.

A large validator set, while theoretically enhancing decentralization, can strain peer-to-peer networking and messaging, potentially causing node failures due to high computational load and bandwidth requirements. Every new validator adds an additional connection, increasing the overall bandwidth needed to maintain consensus. As the network grows, validator node operators may require more sophisticated hardware to support larger bandwidth and internet speeds, which could lead to centralization if smaller, self-hosted nodes are unable to keep pace.

"Validator churn"—the rate at which validators enter and exit the network—also impacts performance. Frequent reconfigurations and validator shuffling, while crucial for security against shard takeover attacks, introduce overhead. Each time validators are shuffled, new committees must be formed, and nodes must synchronize their state, potentially leading to the TPS-Degradation issue discussed earlier. While LogaShard's dual committee framework aims to mitigate this, the underlying challenge of efficiently managing a dynamic and large validator set remains. Future upgrades that require validators to send and receive more data (e.g., proto-danksharding on Ethereum) could exacerbate latency issues

between nodes, further contributing to centralization risk as only high-performance setups can participate effectively.

7.4. Challenges in Implementing Reputation Systems

LogaShard's reliance on a reputation mechanism within its dual committee framework introduces its own set of challenges. While reputation systems can enhance security by incentivizing honest behavior and penalizing malicious actions, their implementation in decentralized environments is complex.

- **Scalability:** Calculating and updating reputation scores for a large and dynamic validator set in real-time, especially across multiple shards, can be computationally intensive and impact network performance.
- **Privacy:** Balancing the need for transparency in reputation calculations with user privacy is essential. Revealing too much information about node behavior could compromise privacy, while insufficient data might make reputation scores less reliable.
- **Sybil Attacks:** Preventing "Sybil attacks," where malicious actors create numerous fake identities to manipulate reputation scores, requires robust identity verification mechanisms and sophisticated algorithms to detect and mitigate such attempts.
- **Regulatory Compliance:** Ensuring that reputation systems comply with various data protection laws and consumer rights regulations is crucial for widespread adoption, adding a legal and ethical layer of complexity.

The design of a fair, resilient, and scalable reputation system that accurately reflects validator trustworthiness without introducing new vulnerabilities or centralizing control is an ongoing area of research and development in decentralized systems.

8. Limitations

While LogaShard presents compelling solutions to critical scalability challenges, it is important to acknowledge the inherent limitations and potential trade-offs within its architecture and the broader sharding paradigm.

8.1. Inherent Limitations of Sharding

Despite its promise, sharding fundamentally introduces complexities that must be meticulously managed.

- **Increased Complexity of Implementation:** Designing and implementing a sharded blockchain is inherently more complex than a monolithic one. This complexity extends to the core protocol, the consensus mechanism, and the application layer, requiring sophisticated algorithms for shard creation, node assignment, and cross-shard communication. This can lead to longer development cycles and a higher potential for subtle bugs or vulnerabilities.
- **Security Vulnerabilities (Shard Takeover Risk):** As discussed, sharding divides the network's security budget across smaller validator sets per shard. While LogaShard's reputation mechanism and auto-rotation aim to mitigate this, the fundamental risk of a shard takeover attack remains a theoretical concern if an adversary can concentrate sufficient resources or exploit a flaw in the shuffling mechanism. The challenge is to ensure that the randomization and rotation are truly unpredictable and that the cost of attacking even a single shard remains prohibitively high.
- **Coordination Overhead:** Even with optimized protocols, the need for coordination between shards for cross-shard transactions introduces overhead and potential latency. While LogaShard aims for

atomic composability, the underlying communication and consensus across shard boundaries will always be more resource-intensive than intra-shard operations. This implies that the network's overall efficiency is still influenced by the proportion and complexity of cross-shard interactions.

8.2. Benchmark Conditions and Real-World Applicability

LogaShard's impressive peak throughput of 8300 TPS was achieved "on a single shard with 600 nodes under LAN conditions". This specific context is a crucial limitation when extrapolating these results to real-world, globally distributed blockchain deployments.

- **Network Latency:** LAN environments inherently have extremely low latency and high bandwidth, which are ideal for distributed consensus and data synchronization. In a global Wide Area Network (WAN) setting, geographical distances, internet infrastructure variations, and network congestion introduce significant latency. This increased latency directly impacts the speed of message propagation, consensus formation, and state synchronization across geographically dispersed nodes, potentially reducing the achievable TPS.
- **Hardware Homogeneity:** Benchmarks often assume relatively homogeneous and high-performance hardware for all participating nodes. In a truly permissionless public blockchain, validator nodes will operate on a wide range of hardware specifications and network connections. This heterogeneity can introduce performance disparities and bottlenecks, as slower nodes might struggle to keep up, impacting the overall network's efficiency and potentially compromising decentralization.
- **Adversarial Conditions:** Benchmarks, while rigorous, may not fully capture the complexities of a live network under sustained adversarial attacks (e.g., denial-of-service attacks, data withholding, or Sybil attacks). The true resilience of LogaShard's parallelized mechanisms would be tested under such real-world, less controlled conditions.

These factors suggest that while LogaShard's architectural innovations demonstrate significant theoretical and experimental promise, their performance in a fully decentralized, globally distributed mainnet environment may differ from the reported LAN benchmarks. Further research and real-world deployments will be necessary to quantify these differences and identify additional optimizations.

8.3. Complexity of Dual Committee and Key Management

The very solutions that make LogaShard robust also introduce their own layers of complexity.

- **Dual Committee Management:** Operating and coordinating two parallel committees (consensus and waiting) within each shard, coupled with a reputation mechanism and dynamic shuffling, adds significant complexity to the protocol's state management and operational logic. Ensuring seamless transitions between committees and accurate reputation scoring requires sophisticated algorithms and robust implementation.
- **Key Management for Pre-negotiation:** While parallelized key pre-negotiation avoids Zero-TPS, it necessitates a highly secure and efficient key management system. Generating, distributing, and securely storing cryptographic keys for future epochs, especially with a secret-reuse strategy, introduces intricate security challenges. Any compromise in this key management process could have severe implications for network security. The system must ensure that the pre-negotiated keys remain secure and are activated precisely when needed, without introducing new attack vectors.

These complexities highlight that LogaShard's innovations, while powerful, demand meticulous engineering and continuous auditing to ensure their long-term stability and security in a live network environment.

Conclusion

The analysis of LogaShard's architecture reveals a focused and innovative approach to overcoming critical scalability limitations in sharded blockchain systems. The pervasive "blockchain trilemma" underscores the inherent difficulty in simultaneously achieving scalability, security, and decentralization. While sharding offers a promising pathway to increased throughput, it introduces its own set of complexities, notably the "TPS-Degradation" and "Zero-TPS" issues that arise during dynamic shard reconfigurations.

LogaShard directly confronts these challenges through two primary architectural innovations: the **parallelized dual committee framework** and the **parallelized key pre-negotiation mechanism with a secret-reuse strategy**. The dual committee framework effectively mitigates TPS-Degradation by allowing ledger synchronization for newly assigned nodes to occur in a "waiting committee" parallel to active transaction processing by a "consensus committee." This prevents the "cold start" problem from halting or significantly degrading the network's throughput. The integration of a reputation mechanism further enhances security by ensuring reliable node participation and making shard takeover attacks more difficult.

The parallelized key pre-negotiation mechanism directly addresses the Zero-TPS issue by enabling the cryptographic key generation for future epochs to run concurrently with the current epoch's transaction processing. This foresight eliminates the need for network pauses during key negotiation, ensuring continuous network availability and maintaining high TPS. The secret-reuse strategy further optimizes this process by reducing computational overhead. These mechanisms are not merely empirical fixes but are supported by claims of "theory-guaranteed security," indicating a robust cryptographic foundation.

LogaShard's achieved benchmark of 8300 TPS on a single shard with 600 nodes under LAN conditions demonstrates a significant leap in performance, outperforming state-of-the-art sharding schemes in reducing reconfiguration overheads by at least 90%. This performance, while impressive, necessitates further validation under diverse, real-world WAN conditions to fully assess its practical scalability.

Despite these innovations, sharded blockchains, including LogaShard, continue to face challenges such as the complexity and overhead of cross-shard transactions, the "hot shard problem" due to imbalanced transaction distribution, data availability concerns, and the inherent complexities of managing large, dynamic validator sets and reputation systems.

The future trajectory for LogaShard involves continuous research into extending parallelism, exploring new cryptographic primitives, and ensuring robust security through formal verification. Ecosystem development, including enhanced developer tooling and interoperability solutions, will be crucial for broader adoption. Ultimately, by effectively addressing the specific bottlenecks of TPS-Degradation and Zero-TPS, LogaShard aims to contribute significantly to the ongoing quest for truly scalable, secure, and decentralized blockchain networks, paving the way for the mass adoption of Web3 applications that demand high throughput and uninterrupted service.